

Multithreaded Programming




Shing-Chi Cheung
Computer Science and Engineering
HKUST

Process and Threads

Multiple threads on multiple CPU cores



Multiple threads sharing a single CPU core



A **process** is a program in execution, which may consist of multiple concurrent execution units known as **threads**.

Threads in a process share the **same** memory heap space.

Threads need to be **synchronized** if they have dependency.

Creating Tasks and Threads (6 important steps)

implements

`java.lang.Runnable`  `TaskClass`

```
// Custom task class
public class TaskClass implements Runnable {
    ...
    public TaskClass(...) {
        ...
    }

    // Implement the run method in Runnable
    public void run() {
        // Tell system how to run custom thread
        ...
    }
    ...
}
```

(a)

```
// Client class
public class Client {
    ...
    public void someMethod() {
        ...
        // Create an instance of TaskClass
        TaskClass task = new TaskClass(...);

        // Create a thread
        Thread thread = new Thread(task);

        // Start a thread
        thread.start();
        ...
    }
    ...
}
```

(b)

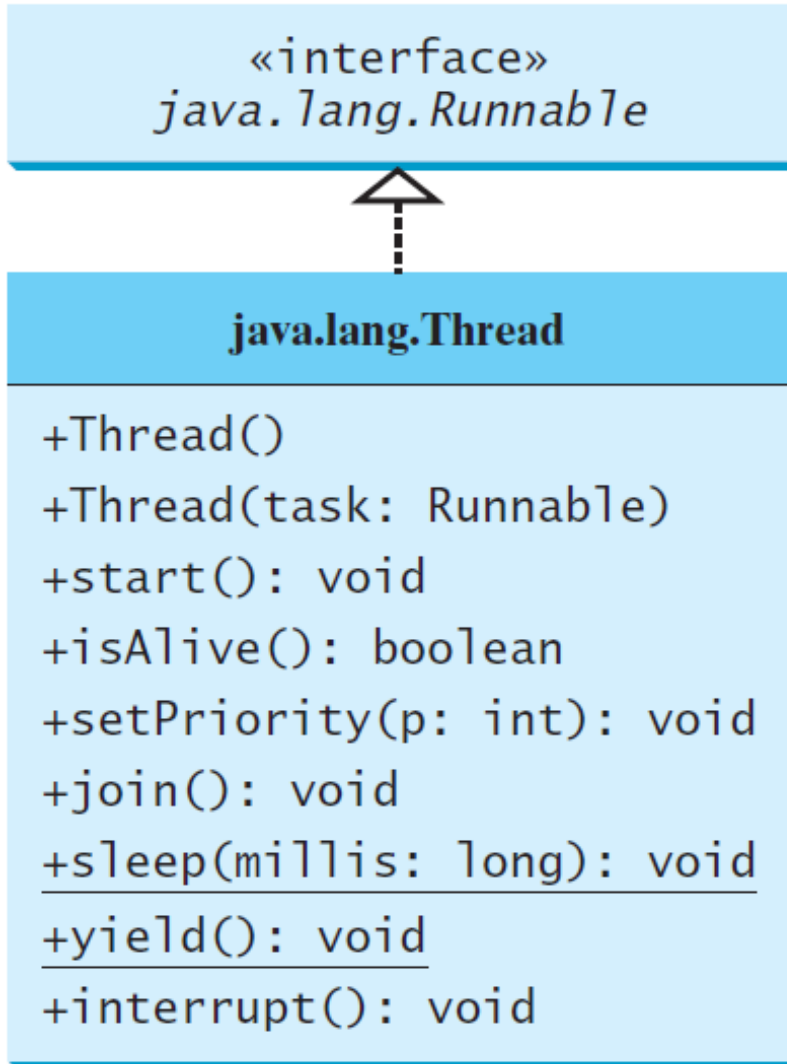
Example: Using the Runnable Interface to Create and Launch Threads

- Objective: Create and run three threads:
 - The first thread prints 'a' 20 times
 - The second thread prints 'b' 20 times
 - The third thread prints 'c' 20 times

```
public void run() {  
    for (var i = 0; i < times; i++) {  
        System.out.print(charToPrint);  
    }  
}
```

TaskThreadDemo.java

The Thread Class



Creates an empty thread.

Creates a thread for a specified task.

Starts the thread that causes the `run()` method to be invoked by the JVM.

Tests whether the thread is currently running.

Sets priority `p` (ranging from 1 to 10) for this thread.

Waits for this thread to finish.

Puts a thread to sleep for a specified time in milliseconds.

Causes a thread to pause temporarily and allow other threads to execute.

Interrupts this thread.

The Static `yield()` Method

- We can use the `yield()` method to release time for other threads. For example, we can modify the `run` method as shown
- The thread yields every time after a character is printed. So, the thread runs relatively slower

```
public void run() {  
    for (var i = 0; i < times; i++) {  
        System.out.print(charToPrint);  
        Thread.yield();  
    }  
}
```

TaskThreadYieldDemo.java

The Static sleep(milliseconds) Method



- The sleep(long mills) method puts the thread to sleep for the specified time in milliseconds. For example, we can modify the run method as shown
- The thread sleeps 100 milliseconds every time after a character is printed

```
public void run() {  
    try {  
        for (var i = 0; i < times; i++) {  
            System.out.print(charToPrint);  
            // sleep 100 milliseconds  
            Thread.sleep(100);  
        }  
    } catch (InterruptedException ex) {  
    }  
}
```

TaskThreadSleepDemo.java

The join() Method

- We can use the join() method to force one thread to wait for another thread to finish. For example, the run method can be modified as shown
- After printing 10 characters, the thread pauses until completion of thread4

```
public void run() {  
    var thread4 = new Thread(  
        new PrintChar('d', 20));  
    thread4.start();  
    try {  
        for (var i = 0; i < times; i++) {  
            System.out.print(charToPrint);  
            if (i >= 10) thread4.join();  
        }  
    } catch (InterruptedException ex) {  
    }  
}
```

TaskThreadJoinDemo.java

Thread States

A thread can be in one of the following four major states

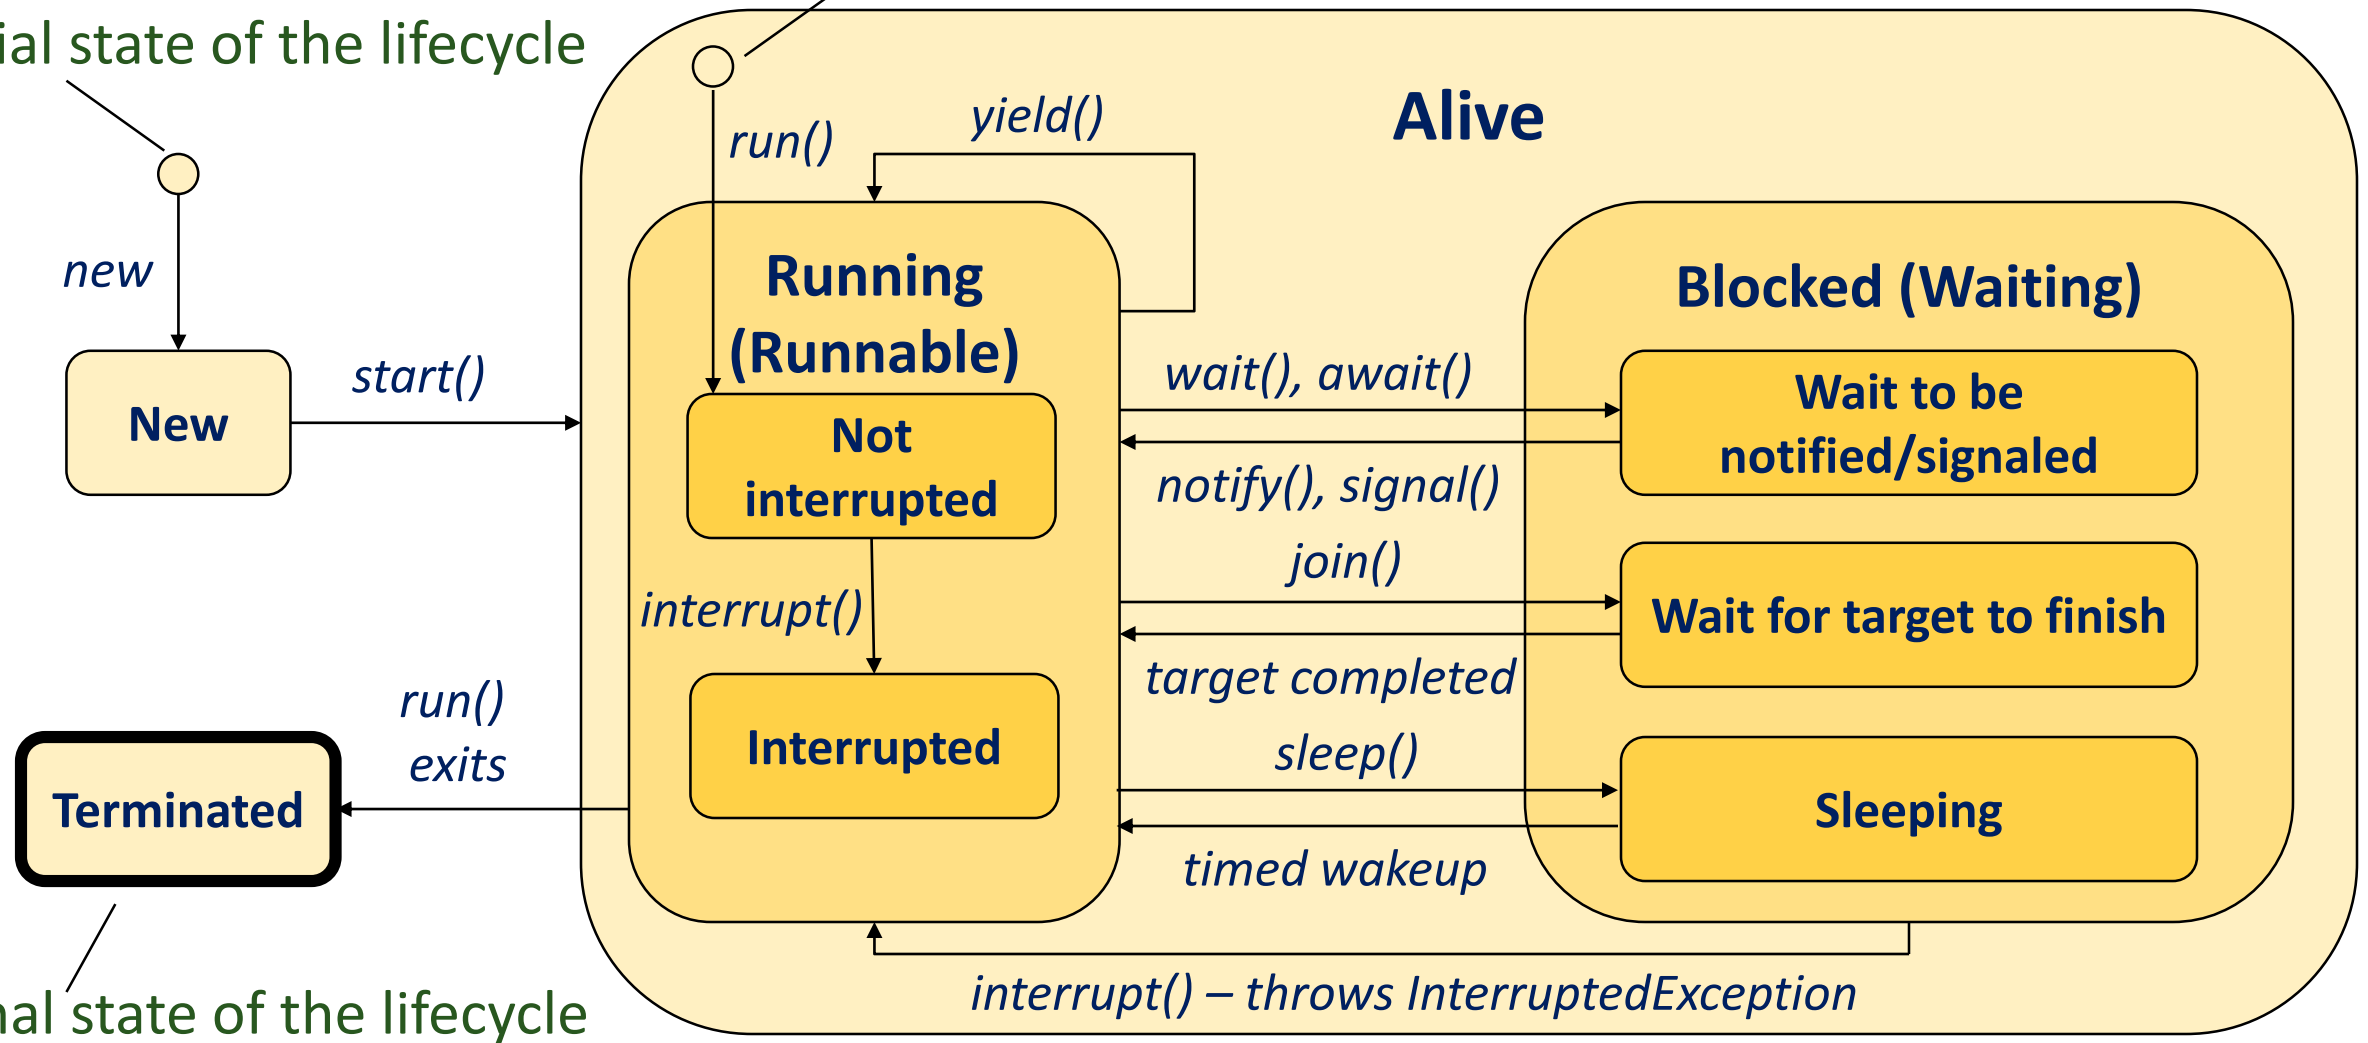
- ❑ **New**: thread is created but not yet started
- ❑ **Running**: thread is either being executed or to be scheduled for execution
- ❑ **Blocked**: thread is either blocked by monitor locks (to be discussed soon) or awaiting some events for a specified waiting time
- ❑ **Terminated**: thread exited

Thread Lifecycle



initial state of the lifecycle

initial state when a thread is alive; it is a transient state



final state of the lifecycle

isAlive(), interrupt(), and isInterrupted()

- **isAlive()** method returns true if the thread has started but not yet terminated; otherwise false
- **interrupt()** method interrupts a thread in the following way:
 - ☐ If the thread is in Running state, its interrupted flag is set
 - ☐ If the thread is in Blocked state; it enters Running state and throws `java.io.Interrupted Exception`
- The **isInterrupted()** method tests whether the current thread's interrupted flag has been set





The deprecated stop(), suspend(), and resume() Methods

- The Thread class also supports the stop(), suspend(), and resume() methods.
- As of Java 2, these methods are deprecated (or outdated) because they are known to be inherently unsafe.
- We should assign null to a Thread variable to indicate that it is stopped rather than use the stop() method.
- **Advice:** Never use these three methods

Thread Priority

- Each thread is assigned a **default priority** of `Thread.NORM_PRIORITY`
 - We can reset the priority using **`setPriority(int priority)`**
- Some constants for priorities include
 - `Thread.MIN_PRIORITY`
 - `Thread.MAX_PRIORITY`
 - `Thread.NORM_PRIORITY`

Example: Flashing Text



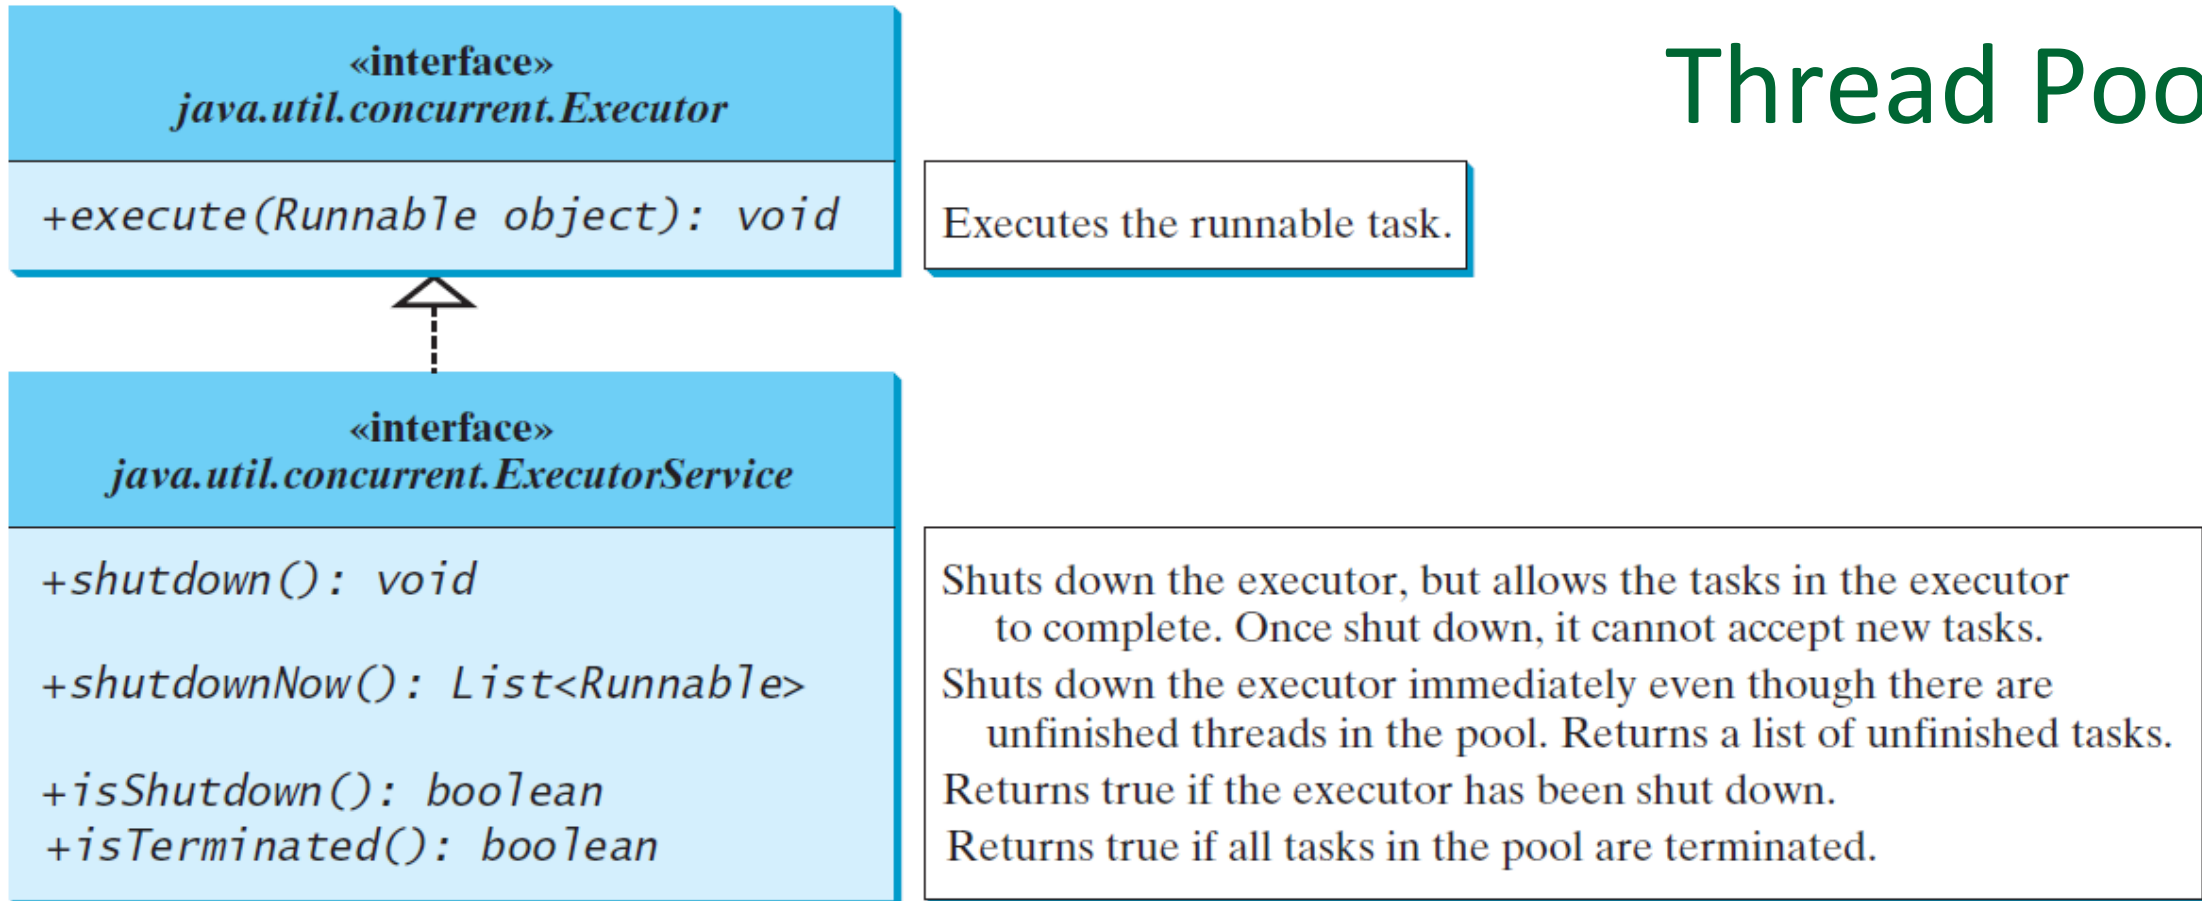
```
public void start(Stage primaryStage) {  
    var pane = new StackPane();  
    var lblText = new Label("COMP3021 is the best COMP course");  
    pane.getChildren().add(lblText);  
    new Thread( () -> {  
        try {  
            while (true) {  
                if ( lblText.getText().trim().length() == 0 )  
                    text = "COMP3021 is the best COMP course";  
                else  
                    text = "";  
                Platform.runLater( () -> lblText.setText(text) );  
                Thread.sleep(200);  
            }  
        }  
        catch (InterruptedException ex) {  
        }  
    }).start();  
}
```

Implementation of run()

FlashText.java

...

Thread Pools



- Starting a new thread for each task limits throughput and causes poor performance. A thread pool is ideal to manage the number of tasks executing concurrently. Use the **Executor** interface for executing tasks in a thread pool and the **ExecutorService** interface for managing and controlling tasks. **ExecutorService** is a subinterface of **Executor**.

Creating Executors

`java.util.concurrent.Executors`

```
+newFixedThreadPool(numberOfThreads:  
    int): ExecutorService
```

```
+newCachedThreadPool():  
    ExecutorService
```

Creates a thread pool with a fixed number of threads executing concurrently. A thread may be reused to execute another task after its current task is finished.


Creates a thread pool that creates new threads as needed, but will reuse previously constructed threads when they are available.

- To create an Executor object, use the static methods in the Executors class.

[ExecutorDemo.java](#)

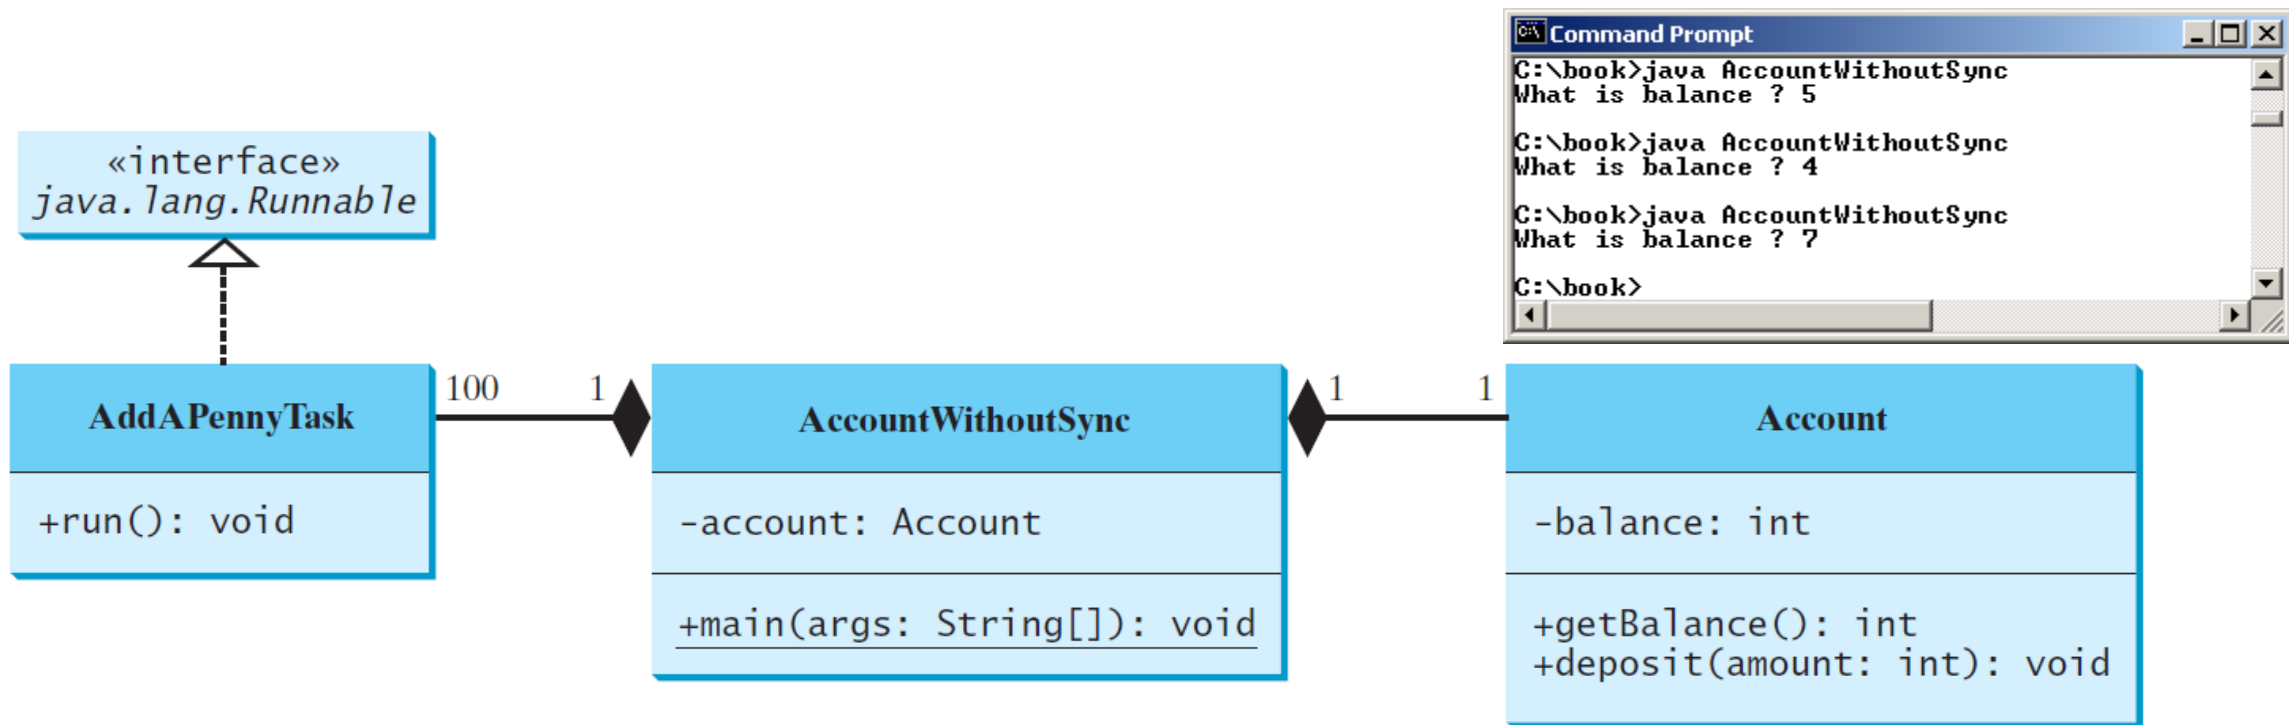
Thread Synchronization

Step	balance	thread[i]	thread[j]
1	0	<code>newBalance = bank.getBalance()+1;</code>	
2	0		<code>newBalance = bank.getBalance()+1;</code>
3	1	<code>bank.setBalance(newBalance);</code>	
4	1		<code>bank.setBalance(newBalance);</code>

- A shared resource may be corrupted if it is accessed simultaneously by multiple threads. 
- For example, two unsynchronized threads accessing the same bank account may cause conflict.

Example: Showing Resource Conflict

- We create and launch one hundred threads, each of which adds a penny to an account that is initially empty



A scenario where conflicts arise

Step	balance	Task 1	Task 2
1	0	newBalance = balance + 1;	
2	0		newBalance = balance + 1;
3	1	balance = newBalance;	
4	1		balance = newBalance;

- ❑ The effect of this scenario is that Task 1 did nothing, because in Step 4 Task 2 overwrites Task 1's result. The problem is that Task 1 and Task 2 are **accessing a common resource** in a way that causes conflicts
- ❑ This is a problem known as **race conditions** in multithreaded programs
- ❑ A class is said to be **thread-safe** if instances of the class **are free from race conditions in the presence of multiple threads**. Account class in the preceding example is **not thread-safe** AccountWithoutSync.java

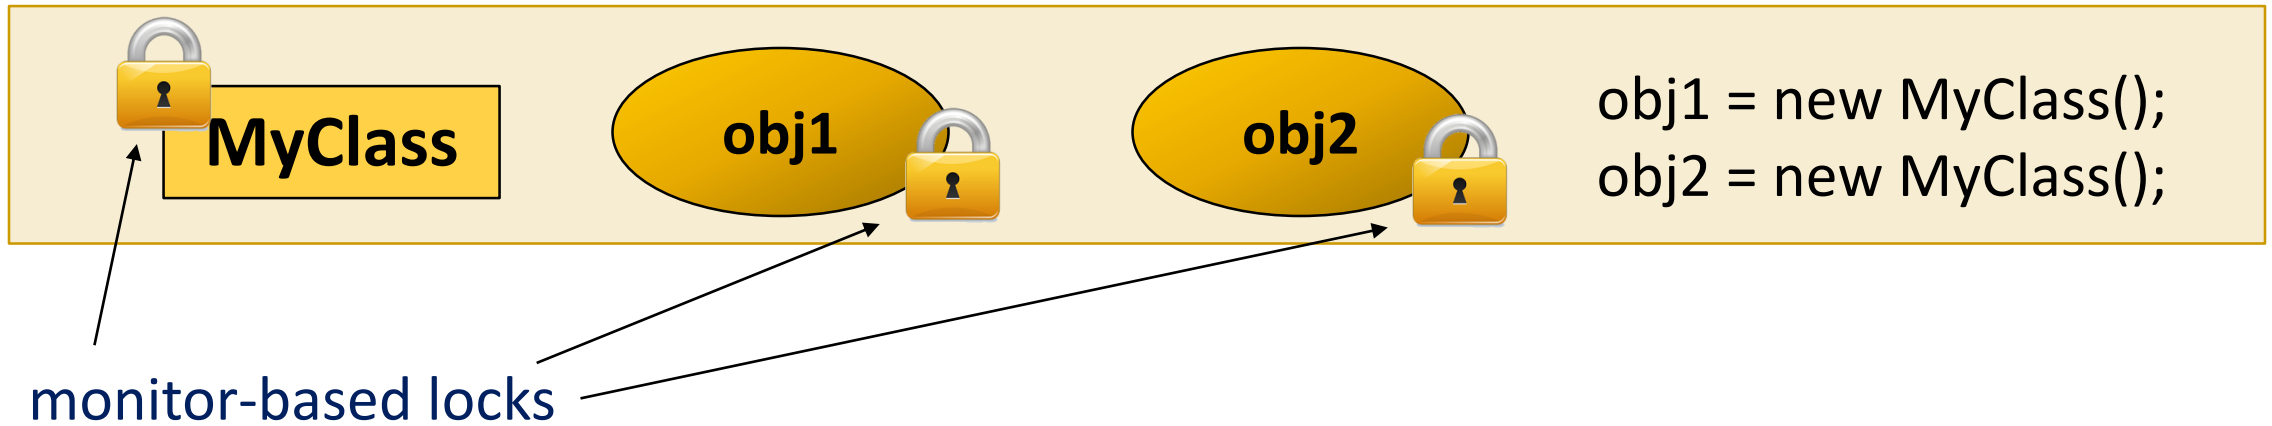
The **synchronized** keyword

- To avoid race conditions, only one thread is allowed to access certain part of the program, known as **critical region**
- The **critical region** is the entire deposit method. We can use the **synchronized** keyword to declare the entire method to be a critical region
- To correct the previous race problem, we can make Account **thread-safe** by adding the **synchronized** keyword in the deposit method:

```
public synchronized void deposit(double amount) { ... }
```

[AccountSync.java](#)

Thread Synchronization using Monitor-based Locks



- JVM assigns a **separate** lock to each class and each (object) instance



- We call them **monitor-based locks** or **implicit locks**

Synchronization using Implicit (Monitor-based) Lock



MyClass



```
obj1 = new MyClass();  
obj2 = new MyClass();
```

1. A thread needs to acquire a lock before it executes a synchronized method
 - ❑ For a synchronized instance method, it acquires the lock of the associated instance
 - ❑ For a synchronized static method, it acquires the lock of the associated class
2. If the thread successfully acquires the lock, it executes the synchronized method and releases the lock when it finishes the method. **The acquisition and release are performed implicitly.** No extra method calls are needed.
3. If the thread fails to acquire the lock, its execution is blocked
 - ❑ It repeats step 1 when the lock is released by another thread

Thread Synchronization using Monitor-based Locks

Step	balance	Task 1	Task 2
1	0	newBalance = balance + 1;	
2	0		newBalance = balance + 1;
3	1	balance = newBalance;	
4	1		balance = newBalance;

- With the deposit method synchronized, the above racing condition cannot happen
- If Task 2 starts to execute the deposit method before Task 1 finishes the method, Task 2 is blocked

Thread Synchronization using Monitor-based Locks

Task 1

Task 2

:Account
deposit()



Thread Synchronization using Monitor-based Locks



Task 1

Task 2

:Account
deposit()

implicit

Acquire a lock on the object account

Thread Synchronization using Monitor-based Locks



Task 1

Task 2

:Account
deposit()

implicit

Acquire a lock on the object account



Execute the deposit method



Wait to acquire the lock

Thread Synchronization using Monitor-based Locks

:Account
deposit()



Task 1

Task 2

implicit

Acquire a lock on the object account



Execute the deposit method

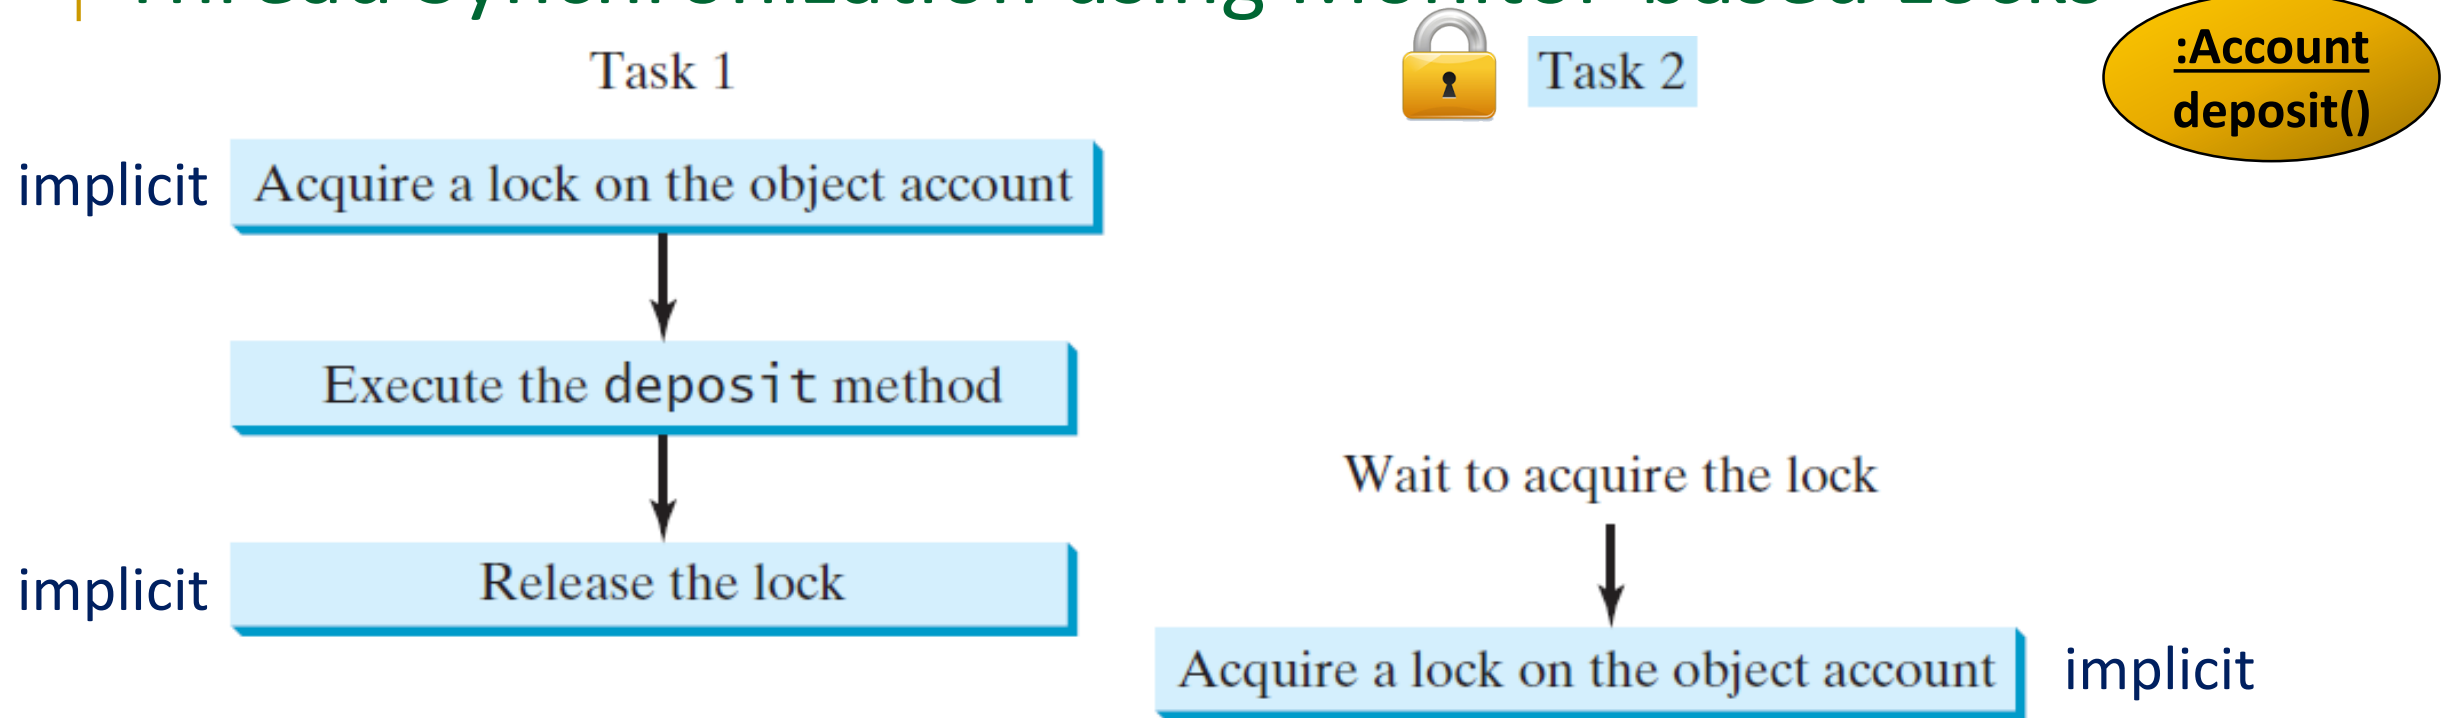


implicit

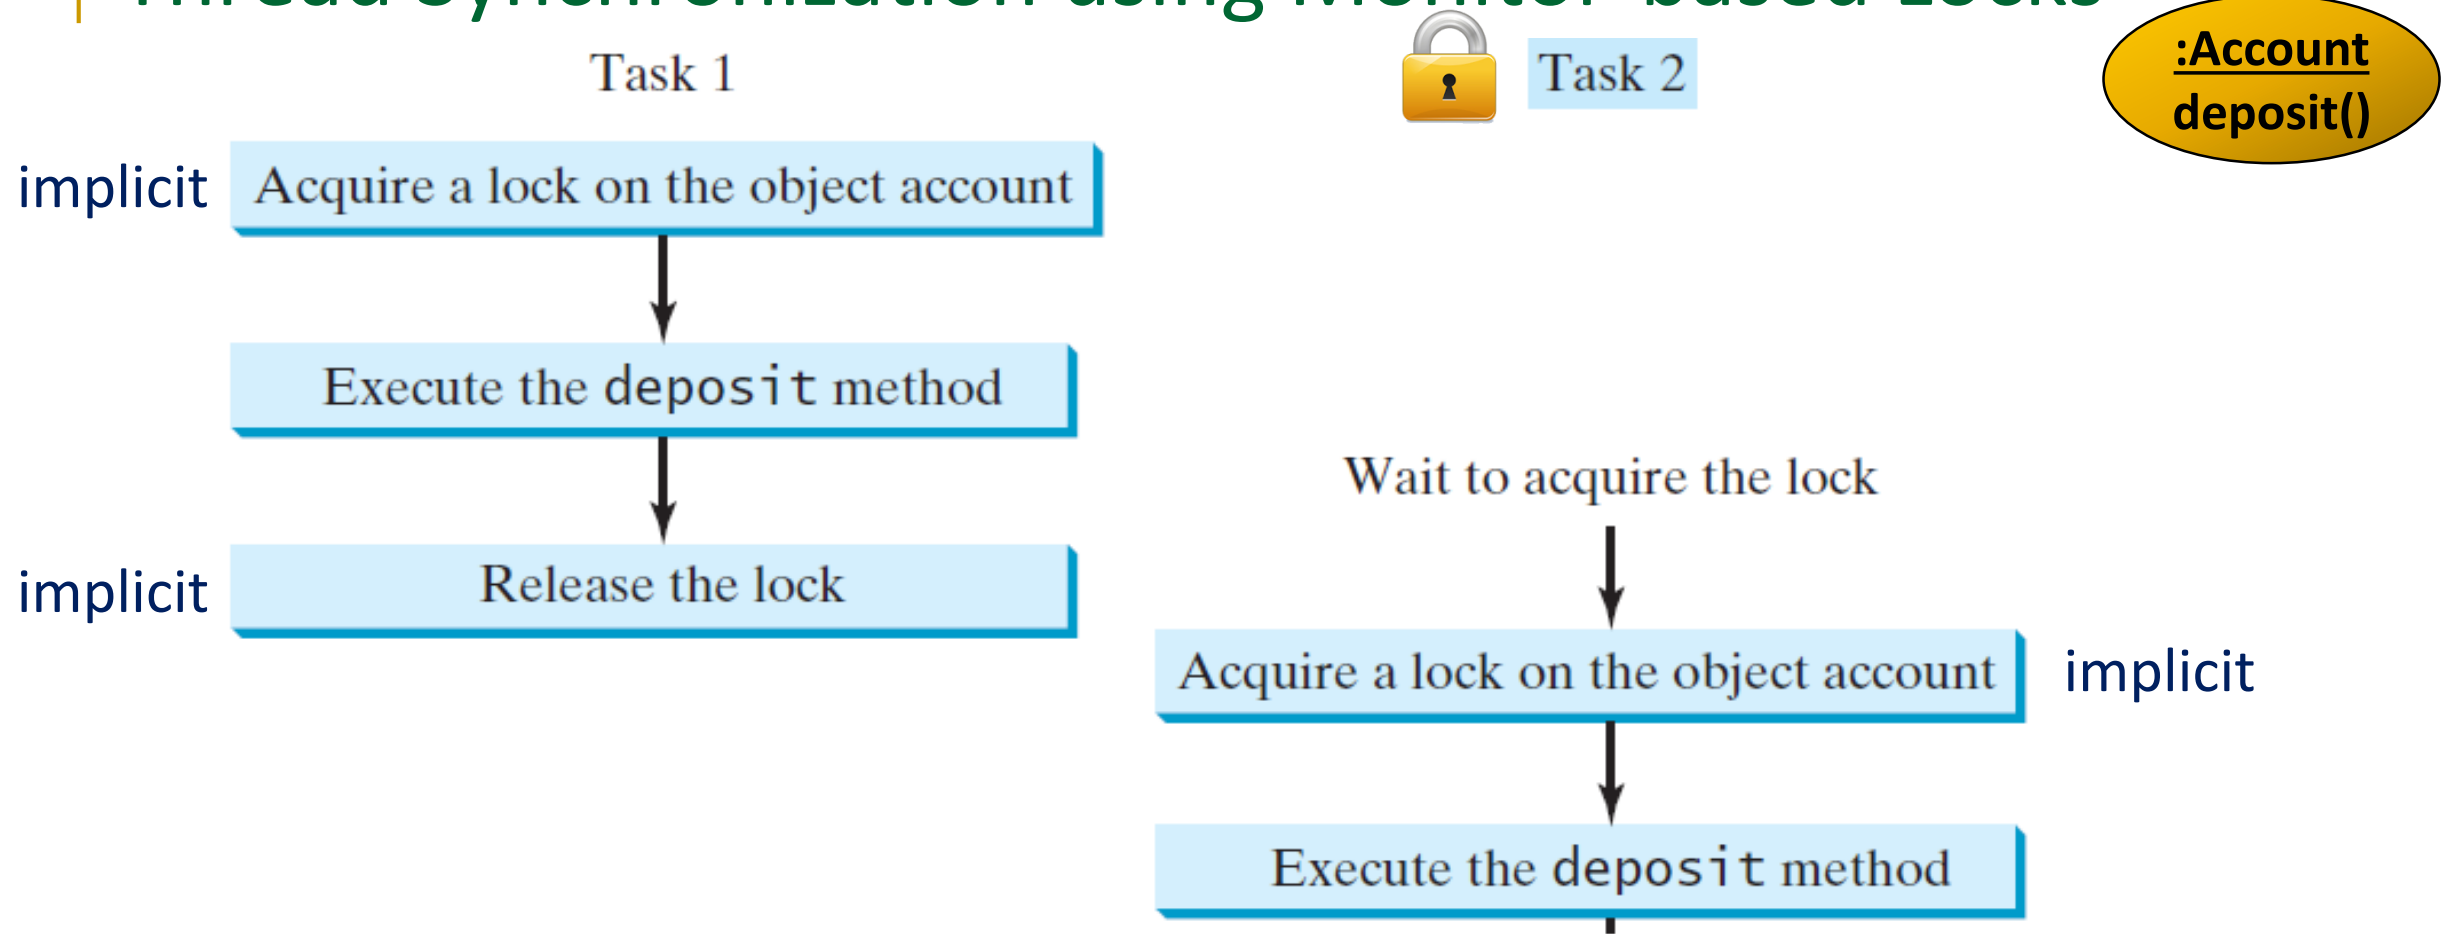
Release the lock

Wait to acquire the lock

Thread Synchronization using Monitor-based Locks

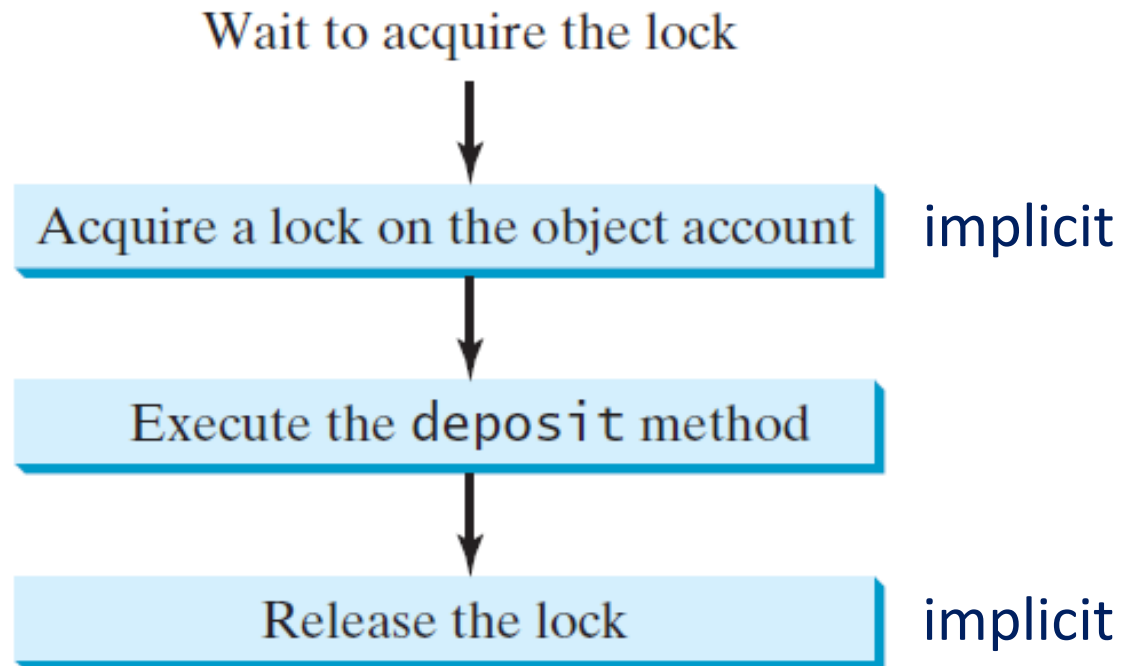
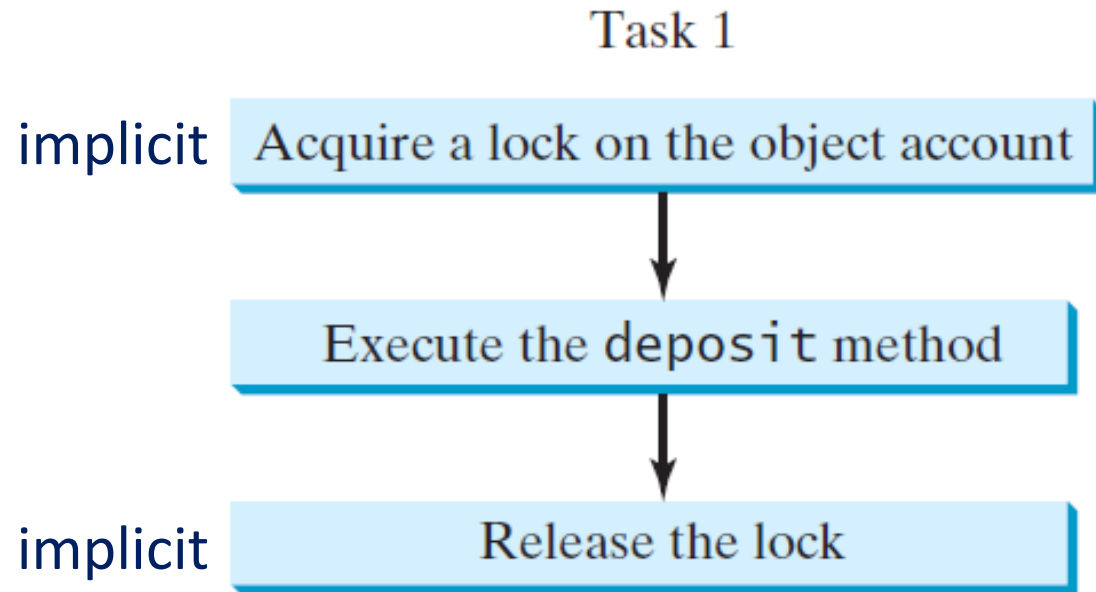


Thread Synchronization using Monitor-based Locks



Thread Synchronization using Monitor-based Locks



:Account
deposit()



Thread Synchronization using Synchronized Statements

```
synchronized (expr) { // expr must be evaluated to an object/class reference  
    statements;  
}
```

identifies a lock of
an object or a class



- A synchronized statement can be used to acquire a lock on a specific object, not limited to **this** object, when executing a block of the code in a method. This block is referred to as a **synchronized block**

- `synchronized (this) { ... }` `// synchronized using the current instance`
- `synchronized (new Account()) { ... }` `// synchronized using a new instance`
- `synchronized (Account.class) { ... }` `// synchronized using a class`

Thread Synchronization using Synchronized Statements

```
synchronized (expr) { // expr must be evaluated to object/class reference  
    statements;  
}
```

identifies a lock of
an object or a class




- The expression **expr** must be evaluated to an **object reference** or a **class reference**. If the object or class is already locked by another thread, the thread is blocked until the lock is successfully acquired
 - When the lock is acquired, the statements in the **synchronized block** are executed, and then the lock is released after these statements have been executed.
- [AccountWithSyncBlock.java](#)

Synchronized Statements vs. Synchronized Methods

- Any **synchronized instance method** can be converted into an instance method with a **synchronized block**

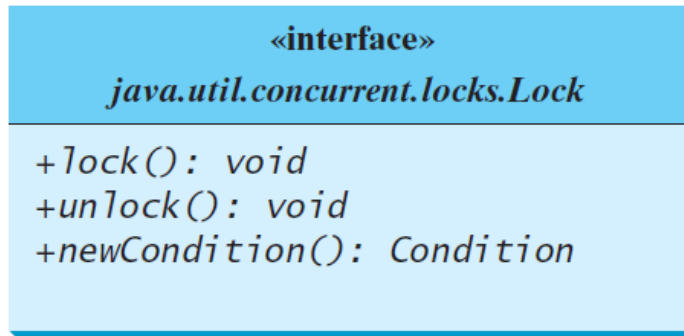



<pre>public synchronized void xMethod() { // method body }</pre>		<pre>public void xMethod() { synchronized (this) { // method body } }</pre>
--	---	---


Q: Why synchronized statements?

A: **Synchronized statements** enable us to synchronize threads at **statement granularity** instead of **method granularity**

Thread Synchronization Using Explicit Locks



Acquires the lock.
Releases the lock.
Returns a new `Condition` instance that is bound to this `Lock` instance.




Same as `ReentrantLock(false)`.
Creates a lock with the given fairness policy. When the fairness is true, the longest-waiting thread will get the lock. Otherwise, there is no particular access order.

Need to
call these
methods
explicitly

- Besides using monitor-based locks, we can synchronize threads by explicit locks
- An explicit lock is an instance of the **Lock** interface, which declares the methods to acquire and release locks. A lock may use the **newCondition()** method to **create** any number of **Condition objects**, which can be used for thread communications

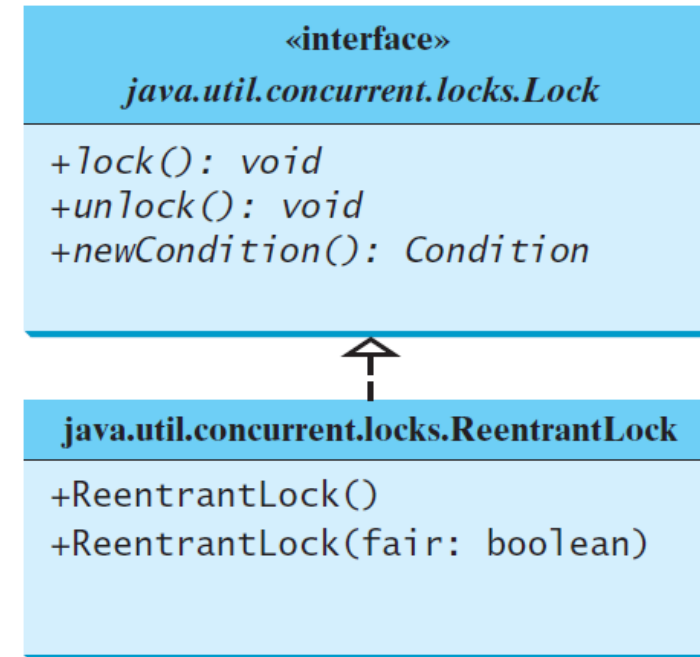
Thread Synchronization Using Explicit Locks



Example: A lock of a n-slot buffer

- c_1 : buffer is non-empty
 - consumer tasks queue for c_1
- c_2 : buffer is not full
 - supplier tasks queue for c_2

In contrast: implicit lock – one lock and one condition per object (more restrictive)



Need to call these methods explicitly

- Besides using monitor-based locks, we can synchronize threads by explicit locks
- An explicit lock is an instance of the **Lock** interface, which declares the methods to acquire and release locks. A lock may use the **newCondition()** method to **create** any number of **Condition objects**, which can be used for thread communications

Fairness Policy

- **ReentrantLock** is a concrete implementation of **Lock** for creating mutual exclusive locks.



- There is no effect to a thread when it acquires a reentrant lock being held
- We can create a lock with the specified **fairness policy**.
 - **True fairness** guarantees the longest-wait thread to obtain the lock first
 - **False fairness** grants a lock to a waiting thread without any access order
 - Programs using fair locks accessed by many threads may have poor overall performance than those using the default setting, but have smaller variances in times to obtain locks and guarantee free from starvation

Using Explicit Locks

```
public class Account {  
    private Lock lock = new ReentrantLock(true); // Create a fair lock
```

...

```
public void deposit(int amount) throws InterruptedException {
```

```
    lock.lock(); // Acquire the lock
```

```
    try {
```

```
        var newBalance = balance + amount;
```

```
        Thread.sleep( 5 );
```

```
        balance = newBalance;
```

```
    } catch (InterruptedException ex) {
```

```
        throws ex;
```

```
    }
```

```
    finally {
```

```
        lock.unlock(); // Release the lock
```



} critical region



Reentrant lock:
lock()
unlock()

[AccountWithSyncUsingLock.java](#)

Why do we need to call unlock in finally block?

Cooperation Among Threads



«interface»

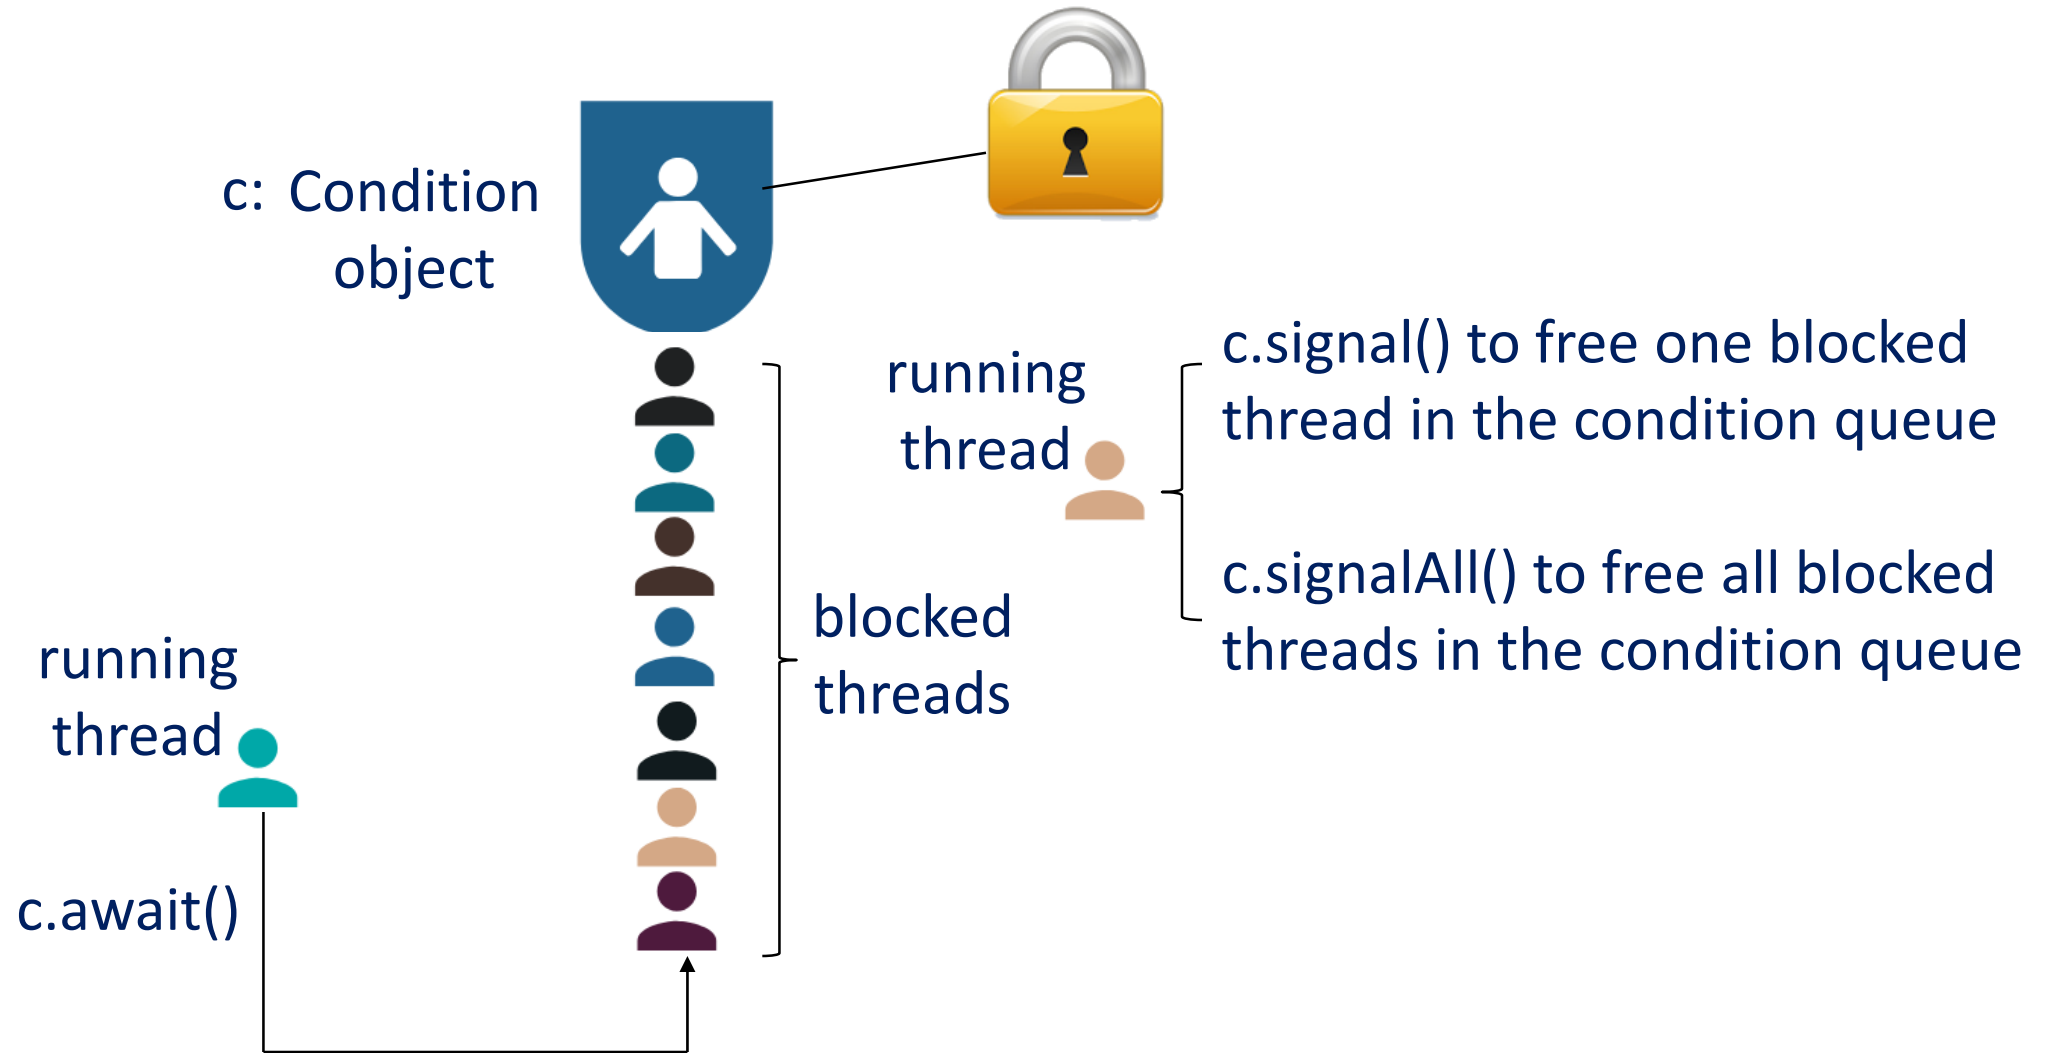
java.util.concurrent.Condition

```
+await(): void  
+signal(): void  
+signalAll(): Condition
```

Causes the current thread to wait until the condition is signaled.
Wakes up one waiting thread.
Wakes up all waiting threads.

- The conditions facilitate communications among threads. A thread can specify what to do under each condition
- Conditions are objects created by invoking the `newCondition()` method on a Lock object. Once a condition is created, we use its `await()`, `signal()`, and `signalAll()` methods for thread communications
- `await()` causes the current thread to wait until the condition is signaled; `signal()` wakes up one waiting thread, and `signalAll()` wakes up all waiting threads

Condition – await(), signal(), signalAll()



Example: Thread Cooperation



Suppose we launch two threads: one deposits to an account, and the other withdraws from the same account.

The withdraw thread awaits if the amount to be withdrawn is over the current account balance. Whenever the deposit thread adds new fund to the account, it signals the withdraw thread to resume. If the amount is still not enough for the withdrawal, the withdraw thread continues to await more fund in the account.

Assume the initial balance is 0 and the amount to deposit and to withdraw is randomly generated.

```
Command Prompt
C:\book>java ThreadCooperation
Thread 1          Thread 2          Balance
Deposit 7         Withdraw 9         7
Deposit 1         Withdraw 4         8
Deposit 10        Withdraw 3         18
                  Withdraw 5         9
                  Withdraw 4         5
Deposit 9         Withdraw 2         2
                  Withdraw 3         11
Deposit 3         Withdraw 5         6
                  Withdraw 2         4
                  Withdraw 2         7
```


Cooperation Among Threads

Withdraw Task

```
lock.lock();
```

```
while (balance < withdrawAmount)  
    newDeposit.await();
```

```
balance -= withdrawAmount
```

```
lock.unlock();
```

```
Condition newDeposit  
= lock.newCondition();
```

withdraw
thread

newDeposit.await()

blocked
threads



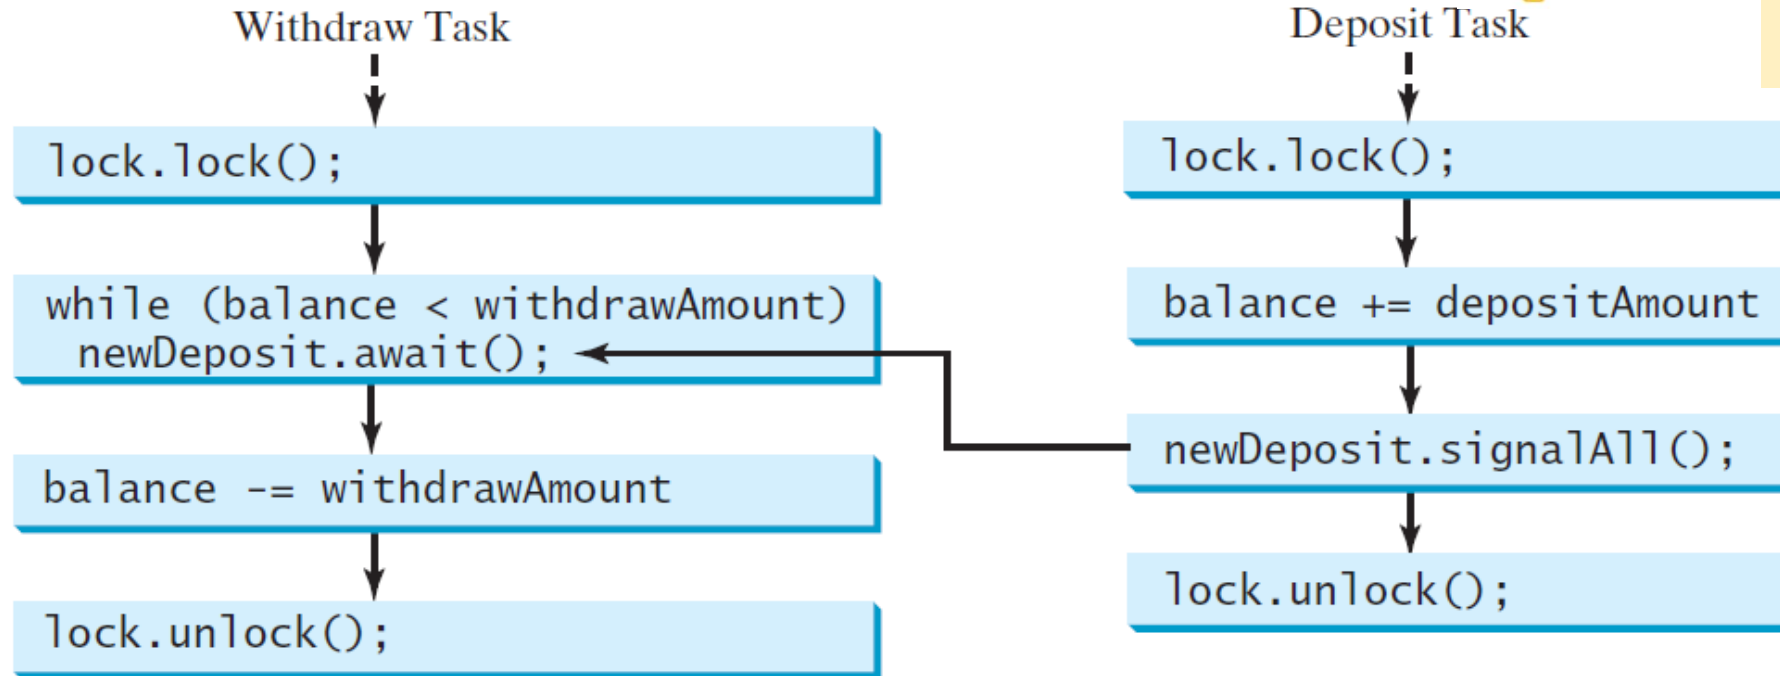
- Use a lock with a condition: newDeposit (to indicate new deposit added to the account)
- If the balance is less than the amount to be withdrawn, the withdraw task will wait for the newDeposit condition. When the deposit task adds money to the account, the task signals the waiting withdraw task to try again

ThreadCooperation.java

Cooperation Among Threads



Condition `newDeposit`
= `lock.newCondition();`



- Use a lock with a condition: `newDeposit` (to indicate new deposit added to the account)
- If the balance is less than the amount to be withdrawn, the withdraw task will wait for the `newDeposit` condition. When the deposit task adds money to the account, the task signals the waiting withdraw task to try again

`ThreadCooperation.java`

Java's Built-in Monitors (Implicit Locks)

- Prior to Java 5, thread communications are programmed using object's built-in monitors.
- Locks and conditions introduced since Java 5 are more powerful and flexible than the built-in monitor, which assumes one condition for each monitor lock
- However, we likely encounter the Java's built-in monitor when working with legacy code.
- A **monitor** is **an object with mutual exclusion and synchronization capabilities**. Only one thread can execute a method at a time in the monitor. A thread enters the monitor by acquiring a lock on the monitor and exits by releasing the lock. Any object can be a monitor. An object becomes a monitor once a thread locks it.
- Locking is implemented using the **synchronized** keyword on a method or a block. A thread must acquire a lock to execute a synchronized method or block. A thread can wait in a monitor if the condition is not right for it to continue executing in the monitor

wait(), notify(), and notifyAll()

- Use the `wait()`, `notify()`, and `notifyAll()` methods to facilitate communication among threads.
- The `wait()`, `notify()`, and `notifyAll()` methods must be called in a **synchronized method** or a **synchronized block** on the calling object of these methods. Otherwise, an **IllegalMonitorStateException** would occur
- The `wait()` method lets the thread wait until some condition occurs and immediately release the monitor lock of the object that owns the `wait()` method
- When the condition occurs, we use the `notify()` or `notifyAll()` methods to notify the waiting threads to resume normal execution. The `notifyAll()` method wakes up **all** waiting threads, while `notify()` picks up only **one** thread from a waiting queue
- When a waiting thread wakes up, it re-acquires the monitor lock that it has released

Example

Thread 1

```
public synchronized void deposit(int amt) {  
    balance += amt;  
    // Notify all threads waiting on the monitor  
    notifyAll();  
}
```


Thread 2


```
public synchronized void withdraw(int amt) {  
    try {  
        while (balance < amt)  
            wait();  
        balance -= amt;  
    } catch (InterruptedException ex) { ... }  
}
```

- When **wait()** is invoked, it **pauses the thread and simultaneously releases the lock** on the object. When the thread **wakes up after being notified**, the lock is **automatically re-acquired**
- The **wait()**, **notify()**, and **notifyAll()** methods on a monitor object are analogous to the **await()**, **signal()**, and **signalAll()** methods on a condition.


ThreadCooperationUsingBuiltInMonitor.java

Why replacing implicit locking by explicit locking?



- Implicit locking assumes **one condition for each lock** 
 - Each implicit lock supports one thread waiting queue (i.e., one condition)
- Explicit locking allows multiple conditions created for one lock
 - Each condition supports one thread waiting queue
- There are many applications where accesses to a critical section are subject to multiple conditions. Example: Producer/Consumer

Producer/Consumer: collaborate on a lock with two conditions



Condition 1:

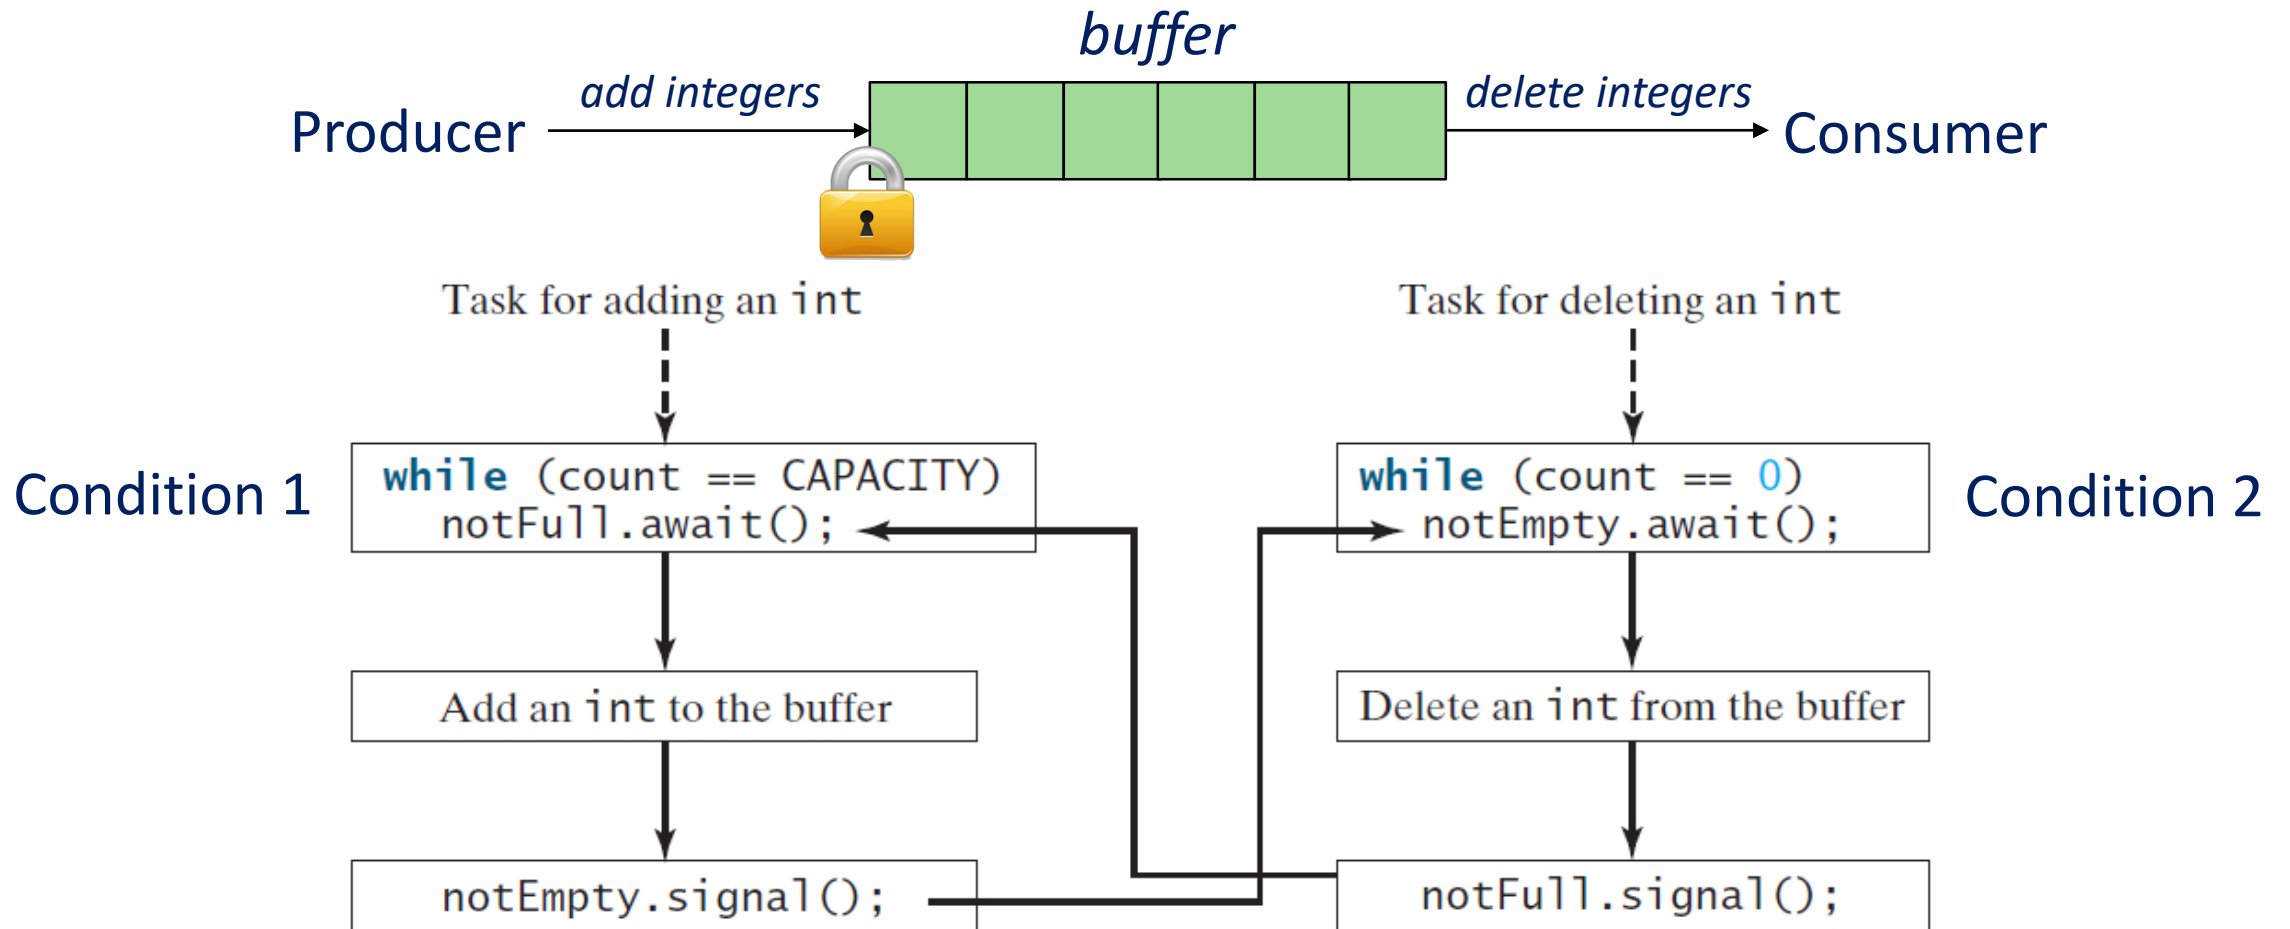
- Producer may add integers only **if the buffer is not full**
- Producers queue at this condition when the buffer is full

- One lock to protect buffer updates:
Only one thread can update the buffer at any time
- Buffer updates are governed by two conditions

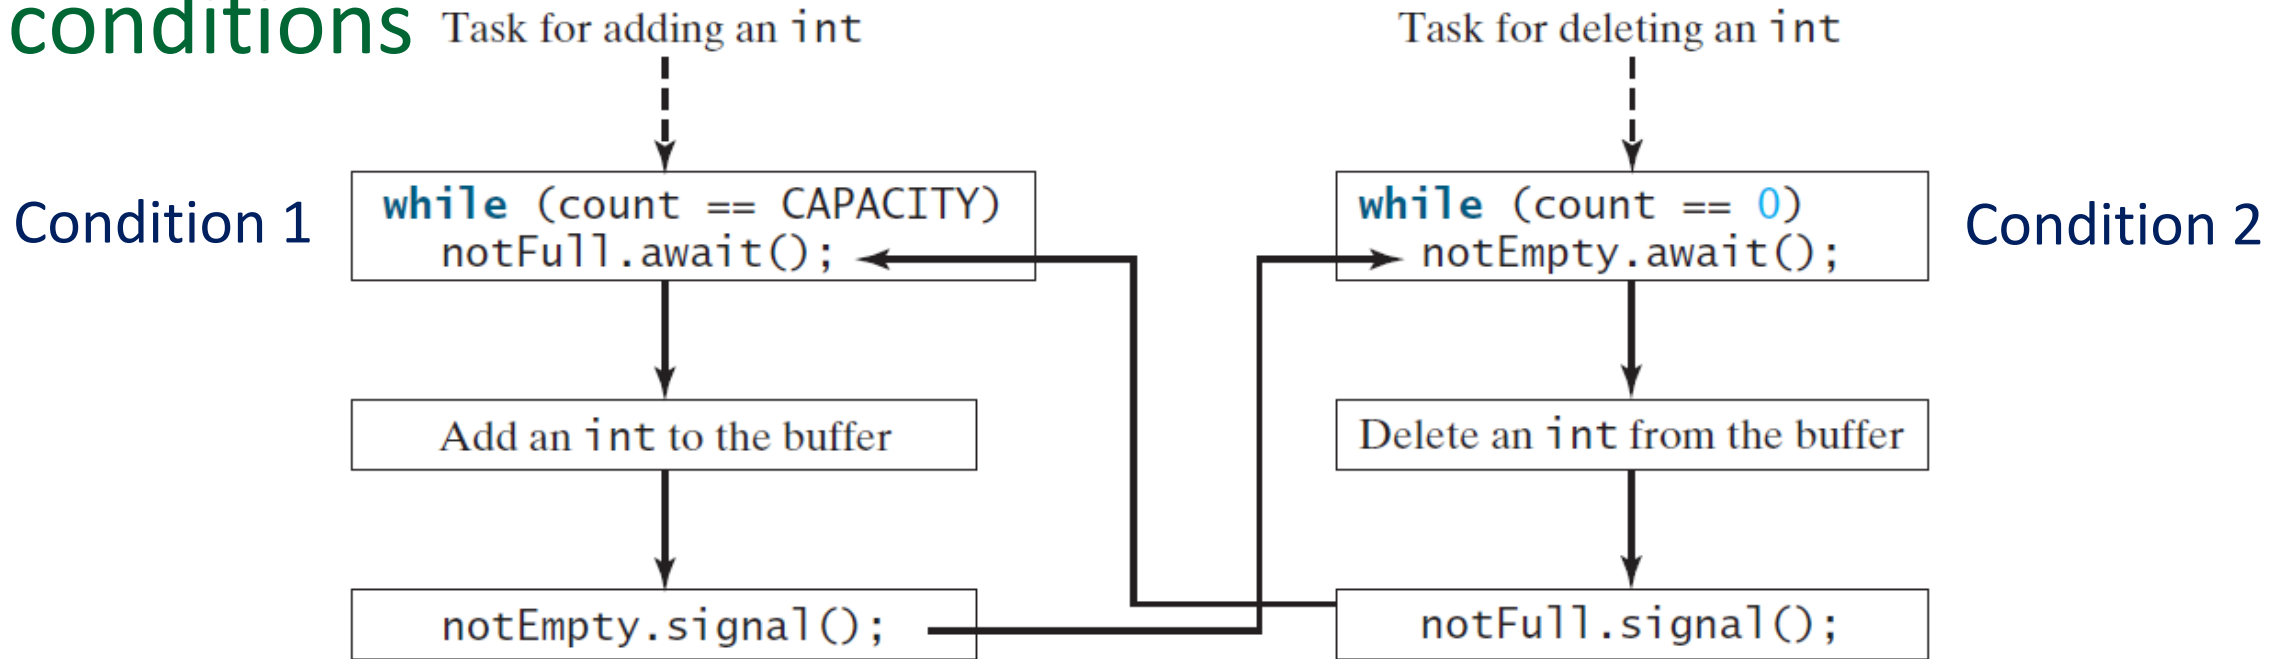
Condition 2:

- Consumer may delete integers only **if the buffer is not empty**
- Consumers queue at this condition when the buffer is empty

Producer/Consumer: collaborate on a lock with two conditions




Producer/Consumer: collaborate on a lock with two conditions



The buffer provides the method `write(int)` to add an `int` value to the buffer and the method `read()` to read and delete an `int` value from the buffer. To synchronize the operations, use a lock with two conditions: `notEmpty` (i.e., buffer is not empty) and `notFull` (i.e., buffer is not full). When a task adds an `int` to the buffer, if the buffer is full, the task will wait for the `notFull` condition. When a task deletes an `int` from the buffer, if the buffer is empty, the task will wait for the `notEmpty` condition.

[ConsumerProducerUsingLock.java](#)

Final Note



- It is sometimes possible to emulate one explicit locks with N conditions by N monitor-based locks
- The solution is usually less elegant and less maintainable
- This problem is more obvious in solutions that require more than one explicit lock
- Advice: Use explicit locking but able to understand monitor-based locking

