

COMP 3311 DATABASE MANAGEMENT SYSTEMS

FINAL REVIEW

INDEXING

A. A page can hold either 3 records or 10 (search-key, pointer) index entries. If a database contains n records, then how many pages are needed to store both the data file and a single-level dense index?

- a) $n/30$
- b) $3n/10$
- c) $10n/3$
- d) $13n/30$

number of records = n $bf_{\text{data file}} = 3$ $bf_{\text{index}} = 10$

Pages needed for the data file: $n/3$

Pages needed for the single-level dense index: $n/10$

Total pages needed: $n/3 + n/10 = \underline{13n/30}$

INDEXING (cont'd)

- B. In a B⁺-tree, if the search-key value is 12 bytes, the page size is 1024 bytes and a pointer is 6 bytes, then the maximum number of search-key values that can be stored in each non-leaf node of the B⁺-tree is:
- a) 54
 - b) 56
 - c) 57
 - d) 58

Each index entry is $12 + 6 = 18$ bytes.

Therefore, $\lfloor 1024 / 18 \rfloor = 56$ index entries (i.e., search-key values) can be stored in a non-leaf node.

INDEXING (cont'd)

C. What is the minimum number of search-keys in any non-root node for a B⁺-tree in which the maximum number of search-keys in a node is 4?

- a) 1
- b) 2
- c) 3
- d) 4

The fan out of the B⁺-tree, n , is 5 (i.e., one more than the number of values).

For internal nodes, the minimum number of values is $\lceil n/2 \rceil - 1 = \lceil 5/2 \rceil - 1 = 2$.

For leaf nodes, the minimum number of values is $\lceil (n-1)/2 \rceil = \lceil 4/2 \rceil = 2$.

Therefore, for any non-root node the minimum number of values is 2.

EXTERNAL SORTING

The relation *Sailor*(sailorId, name, rating, age) is not sorted. Assume each attribute is 25 bytes, the page size is 1,000 bytes and there are 11,000 tuples. For the following questions, apply external sorting using a buffer of 11 pages.

A. How many sorted runs will be produced in pass 0?

- a) 10
- b) 11
- c) 100
- d) 110

tuple size: 100 bytes
 bf_{Sailor} : 10 tuples/page
 B_{Sailor} : 1100 pages
buffer M : 11 pages

In pass zero, 11 pages at a time are sorted in memory.

Therefore, in pass zero $\lceil 1100/11 \rceil = 100$ sorted runs are produced.

EXTERNAL SORTING (cont'd)

B. What is the total number of passes required to sort the relation completely (including pass 0)?

a) 2

b) 3

c) 4

d) 5

tuple size: 100 bytes
 bf_{Sailor} : 10 tuples/page
 B_{Sailor} : 1100 pages
buffer M : 11 pages

After pass zero there are 100 sorted runs.

In pass 1, 10 runs at a time are merged, producing 10 sorted runs.

In pass 2, these 10 sorted runs are merged to produce the final output.

Therefore, the total number of passes = 3.

EXTERNAL SORTING (cont'd)

C. What is the total page I/O cost of sorting the relation?

- a) 4,400
- b) 5,500
- c) 6,600
- d) 7,770

tuple size: 100 bytes
 bf_{Sailor} : 10 tuples/page
 B_{Sailor} : 1100 pages
buffer M : 11 pages

There are a total of three passes.

In each pass the whole relation (1100 pages) is read and written.

Therefore, the total page I/O cost is $3 * 2 * 1100 = \underline{6,600}$

QUERY PROCESSING


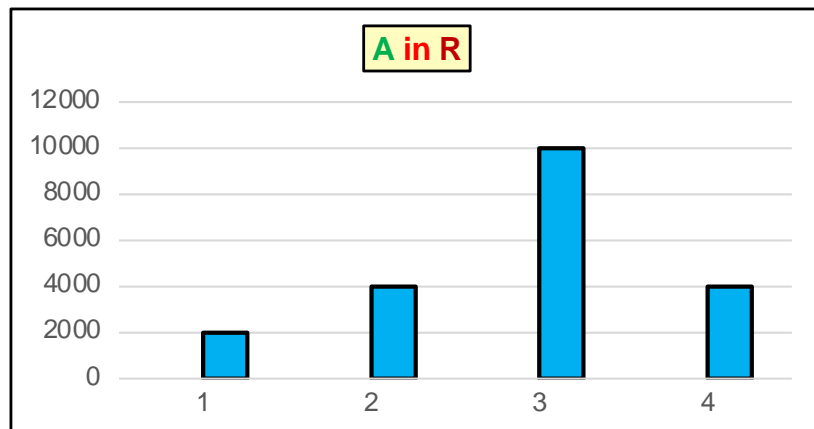
Consider the two relations $R(\underline{B}, \underline{C}, A)$ and $S(\underline{B}, A, Y)$.

R contains 20,000 tuples and S contains 100,000 tuples.

Assume that for both relations, 10 tuples fit per page (i.e., the size of R is 2,000 pages and that of S is 10,000 pages).

The possible values of attribute A in R are $\{1, 2, 3, 4\}$, whereas the possible values of A in S are $\{1, 2, 3, 4, 5\}$.

The following histograms present statistical information about the occurrences of values for A in R and S (e.g., there are 2,000 tuples with $A=1$ in R and 20,000 tuples with $A=1$ in S).





$R(\underline{B}, \underline{C}, A)$
 $S(\underline{B}, A, Y)$

QUERY PROCESSING (CONTD)

Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000

A. How many **tuples** are there in the query result of $(R \text{ JOIN}_{R.A=S.A} S)$ (i.e., what is the cardinality of the output)?

- a) 20,000
- b) 100,000
- c) 320,000,000
- d) 400,000,000



A is not a key for R or S .

According to the histograms, each tuple of R with value $R.A=1$ (of which there are 2,000) can join with 20,000 tuples of S with $S.A=1$. That is, the join result will contain $2,000 * 20,000 = 40 * 10^6$ tuples for $R.A=S.A=1$.

Performing the same computation for A equals 2, 3, 4 and 5 the result will contain $(40 + 100 + 100 + 80 + 0) * 10^6 = \underline{320,000,000}$ tuples.

$R(\underline{B}, C, A)$
 $S(\underline{B}, A, Y)$

QUERY PROCESSING (cont'd)

Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000

- B. What is the **minimum page I/O cost** to compute $(R \text{ JOIN}_{R.A=S.A} S)$ using block nested-loop join; **how many buffer pages** are needed?
- a) Minimum page I/O cost is 12,000; 2,000 buffer pages are needed.
 - b) Minimum page I/O cost is 12,000; 2,002 buffer pages are needed.
 - c) Minimum page I/O cost is 320,000 and 2,002 buffer pages are needed.
 - d) Minimum page I/O cost is 320,000; 12,000 buffer pages are needed.

Since there is no index, at least both relations must be read (i.e., the minimum page I/O cost is $2,000 + 10,000 = \underline{12,000}$ *independent of the algorithm*).

For block nested-loop join, the smaller relation R (2,000 pages) needs to be able to be kept in memory, so that S is scanned only once.

Since 1 page is needed for reading S and 1 page for holding the output, 2,002 buffer pages are needed.

$R(\underline{B}, C, A)$
 $S(\underline{B}, A, Y)$

QUERY PROCESSING (cont'd)

Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000


C. We want to compute $(R \text{ JOIN}_{R.A=S.A} S)$ using block nested-loop join with R as the outer relation. What is the minimum number of buffer pages needed in order to achieve a page I/O cost of 42,000?

- a) 502
- b) 2,000
- c) 2,002
- d) 10,002

Since R (2,000 pages) is the outer relation and the total page I/O cost is 42,000, S needs to be scanned $(42,000 - 2,000) / 10,000 = 4$ times.

If S is scanned 4 times, this means that 4 “blocks” of R need to be read to join with S , and each “block” must be at least $2,000 / 4 = 500$ pages.

Since, 1 page is needed for reading S and 1 page for the output, in total 502 buffer pages are needed.



$R(\underline{B}, \underline{C}, A)$
 $S(\underline{B}, A, Y)$

QUERY PROCESSING (CONTD)

Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000

D. We want to compute $(R \text{ JOIN}_{R.A=S.A} S)$ using hash join with R *as the build input*. How many buckets should be used for partitioning and what is the minimum buffer requirement?

- a) 4 buckets; 202 buffer pages.
- b) 4 buckets; 1,002 buffer pages.
- c) 5 buckets; 202 buffer pages.
- d) 5 buckets; 102 buffer pages.

The join condition is $R.A=S.A$ and there are only four values of A in R . Therefore, 4 buckets should be used.

According to the R histogram, 10,000 tuples of the build input A , have value $R.A=3$. Thus, the bucket size should be $\lceil 10,000 / 10 \rceil = 1,000$ pages ($bf=10$) so that each bucket of R can fit into the buffer.

In addition, we need 1 page for reading S and 1 for the output.

Therefore, 1,002 buffer pages are needed.

$R(\underline{B}, C, A)$
 $S(\underline{B}, A, Y)$

QUERY PROCESSING (cont'd)

Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000

E. Given that $R.B$ is a **not null** foreign key referencing $S.B$, **how many tuples** are in the query result of $(R \text{ JOIN}_{R.B=S.B} S)$?

- a) 20,000
- b) 100,000
- c) 120,000
- d) 320,000

Each tuple of R joins with exactly one tuple in S .

Therefore, the size of the join result is the same as the size of R , namely, 20,000 tuples.

$R(\underline{B}, C, A)$
 $S(\underline{B}, A, Y)$

QUERY PROCESSING (cont'd)


Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000

F. We want to compute $(R \text{ JOIN}_{R.B=S.B} S)$ using indexed nested-loop join with *R as the outer relation*. Assume that there is a hash index on $S.B$ with no overflow buckets (i.e., finding an index entry has page I/O cost 1). What is the *join page I/O cost*?

- a) 6,000
- b) 12,000
- c) 22,000
- d) 42,000

For each tuple of R , the hash index entry with the corresponding value of B is found and the pointer followed to the tuple of S for a page I/O cost of 2 per tuple. This is repeated for all 20,000 tuples of R , for a total page I/O cost of $20,000 * 2 = 40,000$.

In addition, the 2,000 pages of R must be read for a total page I/O cost of $40,000 + 2,000 = 42,000$.



$R(\underline{B}, \underline{C}, A)$
 $S(\underline{B}, A, Y)$

QUERY PROCESSING (CONTD)

Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000


G. How many tuples are there in the result of $((\sigma_{A=1} R) \text{ JOIN}_{R.B=S.B} S)$ and what is the minimum query processing page I/O cost using indexed nested-loop join with R as the outer relation. Assume that the only index is a hash index on S.B with no overflow buckets.

- a) The result has 2,000 tuples; the page I/O cost is 6,000.
- b) The result has 2,000 tuples; the page I/O cost is 22,000.
- c) The result has 20,000 tuples; the page I/O cost is 40,000.
- d) The result has 20,000 tuples; the page I/O cost is 42,000.

According to the R histogram, there are only 2,000 tuples in R with $R.A=1$. Each of these tuples, matches exactly 1 tuple in the join with S ($R.B=S.B$). Thus, the result contains 2,000 tuples.

Finding the matching S tuple has page I/O cost 2 (see previous question) and so finding all matches for the 2,000 tuples has page I/O cost 4,000.

Adding the cost of reading R (2,000 page I/Os), the minimum query processing page I/O cost is 6,000.



$R(\underline{B}, \underline{C}, A)$
 $S(\underline{B}, A, Y)$

QUERY PROCESSING (CONTD)

Relation	R	S
# tuples	20,000	100,000
# pages	2,000	10,000

H. How many tuples are there in the result of $((\sigma_{A=1} R) \text{ JOIN}_{R.B=S.B} (\sigma_{A=3} S))$ and what is the minimum query processing page I/O cost using index nested-loop join with *R as the outer relation*. Assume that the only index is a hash index on *S.B* with no overflow buckets.

- a) The result has 200 tuples; the page I/O cost is 600.
- b) The result has 200 tuples; the page I/O cost is 6,000.
- c) The result has 2,000 tuples; the page I/O cost is 6,000.
- d) The result has 2,000 tuples; the page I/O cost is 42,000.

According to the *S* histogram, among the 2,000 tuples in the result of $((\sigma_{A=1} R) \text{ JOIN}_{R.B=S.B} S)$, only $10,000 / 100,000 = 0.1$ (i.e., 10%) are expected to satisfy the additional condition *S.A=3*. Thus, the result is expected to contain $2,000 * 0.1 = 200$ tuples.


The minimum query processing page I/O cost is 6,000, the same as in the previous question.

TRANSACTION MANAGEMENT

4. Consider the following schedules of two transactions T_1 and T_2 . Indicate for each whether it is serial, (conflict) serializable or not serializable. r denotes a read and w a write operation.

a) Schedule: $r_1(A) w_1(A) r_2(A) w_2(B)$

Serial: $T_1 T_2$



T_1	T_2
$r_1(A)$ $w_1(A)$	$r_2(A)$ $w_2(B)$


A red arrow points from the $w_1(A)$ operation in the T_1 column to the $r_2(A)$ operation in the T_2 column, indicating a write-read conflict.

TRANSACTION MANAGEMENT (cont'd)

4. Consider the following schedules of two transactions T_1 and T_2 . Indicate for each whether it is serial, (conflict) serializable or not serializable. r denotes a read and w a write operation.

b) Schedule: $r_1(A)$ $r_2(A)$ $w_1(A)$ $w_2(B)$

Serializable: $T_2 T_1$




T_1	T_2
read(A)	
	read(A)
write(A)	
	write(B)

TRANSACTION MANAGEMENT (cont'd)

4. Consider the following schedules of two transactions T_1 and T_2 . Indicate for each whether it is serial, (conflict) serializable or not serializable. r denotes a read and w a write operation.

c) Schedule: $r_1(A)$ $r_2(A)$ $w_2(A)$ $w_1(A)$

Not Serializable



T_1	T_2
read(A)	read(A)
write(A)	write(A)

Red arrows indicate dependencies: one from T_2 's read(A) to T_1 's read(A), and another from T_1 's write(A) to T_2 's write(A).

TRANSACTION MANAGEMENT (cont'd)

5. Consider the schedule $r_1(A) w_1(A) r_2(A) w_2(B) c_1 c_2$ (where c_1 and c_2 indicate the commit statements).

a) Is the schedule recoverable? Why?

It is recoverable (T_2 reads database item A written by T_1 , but T_1 commits before T_2).

Is the schedule cascadeless? Why?

It is not cascadeless (T_2 read database item A before T_1 commits).

T_1	T_2
read(A)	
write(A)	
	read(A)
	write(B)
commit	
	commit

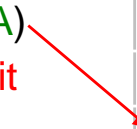
TRANSACTION MANAGEMENT (cont'd)

5. Consider the schedule $r_1(A) w_1(A) r_2(A) w_2(B) c_1 c_2$ (where c_1 and c_2 indicate the commit statements).
- b) Change the time of the commits (c_1, c_2) in the schedule in a) so that it becomes a cascadeless schedule.

$r_1(A) w_1(A) c_1 r_2(A) w_2(B) c_2$

Transaction T_2 reads the value of A written by T_1 , after T_1 commits.

T_1	T_2
read(A) write(A) commit	read(A) write(B) commit



TRANSACTION MANAGEMENT (cont'd)

5. Consider the schedule $r_2(A) \ r_1(A) \ w_1(A) \ w_2(B) \ c_2 \ c_1$ (where c_1 and c_2 indicate the commit statements).

(c) Is the schedule $r_2(A) \ r_1(A) \ w_1(A) \ w_2(B) \ c_2 \ c_1$ recoverable and cascadeless?

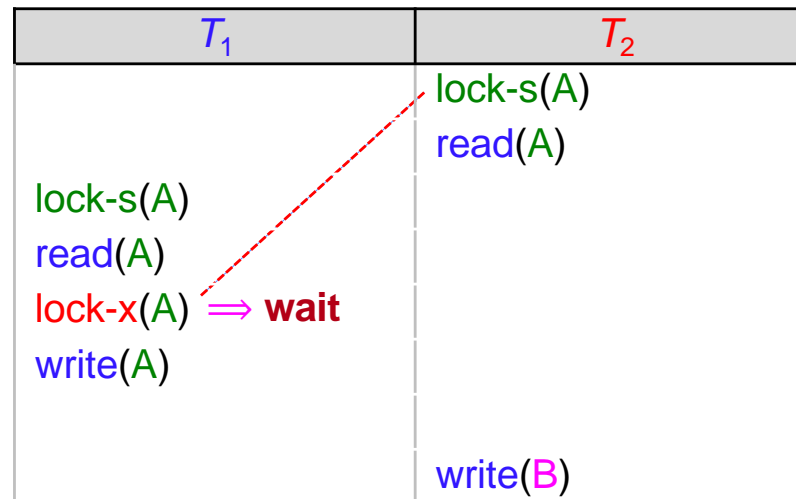
Both recoverable and cascadeless – no transaction reads items after they have been written by the other.

T_1	T_2
$read(A)$ $write(A)$ $commit$	$read(A)$ $write(B)$ $commit$

TRANSACTION MANAGEMENT (cont'd)

6. Modify the schedule $r_2(A)$ $r_1(A)$ $w_1(A)$ $w_2(B)$ according to the 2PL protocol by adding **lock-s**, **lock-x**, **unlock** statements. Explain briefly whether the schedule is allowed by 2PL.

The schedule is allowed by 2PL . T_1 must wait until T_2 unlocks A.



TRANSACTION MANAGEMENT (cont'd)

7. Modify the schedule $r_2(A) \ r_1(A) \ w_1(A) \ w_2(B)$ according to timestamp-ordering protocol by adding **RTS** (read timestamp) and **WTS** (write timestamp) statements. Assume that the timestamps of T_1 and T_2 are 2 and 1, respectively. The initial read and write timestamps of A and B are both 0.

T_1 [TS=2]	T_2 [TS=1]
$\text{read}(A) \text{ RTS}(A)=2; \text{WTS}(A)=0$ $\text{write}(A) \text{ RTS}(A)=2; \text{WTS}(A)=2$	$\text{read}(A) \text{ RTS}(A)=1; \text{WTS}(A)=0$ $\text{write}(B) \text{ RTS}(B)=0; \text{WTS}(B)=1$

Read

If $\text{TS}(T_i) < \text{WTS}(Q)$ **rollback**

If $\text{TS}(T_i) \geq \text{WTS}(Q)$

$\text{RTS}(Q) = \max(\text{TS}(T_i), \text{RTS}(Q))$

Write

If $\text{TS}(T_i) < \text{RTS}(Q)$ **rollback**

If $\text{TS}(T_i) < \text{WTS}(Q)$ **ignore**

Otherwise $\text{WTS}(Q) = \text{TS}(T_i)$

TRANSACTION MANAGEMENT (cont'd)

8. Modify the schedule $r_2(A) \ r_1(A) \ w_1(A) \ w_2(B)$ according to the multi-version timestamp-ordering protocol by adding **RTS** (read timestamp) and **WTS** (write timestamp) statements and specify the versions of the items. Assume that the timestamps of T_1 and T_2 are 1 and 2, respectively and that the initial versions of items are A_0 and B_0 with values 0 for their read and write timestamps. Complete the correct version numbers (e.g., $\text{read}(A_0)$ instead of $\text{read}(A)$).

T_1 [TS=1]	T_2 [TS=2]
$\text{read}(A_0)$ $\text{RTS}(A_0)=2$; $\text{WTS}(A_0)=0$	$\text{read}(A_0)$ $\text{RTS}(A_0)=2$; $\text{WTS}(A_0)=0$
$\text{write}(A)$	$\text{write}(B_1)$ $\text{RTS}(B_1)=2$; $\text{WTS}(B_1)=2$

$\text{TS}(T_1)=1 < \text{RTS}(A_0)=2 \Rightarrow \text{rollback}$

Read

Reads always succeed

set $\text{RTS}(Q_k) = \max(\text{TS}(T_i), \text{RTS}(Q_k))$

Write

If $\text{TS}(T_i) < \text{RTS}(Q_k)$ **rollback**

If $\text{TS}(T_i) = \text{WTS}(Q_k)$

overwrite contents

If $\text{TS}(T_i) > \text{WTS}(Q_k)$

create new version

set $\text{R/WTS}(Q') = \text{TS}(T_i)$

