



COMP 3311: Database Management Systems

Final Review

1. A. A page can hold either 3 records or 10 (key, search-pointer) index entries. If a database contains n records, then how many pages are needed to store both the data file and a single-level dense index?
 - a) $n/30$
 - b) $3n/10$
 - c) $10n/3$
 - d) $13n/30$
 - B. In a B⁺-tree, if the search-key value is 12 bytes, the page size is 1024 bytes and a pointer is 6 bytes, then the maximum number of search-key values that can be stored in each non-leaf node of the B⁺-tree is:
 - a) 54
 - b) 56
 - c) 57
 - d) 58
 - C. What is the minimum number of search-keys in any non-root node for a B⁺-tree in which the maximum number of search-keys in a node is 4?
 - a) 1
 - b) 2
 - c) 3
 - d) 4
-
2. The relation Sailor(sailorId, name, rating, age) is not sorted. Assume each attribute is 25 bytes, the page size is 1,000 bytes and there are 11,000 tuples. For the following questions, apply external sorting using a buffer of 11 pages.
 - A. How many sorted runs will be produced in pass 0?
 - a) 10
 - b) 11
 - c) 100
 - d) 110
 - B. What is the total number of passes required to sort the relation completely (including pass 0)?
 - a) 2
 - b) 3
 - c) 4
 - d) 5
 - C. What is the total page I/O cost of sorting the relation?
 - a) 4,400
 - b) 5,500
 - c) 6,600
 - d) 7,770

3. Consider the two relations $R(\underline{B}, \underline{C}, A)$ and $S(\underline{B}, A, Y)$. R contains 20,000 tuples and S contains 100,000 tuples. Assume that for both relations 10 tuples fit per page (i.e., the size of R is 2,000 pages and that of S is 10,000 pages). The possible values of attribute A in R are $\{1, 2, 3, 4\}$, whereas the possible values of A in S are $\{1, 2, 3, 4, 5\}$. The following histograms present statistical information about the occurrences of values for A in R and S (e.g., there are 2,000 tuples with $A=1$ in R and 20,000 tuples with $A=1$ in S).



- A. How many tuples are there in the query result of $(R \text{ JOIN}_{R.A=S.A} S)$ (i.e., what is the cardinality of the output)?
- 20,000
 - 100,000
 - 320,000,000
 - 400,000,000
- B. What is the minimum page I/O cost to compute $(R \text{ JOIN}_{R.A=S.A} S)$ using block nested-loop join; how many buffer pages are needed?
- Minimum page I/O cost is 12,000; 2,000 buffer pages are needed.
 - Minimum page I/O cost is 12,000; 2,002 buffer pages are needed.
 - Minimum page I/O cost is 320,000; 2,002 buffer pages are needed.
 - Minimum page I/O cost is 320,000; 12,000 buffer pages are needed.
- C. We want to compute $(R \text{ JOIN}_{R.A=S.A} S)$ using block nested-loop join with R as *the outer relation*. What is the minimum number of buffer pages needed in order to achieve a page I/O cost of 42,000?
- 502
 - 2,000
 - 2,002
 - 10,002
- D. We want to compute $(R \text{ JOIN}_{R.A=S.A} S)$ using hash join with R as *the build input*. How many buckets should be used for partitioning and what is the minimum buffer requirement?
- 4 buckets; 202 buffer pages.
 - 4 buckets; 1,002 buffer pages.
 - 5 buckets; 202 buffer pages.
 - 5 buckets; 102 buffer pages.
- E. Given that $R.B$ is a not null foreign key referencing $S.B$, how many tuples are in the query result of $(R \text{ JOIN}_{R.B=S.B} S)$?
- 20,000
 - 100,000
 - 120,000
 - 320,000

F. We want to compute $(R \text{ JOIN}_{R.B=S.B} S)$ using indexed nested-loop join with R as *the outer relation*. Assume that there is a hash index on $S.B$ with no overflow buckets (i.e., finding an index entry has page I/O cost 1). What is the join page I/O cost?

- a) 6,000
- b) 12,000
- c) 22,000
- d) 42,000

G. How many tuples are there in the result of $((\sigma_{A=1}R) \text{ JOIN}_{R.B=S.B} S)$ and what is the minimum query processing page I/O cost using indexed nested-loop join with R as *the outer relation*. Assume that the only index is a hash index on $S.B$ with no overflow buckets.

- a) The result has 2,000 tuples; the page I/O cost is 6,000.
- b) The result has 2,000 tuples; the page I/O cost is 22,000.
- c) The result has 20,000 tuples; the page I/O cost is 40,000.
- d) The result has 20,000 tuples; the page I/O cost is 42,000.

H. How many tuples are there in the result of $((\sigma_{A=1}R) \text{ JOIN}_{R.B=S.B} (\sigma_{A=3}S))$ and what is the minimum query processing page I/O cost using indexed nested-loop join with R as *the outer relation*. Assume that the only index is a hash index on $S.B$ with no overflow buckets.

- a) The result has 200 tuples; the page I/O cost is 600.
- b) The result has 200 tuples; the page I/O cost is 6,000.
- c) The result has 2,000 tuples; the page I/O cost is 6,000.
- d) The result has 2,000 tuples; the page I/O cost is 42,000.

4. Consider the following schedules of two transactions T_1 and T_2 . Indicate for each whether it is serial, (conflict) serializable or not serializable. r denotes a read and w a write operation.

a) Schedule: $r_1(A) w_1(A) r_2(A) w_2(B)$

T_1	T_2
read(A)	
write(A)	
	read(A)
	write(B)

b) Schedule: $r_1(A) r_2(A) w_1(A) w_2(B)$

T_1	T_2
read(A)	
	read(A)
write(A)	
	write(B)

c) Schedule: $r_1(A) r_2(A) w_2(A) w_1(A)$

T_1	T_2
read(A)	
	read(A)
	write(A)
write(A)	

5. Consider the schedule $r_1(A) w_1(A) r_2(A) w_2(B) c_1 c_2$ (where c_1 and c_2 indicate the commit statements).

- a) Is the schedule recoverable and cascadeless? Why?
- b) Change the time of the commits (c_1, c_2) in the schedule in a) so that it becomes a cascadeless schedule.
- c) Is the schedule $r_2(A) r_1(A) w_1(A) w_2(B) c_2 c_1$ recoverable and cascadeless?

6. Modify the schedule $r_2(A) r_1(A) w_1(A) w_2(B)$ according to the 2PL protocol by adding lock-s, lock-x, unlock statements. Explain briefly whether the schedule is allowed by 2PL.

T_1	T_2
read(A) write(A)	read(A) write(B)

7. Modify the schedule $r_2(A) r_1(A) w_1(A) w_2(B)$ according to timestamp-ordering protocol by adding RTS (read timestamp) and WTS (write timestamp) statements. Assume that the timestamps of T_1 and T_2 are 2 and 1, respectively. The initial read and write timestamps of A and B are both 0.

T_1 [TS=2]	T_2 [TS=1]
read(A) write(A)	read(A) write(B)

8. Modify the schedule $r_2(A) r_1(A) w_1(A) w_2(B)$ according to the multi-version timestamp-ordering protocol by adding RTS (read timestamp) and WTS (write timestamp) statements and specify the versions of the items. Assume that the timestamps of T_1 and T_2 are 1 and 2, respectively and that the initial versions of items are A_0 and B_0 with values 0 for their read and write timestamps. Complete the correct version numbers (e.g., read(A_0) instead of read(A)).

T_1 [TS=1]	T_2 [TS=2]
read(A) write(A)	read(A) write(B)