

COMP 3311: Database Management Systems

Lecture 14 Exercises Query Processing: Introduction

Exercise 1: For the relation Apartment(code, year, price, area), we want to process the query

Find the apartments in a given set of areas.

There are 1,000,000 apartments uniformly distributed over 100 areas. 10 Apartment records fit in one page. The records are ordered on price. There is a hash index on code and a dense, non-clustering B⁺-tree index on area. Each B⁺-tree node can fit 20 pointers (19 values). 100 record pointers can fit on a page. Assume that each non-leaf B⁺-tree node holds the minimum number of values.

Is it advantageous, in terms of page I/Os, to use the B⁺-tree on area to answer the query if there are 5 area values?

Cost for file scan:

Cost to use B⁺-tree index:

Exercise 2: How many passes are needed to sort a file that has 108 pages, using only 5 pages of main memory?

Number of passes

Pass 0:

Pass 1:

Pass 2:

Pass 3:

What is the page I/O cost?

Name: (1) _____ / _____ Student#: (1) _____ Date: _____
Last/Family (PRINT) Given/First (PRINT)

Name: (2) _____ / _____ Student#: (2) _____
Last/Family (PRINT) Given/First (PRINT)

NOTE: You are highly encouraged to do this exercise with a partner.

COMP 3311: Database Management Systems

Lecture 14 Exercises Query Processing: Introduction

Exercise 3: Assume that for the relation *Sailor*(*sailorId*, *sName*, *rating*, *age*) each record is 50 bytes, a page can hold 80 records, there are 500 pages of records and the file is unordered. We want to process the following two queries: i. $\sigma_{\text{sailorId} < 50,000}(\text{Sailor})$ ii. $\sigma_{\text{sailorId} = 50,000}(\text{Sailor})$

Estimate the page I/Os for each query given the information in a) and b) below.

a) Assume that we have a B⁺-tree index on the search key *sailorId* with height 5, 50 index pages, the lowest index search-key value is 1 and highest index search-key value is 100,000.

i. $\sigma_{\text{sailorId} < 50,000}(\text{Sailor})$

Page I/Os:

ii. $\sigma_{\text{sailorId} = 50,000}(\text{Sailor})$

Page I/Os:

b) Assume that we have a hash index on the search key *sailorId* with 50 index pages, no overflow, the lowest index search-key value is 1 and highest index search-key value is 100,000.

i. $\sigma_{\text{sailorId} < 50,000}(\text{Sailor})$

Page I/Os:

ii. $\sigma_{\text{sailorId} = 50,000}(\text{Sailor})$

Page I/Os:

Exercise 4: How many passes are needed to sort a file that has 600,000 pages, using 150 pages of main memory?

What is the page I/O cost?

Exercise 5: Let *B* be the size of a file in pages. How many memory pages *M* do you need to sort the file using one sorting pass and one merge pass (i.e., 2 passes total)? Show your calculations.