

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

LECTURE 15 EXERCISES

QUERY PROCESSING:

JOIN OPERATION

EXERCISE 1

Use the following information about the relations to estimate the page I/O cost to compute the query result using the stated join strategies.

Sailor(sailorId, sName, rating, age)

Reserves(sailorId, boatId, rDate)

Query: `select *`
`from Sailor natural join Reserves;`

- Page size: 1000 bytes; buffer size $M = 100$ pages.
- Each attribute (and pointer where applicable) is 20 bytes.
- Each Sailor tuple is 80 bytes, and each Reserves tuple is 60 bytes.
- Sailor tuples = 10,000; $bf_{Sailor} = \lfloor 1000 / 80 \rfloor = 12$ Sailor tuples per page.
- Reserves tuples = 40,000; $bf_{Reserves} = \lfloor 1000 / 60 \rfloor = 16$ Reserves tuples per page.
- Sailor requires $\lceil 10000 / 12 \rceil = 834$ pages and Reserves $\lceil 40000 / 16 \rceil = 2500$ pages.
- Since there are 10,000 Sailor tuples and 40,000 Reserves tuples, a sailor has, on average, 4 reservations (but it is possible that some sailor's have more, while some have none).

```
select *  
from Sailor natural join Reserves;
```

EXERCISE I (cont'd)

Sailor tuples: 10,000
 bf_{Sailor} : 12
Sailor pages: 834
Reserves tuples: 40,000
 $bf_{Reserves}$: 16
Reserves pages: 2500
 M pages: 100

a) i. block nested-loop join - using Sailor as outer relation

- Read 98 pages of the Sailor relation into the buffer at a time (there are $\lceil 834 / 98 \rceil = 9$ “blocks” of Sailor pages). (*Requires 98 buffer pages*)
- For each “block” of Sailor pages we scan the Reserves relation (page-by-page) to find matching tuples. (*Requires 1 buffer page*)
- One buffer page is allocated for the output.

Join page I/O cost: $9 * 2500 + 834 = 23,334$

ii. block nested-loop join - using Reserves as outer relation

- Read 98 pages of the Reserves relation into the buffer at a time (there are $\lceil 2500 / 98 \rceil = 26$ “blocks” of Reserves pages). (*Requires 98 buffer pages*)
- For each “block” of Reserves pages we scan the Sailor relation (page-by-page) to find matching tuples. (*Requires 1 buffer page*)
- One buffer page is allocated for the output.

Join page I/O cost: $26 * 834 + 2500 = 24,184$


```
select *  
from Sailor natural join Reserves;
```

EXERCISE 1 (cont'd)

Sailor tuples: 10,000
 bf_{Sailor} : 12
Sailor pages: 834
Reserves tuples: 40,000
 $bf_{Reserves}$: 16
Reserves pages: 2500
 M pages: 100

b) indexed nested-loop join with hash index
on Reserves.sailorId (assume no overflow).

- For each Sailor tuple, find the corresponding entry in the hash index on Reserves.sailorId.
- This takes 1 page I/O per Sailor tuple (since we assume no overflow).
- Since each Sailor tuple has on average 4 reservations, and, since the hash index is non-clustering (secondary), we expect each sailorId to have 4 matching tuples in the Reserves relation. Therefore, we need 4 page I/Os per Sailor tuple to retrieve the Reserves records.
- We also need 1 page I/O per sailorId to retrieve the indirection pointers.



Join page I/O cost: cost of reading Sailor + #tuples in Sailor * 6
= 834 + 10,000 * 6 = 60,834


```
select *
from Sailor natural join Reserves;
```

EXERCISE I (cont'd)

Sailor tuples: 10,000
 bf_{Sailor} : 12
 Sailor pages: 834
 Reserves tuples: 40,000
 $bf_{Reserves}$: 16
 Reserves pages: 2500
 M pages: 100

c) merge join


Sort Sailor on sailorId



Since there are 100 buffer pages, for each sorting run read in 100 Sailor pages and write 100 pages. There are $\lceil 834/100 \rceil = 9$ sorted runs created 8 of size 100 pages and 1 of size 34 pages. Totally $834 \times 2 = 1668$ pages are read and written.

PASS 1 (merge)

The Sailor relation can be merged in 1 pass since there are 9 sorted runs, which require only 9 pages.



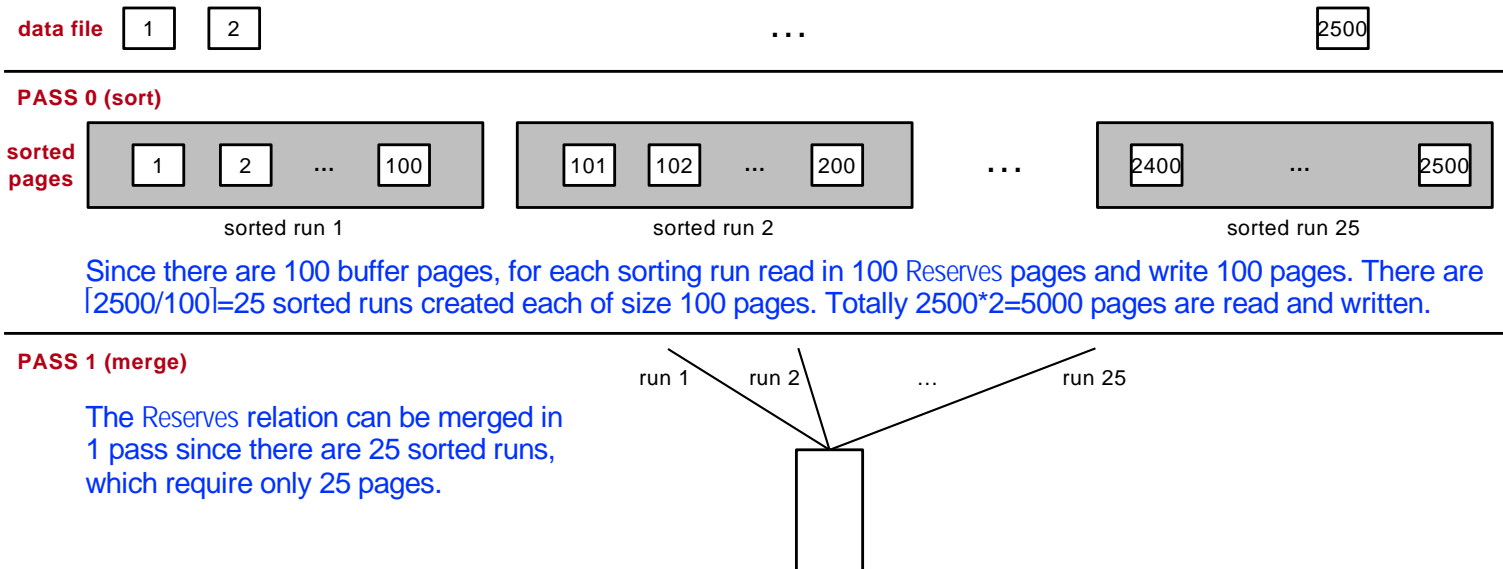
Sailor sorting page I/O cost: $\underset{r}{834} + \underset{w}{834}$ (Pass 0) + $\underset{r}{834} + \underset{w}{834}$ (Pass 1) = 3,336

```
select *
from Sailor natural join Reserves;
```

Sailor tuples: 10,000
 bf_{Sailor} : 12
 Sailor pages: 834
 Reserves tuples: 40,000
 $bf_{Reserves}$: 16
 Reserves pages: 2500
 M pages: 100

EXERCISE I (cont'd)

Sort Reserves on sailorId



At Pass 1 (merge) of the sort, as each sorted page of Reserves is generated, we can read the Sailor pages and directly find the joining tuples of Sailor. Consequently, we avoid writing the result of Pass 1 of the sort (i.e., 2500 pages) to a temporary file and reading it again for the merge-join phase.

Reserves sorting page I/O cost: $\underset{r}{2500} + \underset{w}{2500}$ (Pass 0) + $\underset{r}{2500}$ (Pass 1) = 7,500

Thus, for joining we only need to read and scan the 834 sorted Sailor pages.

Join page I/O cost: $3,336 + 7,500 + 834 = \underline{11,670}$

```
select *  
from Sailor natural join Reserves;
```

EXERCISE 1 (cont'd)

Sailor tuples: 10,000
 bf_{Sailor} : 12
Sailor pages: 834
Reserves tuples: 40,000
 $bf_{Reserves}$: 16
Reserves pages: 2500
 M pages: 100


d) hash join (assume no overflow)

- Use the smaller relation (Sailor = 834 pages) as the **build input**.
- We should choose the number of partitions n such that, when doing the **join phase**, all the pages of **each build partition** of Sailor **fits in the buffer**. For example, we can use 10 partitions so that the size of each partition is $\lceil 834/10 \rceil = 84$ pages, which fit into the buffer.
- First read and partition Sailor using 10 buffer pages, one for each partition.
Page I/O cost: $834 + 834 = 1,668$ page I/Os.
- Then read and partition Reserves using 10 buffer pages, one for each partition.
Page I/O cost: $2500 + 2500 = 5,000$ page I/Os.
 - Note that each partition of Reserves occupies 250 pages, which is more than the available buffer of 100 pages. However, we don't care since only the partitions of Sailor need to fit in the buffer.

```
select *
from Sailor natural join Reserves;
```

EXERCISE I (cont'd)

Sailor tuples: 10,000
 bf_{Sailor} : 12
 Sailor pages: 834
 Reserves tuples: 40,000
 $bf_{Reserves}$: 16
 Reserves pages: 2500
 M pages: 100



- Finally, we read, in turn, each partition (i.e., 84 pages) of Sailor, read the corresponding partition of Reserves page by page, match it (i.e., by hashing) against the Sailor partition and output the matching tuples.

Cost: $834 + 2500 = \underline{3,334}$ page I/Os.

Join page I/O cost: $3 * (2500 + 834) = \underline{10,002}$

EXERCISE 2

The relations $R_1(A, B, C)$ and $R_2(C, D, E)$ have the following properties:

- R_1 has 20,000 tuples
- R_2 has 45,000 tuples
- $bf_{R_1} = 25$ tuples
- $bf_{R_2} = 30$ tuples

R_1 requires 800 pages
 R_2 requires 1500 pages

Assuming that there are 800 buffer pages available for processing a join, estimate the page I/O cost for each of the following join strategies for $R_1 \text{ JOIN } R_2$.

- nested-loop join
- block nested-loop join
- merge join (assume that the relations are not sorted initially)
- hash join (assume no overflow occurs)

EXERCISE 2 (CONT'D)

a) nested-loop join: $n_r * B_s + B_r$

i. using R₁ as the outer relation

Join page I/O cost: $20000 * 1500 + 800 = \underline{30,000,800}$

ii. using R₂ as the outer relation

Join page I/O cost: $45000 * 800 + 1500 = \underline{36,001,500}$

b) block nested-loop join: $\lceil B_r / (M-2) \rceil * B_s + B_r$

i. using R₁ as the outer relation

Join page I/O cost: $\lceil 800/798 \rceil * 1500 + 800 = \underline{3,800}$

ii. using R₂ as the outer relation

Join page I/O cost: $\lceil 1500/798 \rceil * 800 + 1500 = \underline{3,100}$

EXERCISE 2 (CONT'D)

c) merge join (assume that both relations are not sorted initially)

Need to use external sorting.

👉 External sorting cost (sort & merge): $2 * B_r * (1 + \lceil \log_{M-1}(B_r/M) \rceil)$

Sorting cost

Page I/O cost to sort R₁: $2 * 800 * (1 + \lceil \log_{799}(800/800) \rceil) = 2 * 800 * (1 + 0) =$
1,600

Page I/O cost to sort R₂: $2 * 1500 * (1 + \lceil \log_{799}(1500/800) \rceil) =$
 $2 * 1500 * (1 + 1) =$ 6,000

Total page I/O cost to sort: 1600 + 6000 = 7,600

Merge cost (join phase)

Total page I/O cost to merge: 1500 + 800 = 2,300

Join page I/O cost: 7600 + 2300 = 9,900

EXERCISE 2 (CONT'D)

R_1 tuples = 20,000
 R_1 tuples/page = 25
 R_1 pages = 800
 R_2 tuples = 45,000
 R_2 tuples/page = 30
 R_2 pages = 1500
 M = 800 pages

d) hash join (assume no overflow occurs)

Use R_1 as the build input since it is smaller.

- Note that there is no need for recursive partitioning since the number of partitions is less than the number of buffer pages M (i.e.,
 $M > \sqrt{B_r} = \sqrt{800} = 28.3$)

Join page I/O cost: $3 * (1500 + 800) = \underline{6,900}$

