

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

LECTURE 17 EXERCISES

QUERY OPTIMIZATION

EXAMPLE RELATIONAL DATABASE

Sailor(sailorId, sName, rating, age)

Boat(boatId, bName, color)

Reserves(sailorId, boatId, rDate)

- There are 10,000 Sailor tuples, 100,000 Reserves tuples and 1,000 Boat tuples.
- For all files, $bf = 10$ tuples per page.
👉 $B_{\text{Sailor}} = 1,000$ pages; $B_{\text{Reserves}} = 10,000$ pages; $B_{\text{Boat}} = 100$ pages
- For a join result $bf = 5$ tuples per page.
- The buffer $M = 100$ pages.
- There are the following indexes:
 - hash index on sailorId for Sailor (no overflow buckets).
 - clustering B⁺-tree index on rDate for Reserves (2 levels).
 - hash index on boatId for Boat (no overflow buckets).
- Our goal is to process the query:

```
select *  
from Sailor natural join Reserves natural join Boat  
where rDate = '01-JAN-2020'  
and color = 'red';
```

Some useful statistics:

- Reserves has 1,000 unique rDates.
- 10% of boats are red.
- A sailor has on average 10 reservations.

EXERCISE 1

Estimate the minimum page I/O cost to process the query using materialization and the join order

$(\text{Sailor JOIN } \sigma_{\text{rDate}='01\text{-JAN-2020'}} \text{Reserves}) \text{ JOIN } \sigma_{\text{color}='red'} \text{Boat}$

C_1 : Cost of computing $\text{Temp}_1 = (\text{Sailor JOIN } \sigma_{\text{rDate}='01\text{-JAN-2020'}} \text{Reserves})$

C_2 : Cost of computing $\text{Temp}_2 = \sigma_{\text{color}='red'} \text{Boat}$ (no index on color)

C_3 : Cost of $\text{Temp}_1 \text{ JOIN } \text{Temp}_2$

Query processing page I/O cost: $C_1 + C_2 + C_3$

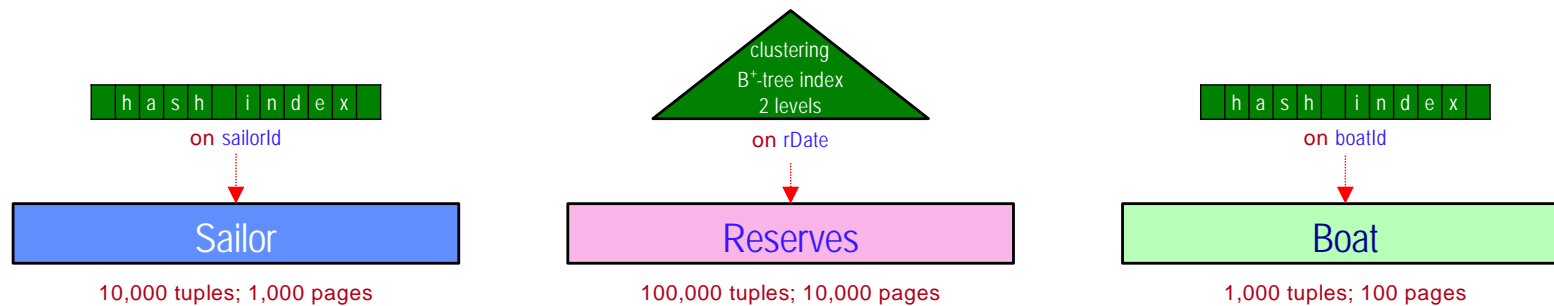
EXERCISE I (cont'd)

Sailor(sailorId, sName, rating, age)

Boat(boatId, bName, color)

Reserves(sailorId, boatId, rDate)

bf: 10 tuples/page
 B_{Sailor} : 1,000 pages
 B_{Reserves} : 10,000 pages
 B_{Boat} : 100 pages
 M : 100 pages
Hash index: Sailor.sailorId
B⁺-tree index: Reserves.rDate (2 levels)
Hash index: Boat.boatId
Unique Reserves rDates: 1000
Red boats: 10% (100 tuples)
Reservations/sailor: 10



```
select *  
from Sailor natural join Reserves natural join Boat  
where rDate = '01-JAN-2020'  
and color = 'red';
```

Some useful statistics:

- Reserves has 1,000 unique rDates.
- 10% of boats are red.
- A sailor has on average 10 reservations

```
select *
from Sailor natural join Reserves natural join Boat
where rDate = '01-JAN-2020'
and color = 'red';
```

EXERCISE I (CONT'D)

bf: 10 tuples/page
 B_{Sailor} : 1,000 pages
 B_{Reserves} : 10,000 pages
 B_{Boat} : 100 pages
 M : 100 pages
Hash index: Sailor.sailorId
B+-tree index: Reserves.rDate (2 levels)
Hash index: Boat.boatId
Unique Reserves rDates: 1000
Red boats: 10% (100 tuples)
Reservations/sailor: 10

C_1 : Cost of computing $\text{Temp}_1 = (\text{Sailor JOIN } \sigma_{\text{rDate}='01\text{-JAN-2020'}} \text{Reserves})$

In order to estimate C_1 , we need to determine the best evaluation plan (sub-plan) for computing Temp_1 .

Some alternative join strategies:

- block nested-loop – Sailor as outer relation: $\lceil B_r / (M-2) \rceil * B_s + B_r$.
- block nested-loop – Reserves as outer relation: $\lceil B_r / (M-2) \rceil * B_s + B_r$.
- sort-merge join: $B_r + B_s + \text{sorting cost: } 2 * B_r * (1 + \lceil \log_{M-1}(B_r/M) \rceil)$.
- hash join: $3 * (B_r + B_s)$.
- indexed nested-loop – Reserves as outer relation: $B_r + n_r * c$.

```

select *
from Sailor natural join Reserves natural join Boat
where rDate = '01-JAN-2020'
and color = 'red';

```

EXERCISE I (CONT'D)

bf: 10 tuples/page
 B_{Sailor} : 1,000 pages
 B_{Reserves} : 10,000 pages
 B_{Boat} : 100 pages
 M : 100 pages
 Hash index: Sailor.sailorId
 B+-tree index: Reserves.rDate
 (2 levels)
 Hash index: Boat.boatId
 Unique Reserves rDates: 1000
 Red boats: 10% (100 tuples)
 Reservations/sailor: 10

C_1 : Cost of computing $\text{Temp}_1 = (\text{Sailor JOIN } \sigma_{\text{rDate}='01\text{-JAN-2020'}} \text{Reserves})$

a) Strategy 1: block nested-loop – Sailor as outer relation.

$$\begin{aligned}
 \text{Join page I/O cost: } & \lceil B_r / (M-2) \rceil * B_s + B_r \\
 & = \lceil 1000 / (100-2) \rceil * 10000 + 1000 = \underline{111,000}
 \end{aligned}$$

b) Strategy 2: block nested-loop – Reserves as outer relation.

$$\begin{aligned}
 \text{Join page I/O cost: } & \lceil B_r / (M-2) \rceil * B_s + B_r \\
 & = \lceil 10000 / (100-2) \rceil * 1000 + 10000 = \underline{113,000}
 \end{aligned}$$

```
select *
from Sailor natural join Reserves natural join Boat
where rDate = '01-JAN-2020'
and color = 'red';
```

EXERCISE I (CONT'D)

bf: 10 tuples/page
 B_{Sailor} : 1,000 pages
 B_{Reserves} : 10,000 pages
 B_{Boat} : 100 pages
 M : 100 pages
 Hash index: Sailor.sailorId
 B+-tree index: Reserves.rDate
 (2 levels)
 Hash index: Boat.boatId
 Unique Reserves rDates: 1000
 Red boats: 10% (100 tuples)
 Reservations/sailor: 10

c) Strategy 3: sort-merge join: $B_r + B_s + \text{sorting cost}$.

$$\begin{aligned} \text{Sort Sailor page I/O cost: } & 2 * B_r * (1 + \lceil \log_{M-1}(B_r/M) \rceil) \\ & = 2 * 1000 * (1 + \lceil \log_{100-1}(1000/100) \rceil) = 4,000 \end{aligned}$$

$$\begin{aligned} \text{Sort Reserves page I/O cost: } & 2 * B_r * (1 + \lceil \log_{M-1}(B_r/M) \rceil) \\ & = 2 * 10000 * (1 + \lceil \log_{100-1}(10000/100) \rceil) = 60,000 \end{aligned}$$

$$\text{Merge page I/O cost: } B_r + B_s = 1000 + 10000 = 11,000$$

$$\text{Join page I/O cost: } 4000 + 60000 + 11000 = \underline{75,000}$$

d) Strategy 4: hash join: $3 * (B_r + B_s)$.

$$\text{Join page I/O cost: } 3 * (1000 + 10000) = \underline{33,000}$$

```


select *
from Sailor natural join Reserves natural join Boat
where rDate = '01-JAN-2020'
and color = 'red';

```

EXERCISE I (CONT'D)

bf: 10 tuples/page
 B_{Sailor} : 1,000 pages
 B_{Reserves} : 10,000 pages
 B_{Boat} : 100 pages
 M : 100 pages
 Hash index: Sailor.sailorId
 B⁺-tree index: Reserves.rDate (2 levels)
 Hash index: Boat.boatId
 Unique Reserves rDates: 1000
 Red boats: 10% (100 tuples)
 Reservations/sailor: 10

e) Strategy 5: indexed nested-loop – Reserves as outer relation.



Sailor contains a hash index on the join attribute sailorId. Furthermore, we have a selective condition ($\sigma_{\text{rDate}='01\text{-JAN-2020'}}(\text{Reserves})$) and a clustering index on Reserves.rDate.

How many reservations do we need to access?

Since there are 100,000 reservations and 1,000 unique rDates, on a given rDate there are $100000 / 1000 = 100$ reservations. Therefore, we need to access 100 Reserves tuples for 01-JAN-2020.


```
select *
from Sailor natural join Reserves natural join Boat
where rDate = '01-JAN-2020'
and color = 'red';
```

EXERCISE I (CONT'D)

bf: 10 tuples/page
*B*_{Sailor}: 1,000 pages
*B*_{Reserves}: 10,000 pages
*B*_{Boat}: 100 pages
M: 100 pages
 Hash index: Sailor.sailorId
 B+-tree index: Reserves.rDate (2 levels)
 Hash index: Boat.boatId
 Unique Reserves rDates: 1000
 Red boats: 10% (100 tuples)
 Reservations/sailor: 10

How many pages need to be accessed, to find these 100 reservations?

The index on Reserves.rDate has 2 levels and is ordered on rDate. Since, 10 Reserves tuples fit per page, to retrieve 100 reservations 2 index pages and 10 Reserves pages need to be accessed = 12 page I/Os. (Assumes all 100 Reserves tuples in the result are in 2 pages.)

For each of the 100 Reserves tuples, the corresponding Sailor tuple is retrieved using the hash index on Sailor.sailorId, with cost $2 * 100 = 200$ page I/Os. The join result will contain 100 tuples.

Join page I/O cost: $12 + 200 = 212$

Strategy 1	Strategy 2	Strategy 3	Strategy 4	Strategy 5
111,000	113,000	75,000	33,000	212

Since *bf*=5 for a join result, Temp₁ will occupy $\lceil 100 / 5 \rceil = 20$ pages.

Page I/O cost to write Temp₁: 20

Total page I/O cost for C₁: $212 + 20 = 232$

```
select *
from Sailor natural join Reserves natural join Boat
where rDate = '01-JAN-2020'
and color = 'red';
```

EXERCISE I (CONT'D)

bf: 10 tuples/page
 B_{Sailor} : 1,000 pages
 B_{Reserves} : 10,000 pages
 B_{Boat} : 100 pages
 M : 100 pages
 Hash index: Sailor.sailorId
 B+-tree index: Reserves.rDate (2 levels)
 Hash index: Boat.boatId
 Unique Reserves rDates: 1000
 Red boats: 10% (100 tuples)
 Reservations/sailor: 10

C_2 : Cost of computing $\text{Temp}_2 = \sigma_{\text{color}='red'} \text{Boat}$ (no index on color)

Strategy: file scan – linear search

For **Boat**, there is only a hash index on **boatId**. Therefore, to find red boats, a scan of the entire relation (100 pages) is needed.

File scan page I/O cost: 100

Since only 10% of the boats are red, we expect to retrieve $10\% * 1000 = 100$ Boat tuples, which can fit on $\lceil 100 / 10 \rceil = \underline{10}$ pages.

Page I/O cost to write Temp_2 : 10

Total page I/O cost for C_2 : $\underset{r}{100} + \underset{w}{10} = \underline{110}$ page I/Os

```

select *
from Sailor natural join Reserves natural join Boat
where rDate = '01-JAN-2020'
and color = 'red';

```

EXERCISE I (CONT'D)

bf: 10 tuples/page
 B_{Sailor} : 1,000 pages
 B_{Reserves} : 10,000 pages
 B_{Boat} : 100 pages
 M : 100 pages
 Hash index: Sailor.sailorId
 B+-tree index: Reserves.rDate (2 levels)
 Hash index: Boat.boatId
 Unique Reserves rDates: 1000
 Red boats: 10% (100 tuples)
 Reservations/sailor: 10

C₃: Cost of Temp₁ JOIN Temp₂ (Temp₁ is 20 pages; Temp₂ is 10 pages)

Strategy: block nested-loop join

Read the 20 pages of Temp₁.

Page I/O cost to read Temp₁: 20

Read the 10 pages of Temp₂.

Page I/O cost to read Temp₂: 10

Since there are 100 buffer pages, the join can be done in memory after reading both relations.

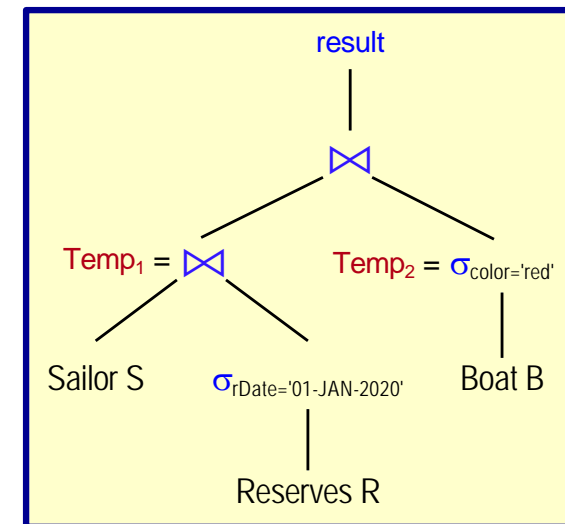
Join page I/O cost **C₃**: $20_r + 10_r = 30$

Query processing page I/O cost:

$$C_1 + C_2 + C_3 = 232 + 110 + 30 = 372$$

Note: As usual, we do not include the cost to write the final output.

Relational Algebra Tree



EXERCISE 2

Estimate the **minimum page I/O cost** to process the query **using materialization** and the join order

Sailor JOIN ($\sigma_{\text{rDate}='01\text{-JAN-2020'}$ Reserves JOIN $\sigma_{\text{color}='red'}$ Boat)

C_1 : Cost of computing $\text{Temp}_1 = \sigma_{\text{rDate}='01\text{-JAN-2020'}$ Reserves

C_2 : Cost of computing $\text{Temp}_2 = \text{Temp}_1$ JOIN $\sigma_{\text{color}='red'}$ Boat (no index on color)

C_3 : Cost of Sailor JOIN Temp_2

Query processing page I/O cost: $C_1 + C_2 + C_3$

```
select *  
from Sailor natural join Reserves natural join Boat  
where rDate = '01-JAN-2020'  
and color = 'red';
```

EXERCISE 2 (cont'd)

bf: 10 tuples/page
 B_{Sailor} : 1000
 B_{Reserves} : 10000
 B_{Boat} : 100
 M : 100 pages
Hash index: Sailor.sailorId
B+-tree index: Reserves.rDate
Hash index: Boat.boatId
Unique Reserves rDates: 1000
Red boats: 10% (100 tuples)
Reservations/sailor: 10

C_1 : Cost of computing $\text{Temp}_1 = \sigma_{\text{rDate}='01\text{-JAN-2020'}} \text{Reserves}$

Strategy 1: file scan – linear search

File scan page I/O cost: 10,000

Strategy 2: index lookup using B+-tree index on Reserves.rDate

As previously determined, to find all reservations on '01-JAN-2020' we retrieve 2 index pages and 100 Reserves tuples which occupy $\lceil 100 / 10 \rceil = 10$ pages in the Reserves relation (recall that Reserves is ordered on rDate).

Index lookup page I/O cost: $2 + 10 = 12$


The 100 result Reserves tuples occupy 10 pages.

Page I/O cost to write Temp_1 : 10

Total page I/O cost for C_1 : $12_r + 10_w = 22$

Question: Will all these Reserves tuples be for red boats? **No!**

How many are expected to be for red boats on average? 10%



EXERCISE 2 (cont'd)

bf: 10 tuples/page
 B_{Sailor} : 1000
 B_{Reserves} : 10000
 B_{Boat} : 100
 M : 100 pages
Hash index: Sailor.sailorId
B+-tree index: Reserves.rDate
Hash index: Boat.boatId
Unique Reserves rDates: 1000
Red boats: 10% (100 tuples)
Reservations/sailor: 10

C_2 : Cost of computing $\text{Temp}_2 = \text{Temp}_1 \text{ JOIN } \sigma_{\text{color}='red'} \text{Boat}$ (no index on color)

Strategy: block nested-loop join

Join Temp_1 with Boat and discard tuples where the color is not red. Temp_1 fits in memory so join cost is $B_{\text{Temp}_1} + B_{\text{Boat}}$.

Page I/O cost to read Temp_1 : 10

Page I/O cost to read Boat: 100

Join page I/O cost: $10 + 100 = 110$

Of the 100 Temp_1 tuples that join with a Boat tuple, $10\% * 100 = 10$ tuples are expected to join with a Boat tuple whose color is red since 10% of boats are red.

Since $bf=5$ for a join result, Temp_2 will occupy $\lceil 10 / 5 \rceil = 2$ pages.

Page I/O cost to write Temp_2 : 2

Total page I/O cost for C_2 : $110 + 2 = 112$

EXERCISE 2 (cont'd)

bf: 10 tuples/page
 B_{Sailor} : 1000
 B_{Reserves} : 10000
 B_{Boat} : 100
 M : 100 pages
 Hash index: Sailor.sailorId
 B+-tree index: Reserves.rDate
 Hash index: Boat.boatId
 Unique Reserves rDates: 1000
 Red boats: 10% (100 tuples)
 Reservations/sailor: 10

C_3 : Cost of Sailor JOIN Temp_2

Strategy: indexed nested-loop join using Sailor.sailorId hash index

For each of the 10 Temp_2 tuples, the hash index on sailorId is used to find the information about the sailor. The cost is 2 page I/Os (1 to the hash index and 1 to access the Sailor tuple) per reservation.

Page I/O cost to read Temp_2 : 2

Page I/O cost to use hash index on Sailor:
 $2 * 10 = 20$

Total page I/O cost for C_3 : $2 + 20 = 22$

Query processing page I/O cost:

$$C_1 + C_2 + C_3 = 22 + 112 + 22 = 156$$

Note: As usual, we do not count the cost to write the final output.

Relational Algebra Tree

