

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

LECTURE 19 EXERCISES


TRANSACTIONS

EXERCISE 1

Indicate which of the following schedules involving T_1 and T_2 is **serial**, **serializable** or **not serializable**. r_i denotes a read (of transaction T_i) and w_i is a write (of transaction T_i).

a) $r_1(A) w_1(A) r_2(A) w_2(A)$

Serial



T_1	T_2
$r_1(A)$ $w_1(A)$	$r_2(A)$ $w_2(A)$


Red arrows indicate dependencies: from $r_1(A)$ to $r_2(A)$ and from $w_1(A)$ to $w_2(A)$.

b) $r_1(A) r_2(A) w_1(A) w_2(B)$

Serializable

What is the equivalent serial schedule?

$T_2 T_1$




T_1	T_2
$r_1(A)$ $w_1(A)$	$r_2(A)$ $w_2(B)$

A red arrow indicates a dependency from $r_2(A)$ to $w_1(A)$.

EXERCISE I (cont'd)

c) $r_1(A) \ r_2(A) \ w_1(A) \ w_2(A)$


Not Serializable



T_1	T_2
read(A)	read(A)
write(A)	write(A)

d) $r_2(A) \ r_1(A) \ w_2(B) \ w_1(A)$

Serializable



T_1	T_2
read(A)	read(A)
write(A)	write(B)

What is the equivalent serial schedule?


$T_2 T_1$

EXERCISE 2

For each of the following schedules, state whether it is **serializable**, **recoverable** and **cascadeless**. Justify your answers. r_i denotes a read (of transaction T_i) and w_i a write (of transaction T_i).

a) $w_1(X) \ r_2(X) \ w_1(X) \ c_2 \ a_1$

T_1	T_2
write(X)	
	read(X)
write(X)	
	commit
abort	



Serializable? No

Justification: The precedence graph has a cycle.

Recoverable? No

Justification: T_2 reads data item X and commits, but then T_1 aborts.


Cascadeless? No

Justification: If T_1 aborts, T_2 must also abort.

EXERCISE 2 (cont'd)

b) $r_2(X) w_3(X) c_3 w_1(Y) c_1 r_2(Y) w_2(Z) c_2$

T_1	T_2	T_3
	read(X)	write(X) commit
write(Y) commit	read(Y) write(Z) commit	



Serializable? Yes

Justification: There is no cycle in the precedence graph.

The equivalent serial schedule is $T_1 T_2 T_3$.

Recoverable? Yes

Justification: T_2 reads a data item written by T_1 and commits after T_1 .


Cascadeless? Yes

Justification: The commit of T_1 appears before the commit of T_2 .

EXERCISE 2 (cont'd)

c) $r_1(X) \ w_2(X) \ c_2 \ w_1(X) \ c_1 \ r_3(X) \ c_3$

T_1	T_2	T_3
read(X)		
	write(X)	
	commit	
write(X)		
commit		
		read(X)
		commit



Serializable? No

Justification: The precedence graph has a cycle.

Recoverable? Yes

Justification: T_3 reads a data item written by T_1 and commits after T_1 .

Cascadeless? Yes

Justification: The commit of T_1 appears before the commit of T_3 .

Why recoverable and cascadeless?


The T_2 write is **useless** (called a **blind write**). It is not preceded by a read and it is immediately overwritten by T_1 . So, its value is never seen by any transaction.

EXERCISE 3

For each of the following schedules, answer the questions.

a) Is the schedule **serializable**?

T_1	T_2	T_3
read(X)	read(Y) write(Y)	
write(X)		write(Z)
	read(X) write(X)	
		read(Y) write(Y)
write(Z)		




The schedule is not serializable since there is a **cycle** T_1 , T_2 , T_3 , T_1 in the precedence graph.

There is no equivalent serial schedule because of the cycle.

EXERCISE 3 (cont'd)

b) Is the schedule **serializable**, **recoverable** and **cascadeless**?

T_1	T_2
read(A) write(A) commit	write(B) read(A) commit



Serializable? **Yes**

There is no cycle in the precedence graph.

The equivalent serial schedule is $T_1 T_2$.

Recoverable? **Yes**

T_2 reads A, written by T_1 and commits after T_1 .

Cascadeless? **No**

If T_1 aborts, T_2 must also abort.

How would you place the commit statements in order to make the schedule cascadeless?

T_1	T_2
read(A) write(A) commit	write(B) read(A) commit

EXERCISE 3 (cont'd)

c) Is the schedule **recoverable** and **cascadeless**?

T_1	T_2
read(A) write(A) commit	read(A) write(B) commit

Recoverable? **Yes**

No transaction reads data items after they have been written by the other.

Cascadeless? **Yes**

No transaction reads data items after they have been written by the other.

EXERCISE 4

Consider the following schedule consisting of three transactions T_1 , T_2 and T_3 . r_i denotes a read (of transaction T_i) and w_i is a write (of transaction T_i).

Schedule: $r_3(Z) \ w_3(Z) \ r_1(X) \ r_2(Y) \ w_2(Y) \ w_1(X) \ r_1(Y) \ r_3(X)$

- a) Show that the schedule is serializable by constructing the precedence graph.
- b) What is the equivalent serial schedule?
- c) Modify the original schedule so it becomes **recoverable**, but not **cascadeless** by adding commit operations to the end of the schedule.
- d) Modify the schedule so it becomes **both recoverable and cascadeless** by adding commit operations in the appropriate locations in the schedule.


EXERCISE 4 (cont'd)

Schedule: $r_3(Z)$, $w_3(Z)$, $r_1(X)$, $r_2(Y)$, $w_2(Y)$, $w_1(X)$, $r_1(Y)$, $r_3(X)$, c_1 , c_2 , c_3 ;

a) Show that the schedule is serializable by constructing the precedence graph.

T_1	T_2	T_3
<p>read(X)</p> <p>write(X)</p> <p>read(Y)</p> <p>commit</p>	<p>read(Y)</p> <p>write(Y)</p> <p>commit</p>	<p>read(Z)</p> <p>write(Z)</p> <p>read(X)</p> <p>commit</p>

Precedence Graph



b) What is the equivalent serial schedule?

$T_2 T_1 T_3$

EXERCISE 4 (cont'd)

- c) Modify the original schedule so it becomes **recoverable**, but not cascadeless by adding commit operations to the end the schedule.

Recall: A schedule is recoverable if the **commit** of a transaction T_j that reads data items *previously written* by a transaction T_i appears after the commit operation of T_i .

T_1 reads Y written by T_2
 $\Rightarrow T_1$ must commit after T_2 .
 T_2 does not read any data items written by other transactions.
 T_3 reads X written by T_1
 $\Rightarrow T_3$ must commit after T_1 .

Schedule:

$r_3(Z)$, $w_3(Z)$, $r_1(X)$, $r_2(Y)$, $w_2(Y)$,
 $w_1(X)$, $r_1(Y)$, $r_3(X)$, c_2 , c_1 , c_3

T_1	T_2	T_3
read(X)		read(Z) write(Z)
write(X) read(Y)	read(Y) write(Y)	
		read(X)
commit	commit	commit



EXERCISE 4 (cont'd)

- d) Modify the original schedule so it becomes **both recoverable and cascadeless** by adding commit operations in the appropriate locations in the schedule.

Recall: A schedule is cascadeless if, for each pair of transactions T_i , T_j such that T_j reads a data item previously written by T_i , the **commit** operation of T_i appears before the **read** operation of T_j .

- The commit of T_1 must appear before the **read**(X) of T_3 .
- The commit of T_2 must appear before the **read**(Y) of T_1 .
- The commit of T_3 must appear at the end of T_3 .

Schedule:

$r_3(Z)$, $w_3(Z)$, $r_1(X)$, $r_2(Y)$, $w_2(Y)$,
 c_2 , $w_1(X)$, $r_1(Y)$, c_1 , $r_3(X)$, c_3

T_1	T_2	T_3
		read(Z)
		write(Z)
read(X)		
	read(Y)	
	write(Y)	
	commit	
write(X)		
read(Y)		
commit		
		read(X)
		commit