

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

LECTURE 20 EXERCISES

CONCURRENCY CONTROL:

LOCK-BASED PROTOCOLS




EXERCISE 1

- a) Is the schedule **conflict serializable**?
If yes, give the equivalent serial schedule.

T_1	T_2	T_3
read(X)	read(Y)	
	write(Y)	
write(X)		write(Z)
	read(X)	
	write(X)	
		read(Y)
		write(Y)
write(Z)		

Precedence graph



The schedule is not conflict serializable because there is a cycle $T_1 T_2 T_3 T_1$ in the precedence graph.


Therefore, the schedule will fail under any protocol that aims at conflict serializability.

EXERCISE 1 (CONTD)

T_1	T_2	T_3
lock-s(X) ✓ read(X)	lock-s(Y) ✓ read(Y) lock-x(Y) ✓ write(Y)	
lock-x(X) ✓ write(X)		lock-x(Z) ✓ write(Z)
	lock-s(X) ✗ cannot be granted – T_2 must wait read(X) write(X)	
		lock-s(Y) ✗ cannot be granted – T_3 must wait read(Y) write(Y)
lock-x(Z) ✗ cannot be granted – T_1 must wait write(Z)		
deadlock!		

b) Modify the schedule according to strict 2PL by adding a **lock-s()** before reading a data item, a **lock-x()** before writing a data item and an **unlock()** after all read/write operations have completed. Is the schedule deadlock free?

Wait-for graph



Strict 2PL

All **exclusive-mode locks** must be **held** until a **transaction commits** (shared-mode locks can be released anytime).

Cannot request any lock after releasing a lock.




EXERCISE 2

Which of the following statements is true about the wait-for graph (circle the correct answer)?

- a) T_4 is waiting for T_3 to release a data item.
- b) The system is in a deadlock state after removing the edge between T_2 and T_4 .
- c) The system is in a deadlock state after removing the edge between T_3 and T_4 .
- d) The system is in a deadlock state when T_1 no longer holds a data item needed by T_4 .

Wait-for graph



- a) is not correct since T_3 is waiting for T_4 .
- b) and d) are not correct since after removing these edges, there does not exist a cycle in the wait-for graph.

EXERCISE 3

Show that the following schedule is conflict serializable according to 2PL by adding **lock-s()**, **lock-x()** and **unlock()** instructions, as necessary, to the schedule. If possible, add the **lock-s()**, **lock-x()** and **unlock()** instructions so that *no transaction is required to wait*.

T_1	T_2	T_3
read(X)		
	read(X)	
		read(Y)
read(Z)		
	read(Y)	
	write(X)	
		read(X)
		write(X)
write(Z)		

EXERCISE 3 (CONT'D)

Do you see any problem with this schedule?

Under 2PL, a transaction cannot request locks after it has released any lock.

T_1	T_2	T_3
lock-s(X) read(X) unlock(X)	lock-s(X) read(X)	lock-s(Y) read(Y)
lock-x(Z) read(Z)	lock-s(Y) read(Y) lock-x(X) write(X) unlock(Y) unlock(X)	lock-s(X) read(X) lock-x(X) write(X) unlock(Y) unlock(X)
write(Z) unlock(Z)		



EXERCISE 3 (cont'd)

T_1	T_2	T_3
lock-s(X) read(X)		
	lock-s(X) read(X)	
		lock-s(Y) read(Y)
lock-x(Z) read(Z)	lock-s(Y) read(Y)	
	lock-x(X) write(X) unlock(Y) unlock(X)	
		lock-s(X) read(X)
		lock-x(X) write(X) unlock(Y) unlock(X)
write(Z) unlock(Z) unlock(X)		

Do you see any problem with this schedule?

This lock cannot be granted; T_2 must wait.
How to add the locks so that T_2 does not have to wait?




T_1	T_2	T_3
lock-s(X) ✓ read(X)		
	lock-s(X) ✓ read(X)	
		lock-s(Y) ✓ read(Y)
lock-x(Z) ✓ unlock(X) ✓ read(Z)		
	lock-s(Y) ✓ read(Y)	
	lock-x(X) ✓ write(X)	
	unlock(Y) ✓ unlock(X) ✓	
		lock-s(X) ✓ read(X)
		lock-x(X) ✓ write(X)
		unlock(Y) ✓ unlock(X) ✓
write(Z) unlock(Z) ✓		

EXERCISE 3 (cont'd)

Recall: Under 2PL, a transaction cannot request locks after it has released any lock.

Precedence graph



The schedule is conflict serializable.
The equivalent serial schedule is $T_1 T_2 T_3$.



EXERCISE 4

In which positions, A to E, can an **unlock(X)** instruction be inserted if the schedule is according to:

a) **strict 2PL** (circle the correct answer)

⇒ all **x-locks** must be held until a transaction commits

⇒ **X** has only s-locks in the example

i. {A} {B} {C} {D}

ii. {A} {B} {C} {D} {E}

iii. {A} {C} {D}

iv. {B} {E}

v. {A} {C} {D} {E}

T_1	T_2
lock-s(X)	
read(X)	
	lock-s(X)
lock-x(Y)	
{A}	
read(Y)	
write(Y)	
	read(X)
	{C}
commit	
unlock(Y)	
{B}	
	{D}
	commit
	{E}



EXERCISE 4 (control)

In which positions, A to E, can an **unlock(X)** instruction be inserted if the schedule is according to:

b) **rigorous 2PL** (circle the correct answer)
 ⇒ all locks must be held until a transaction commits

- i. {A} {B} {C} {D}
- ii. {A} {B} {C} {D} {E}
- iii. {A} {C} {D}
- iv. {B} {E}
- v. {A} {C} {D} {E}

T_1	T_2
lock-s(X)	
read(X)	
	lock-s(X)
lock-x(Y)	
{A}	
read(Y)	
write(Y)	
	read(X)
	{C}
commit	
unlock(Y)	
{B}	
	{D}
	commit
	{E}

EXERCISE 5

- Is the schedule conflict serializable? If yes, give the equivalent serial schedule.
- If T_3 aborts after `write(Y)`, which other transactions will be rolled back?
- If T_1 aborts after `write(X)`, which other transactions will be rolled back?
- Construct the wait-for graph that results from this schedule if all locks are only exclusive-locks (`lock-x`), no locks are released, and the execution process runs to the point of `lock-x(Y)` in T_1 .
- Add `lock-s()`, `lock-x()` and `unlock()` instructions to the schedule according to strict 2PL. Indicate which transactions, if any, are required to wait and which instructions cause the wait.


T_1	T_2	T_3	T_4
read(X)			
write(X)			
	read(X)		
		read(Y)	
		write(Y)	
	write(X)		
			read(Y)
write(Y)			

EXERCISE 5 (control)

a) Is the following schedule conflict serializable? If yes, give the equivalent serial schedule.

T_1	T_2	T_3	T_4
read(X) write(X)			
	read(X)		
		read(Y) write(Y)	
	write(X)		
			read(Y)
write(Y)			

Precedence graph



The schedule is conflict serializable.
The equivalent serial schedule is
 $T_3 T_4 T_1 T_2$.

EXERCISE 5 (control)

- b) If T_3 aborts at the end of this schedule, which other transactions will be rolled back?

T_4 because after $\text{write}(Y)$ in T_3 , the Y read by T_4 is corrupted. Note that the $\text{write}(Y)$ of T_1 is not affected by T_3 as it is a blind write (i.e., no read before write) and in the serialization order the $\text{write}(Y)$ of T_1 would come after the $\text{write}(Y)$ of T_3 and overwrite it.

T_1	T_2	T_3	T_4
read(X)			
write(X)	read(X)		
		read(Y)	
	write(X)	write(Y)	
write(Y)			read(Y)

The equivalent serial schedule is $T_3 T_4 T_1 T_2$.

- c) If T_1 aborts at the end of this schedule, which other transactions will be rolled back?


T_2 because after $\text{write}(X)$ in T_1 , the X read by T_2 is corrupted.

EXERCISE 5 (control)

- d) Construct the wait-for graph that results from this schedule if all locks are only exclusive-locks (**lock-x**), no locks are released, and the execution process runs to the point of **lock-x(Y)** in T_1 .

T_1	T_2	T_3	T_4
lock-x(X) ✓ read(X) write(X)	lock-x(X) ✗ cannot be granted – T_2 must wait read(X)	lock-x(Y) ✓ read(Y) write(Y)	lock-x(Y) ✗ cannot be granted – T_3 must wait read(Y)
lock-x(Y) ✗ cannot be granted – T_1 must wait write(Y)	write(X)		

Wait-for Graph



EXERCISE 5 (control)

- e) Add **lock-s()**, **lock-x()** and **unlock()** instructions to the schedule according to strict 2PL. Indicate which transactions, if any, are required to wait and which instructions cause the wait.

Recall: Under strict 2PL, a transaction must hold all its ~~current~~ locks until it commits.

Normally, a transaction will request a **lock-s** before it reads a data item and upgrade to a **lock-x** later if it writes the data item.

