

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

LECTURE 3

ENTITY-RELATIONSHIP (E-R) MODEL

AND DATABASE DESIGN

E-R MODEL: CONSTRAINTS


A constraint is a logical restriction or property of data that for any set of data values:

- we can determine whether the constraint is **true** or **false**;
- we expect the constraint to be **always true**;
- we can **enforce** the constraint.

Examples:


on attributes

salary is between
\$0 and \$100,000




on relationships

every employee works in
at most one department



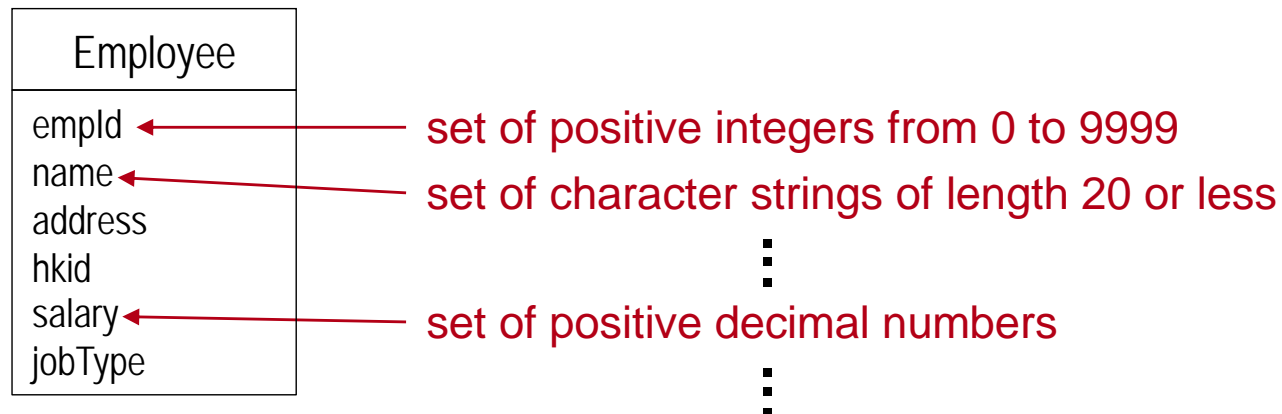
not every employee
is a manager



👉 **Constraints add additional semantics (meaning) to data**
(so as to more accurately reflect the application requirements).

ATTRIBUTE CONSTRAINTS: DOMAIN

A domain constraint restricts an attribute to only have certain values.



A domain constraint can be specified as a type for the attribute and/or a logical predicate that restricts the values.

ATTRIBUTE CONSTRAINTS: KEY

- If the values of some attributes **uniquely identify an entity instance**, then they are a **key** of the entity.
- A **candidate key** is a **minimal set of attributes** (i.e., all attributes are needed) that **uniquely identifies** an entity instance.

Employee
<u>empld</u>
name
address
hkid
salary
jobType

☞ An entity may have **more than one candidate key**.

- One candidate key is selected by the database designer to be the **primary key**.

☞ **This has enforcement implications for implementation.**

primary key \Rightarrow uniqueness is **automatically enforced** by a DBMS

other candidate keys \Rightarrow uniqueness is **not automatically enforced** by a DBMS

- A candidate/primary key can be composed of a set of attributes \Rightarrow **composite key**.

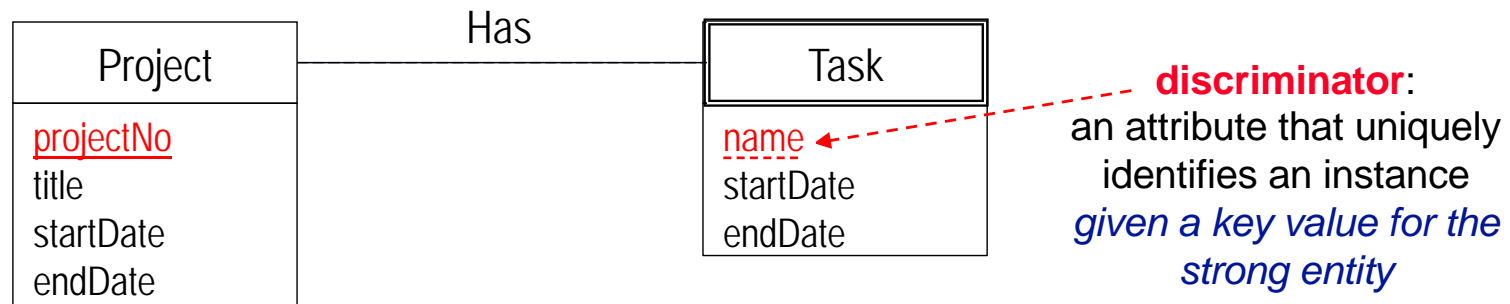
EnrollsIn
<u>studentId</u>
<u>courseCode</u>
grade



STRONG ENTITY VS. WEAK ENTITY: KEY

Strong entity: An entity that has a primary key.

Weak entity: An entity that does not have a primary key.



- A weak entity must be associated with a strong entity, called the **identifying entity**, to be meaningful.
 - ✎ A weak entity depends on its identifying entity for its existence.
- The relationship associating the weak entity to the strong entity is called the **identifying relationship** (shown as a solid line).
- A **discriminator**, *if present*, uniquely identifies a weak entity instance within its identifying relationship.


ENTITY GENERALIZATION CONSTRAINTS: COVERAGE

Disjointness

(a) overlapping

A superclass instance can relate to **more than one** subclass.


E.g., a given company can be both a customer and a supplier at the same time.



b) disjoint

A superclass instance can relate to **at most one** subclass.

E.g., a given employee can be either a secretary or an engineer, but not both at the same time.



ENTITY GENERALIZATION CONSTRAINTS:


COVERAGE (CONT'D)

Completeness

a) partial

A superclass instance **does not need to** relate to any of the subclasses.


E.g., a loan does not need to be a mortgage (loan) or a tax (loan) or a car (loan)—there are other kinds of loans.



(b) total

A superclass instance **must** relate to at least one of the subclasses.


E.g., a given customer must be either a personal or a business customer.




ENTITY GENERALIZATION CONSTRAINTS:

COVERAGE (CONT'D)


overlapping, partial




disjoint, partial



overlapping, total




disjoint, total



👉 Coverage is specified as one from disjointness (when there is more than one subclass) and one from completeness.

RELATIONSHIP CONSTRAINTS: CARDINALITY & PARTICIPATION

Cardinality specifies the maximum number and **participation** specifies the minimum number of relationship instances in which an entity may participate.




For a given project, how many employees can manage it?

☞ Each project is managed by one and only one employee.

For a given employee, how many projects can he/she manage?

☞ An employee does not have to manage any project, but may manage several (i.e., an unknown number of) projects.

RELATIONSHIP CONSTRAINTS: CARDINALITY & PARTICIPATION



minimum cardinality (min-card) \Rightarrow participation constraint

$\min\text{-card}(E_1, R)$: The *minimum* number of relationship instances in which *each entity* of E_1 *must* participate in R .

$\min\text{-card}(E_1, R) = 0 \Rightarrow$ partial participation

$\min\text{-card}(E_1, R) > 0 \Rightarrow$ total participation

maximum cardinality (max-card) \Rightarrow cardinality constraint

$\max\text{-card}(E_1, R)$: The *maximum* number of relationship instances in which *each entity* of E_1 *may* participate in R .

RELATIONSHIP CONSTRAINTS: SPECIAL MAXIMUM CARDINALITIES

$\text{max-card}(E_1, R) = 1$ and $\text{max-card}(E_2, R) = 1$

👉 **one-to-one relationship (1:1)**

$\text{max-card}(E_1, R) = 1$ and $\text{max-card}(E_2, R) = N$

👉 **one-to-many relationship (1:N)**

$\text{max-card}(E_1, R) = N$ and $\text{max-card}(E_2, R) = 1$

👉 **many-to-one relationship (N:1)**

$\text{max-card}(E_1, R) = N$ and $\text{max-card}(E_2, R) = N$


👉 **many to many relationship (N:M)**

RELATIONSHIP CONSTRAINTS: CARDINALITY & PARTICIPATION



An employee does not have to manage a project but can manage several projects.

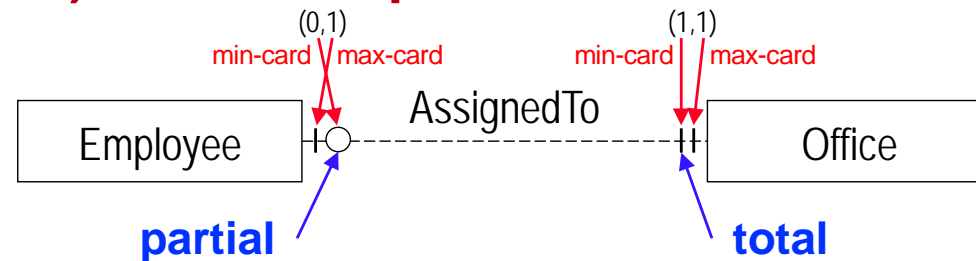
Every project must be managed by an employee and at most one employee.




RELATIONSHIP CONSTRAINTS:

EXAMPLE CARDINALITY & PARTICIPATION

one-to-one (1:1) relationship



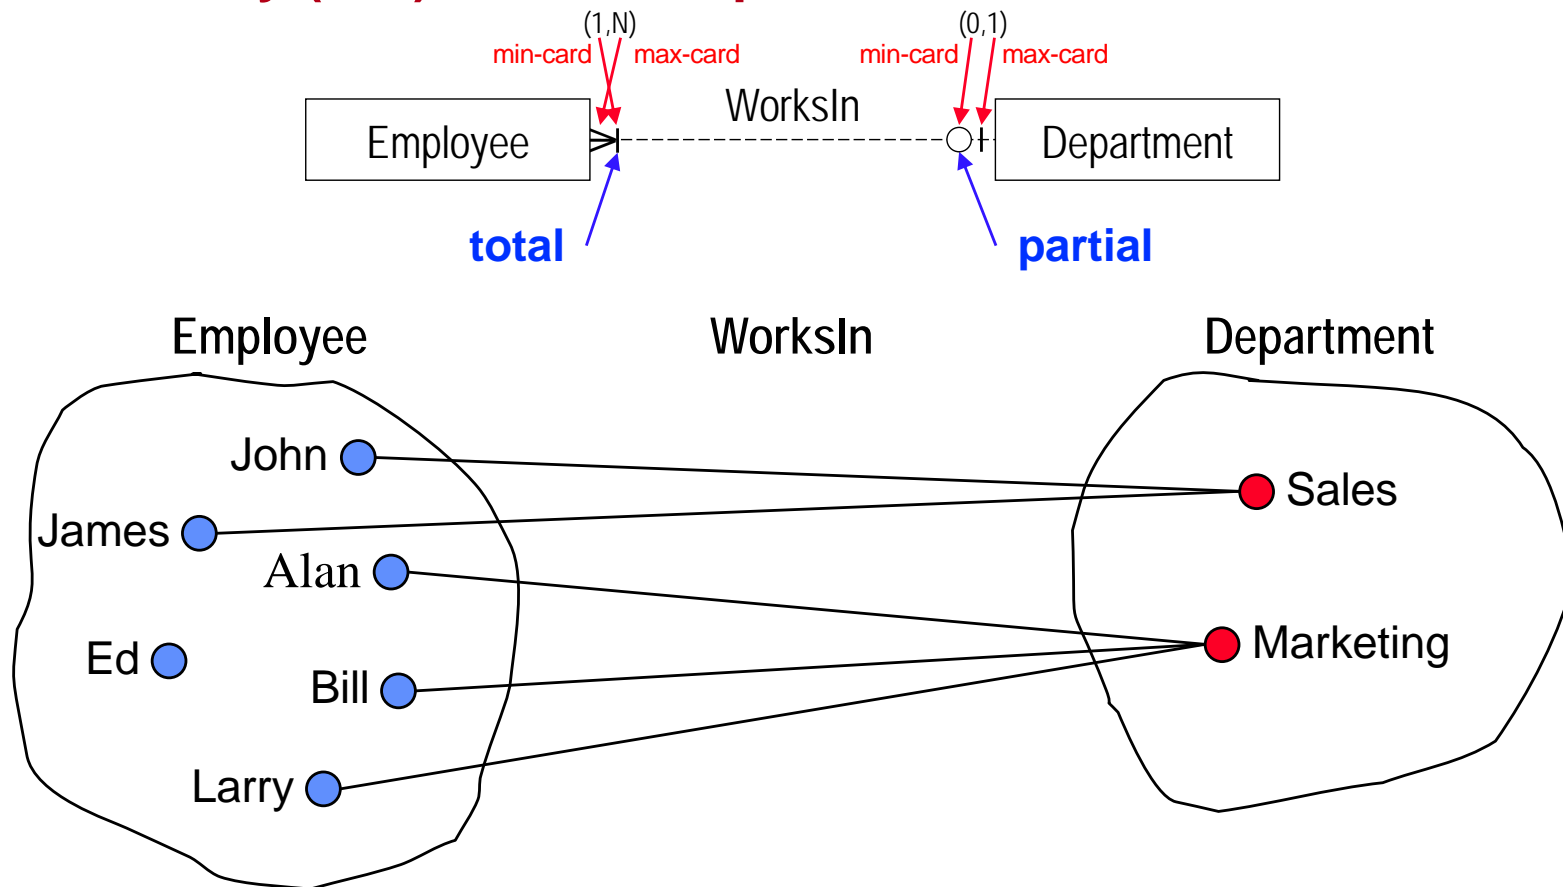
information
engineering
(crow foot)
notation



RELATIONSHIP CONSTRAINTS:

EXAMPLE CARDINALITY & PARTICIPATION

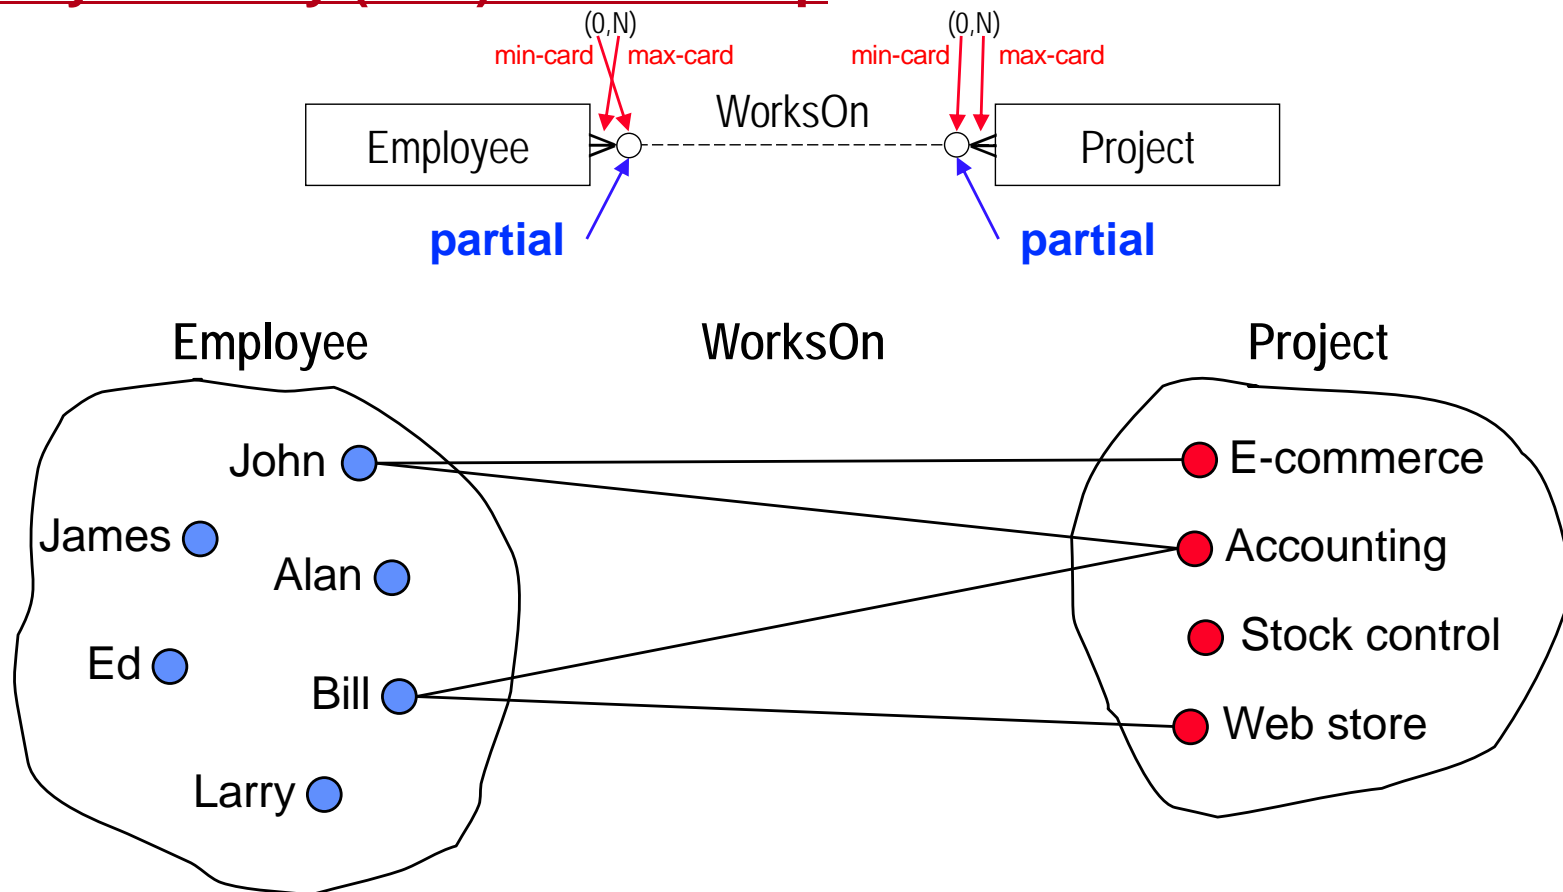
one-to-many (1:N) relationship



RELATIONSHIP CONSTRAINTS:

EXAMPLE CARDINALITY & PARTICIPATION

many-to-many (N:M) relationship

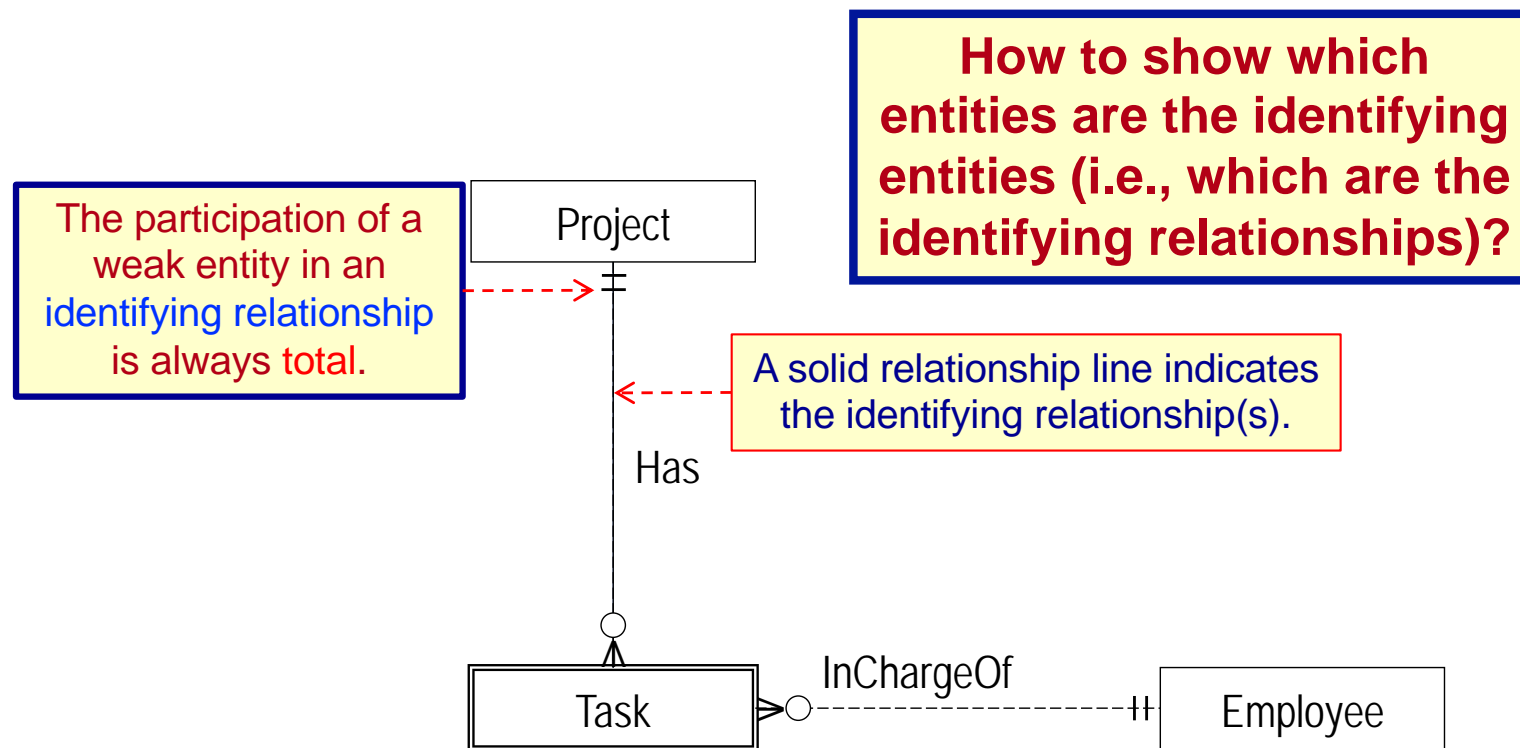


RELATIONSHIP CONSTRAINTS:

WEAK ENTITY PARTICIPATION

- A weak entity may be related to more than one strong entity but may depend on only some of these for its **existence**.


👉 Only some of the strong entities are identifying entities.



RELATIONSHIP CONSTRAINTS: **EXCLUSION**

An *exclusion (XOR) constraint* specifies that **at most one entity instance, among several entity types, can participate in a relationship with a single “root” entity.**


Example: A task can be related to *either* an internal project *or* an external project, *but not both*.



RELATIONSHIP CONSTRAINTS: **EXCLUSION** (CONTD)

An **exclusion (XOR) constraint** specifies that **at most one entity instance, among several entity types, can participate in a relationship with a single “root” entity.**

Example: A task can be related to *either* an internal project *or* an external project, *but not both*.



ANALYZING APPLICATION REQUIREMENTS

1. Identify entities

- What are the major concepts about which data needs to be **permanently** stored?
- Focus on the “**big picture**”, not the details.
 - E.g., student, course not name, address, email, description, credits, etc.

2. Identify relationships between entities

- How are the major concepts related? How do they interact?
- What interactions need to be **permanently** stored.
 - E.g., students enroll in courses not students browse courses

3. Identify properties of entities and relationships


- For each entity and relationship, what information needs to be **permanently** stored.

DESIGN CHOICE: ENTITY VERSUS ATTRIBUTE


entity: When **several properties** can be associated with the concept.

attribute: When the concept has a **simple atomic structure** or **no property** of interest.

Office as attribute




Office as entity



DESIGN CHOICE: PLACING AN ATTRIBUTE

? salary



Where to place salary?

Relationship attributes are usually needed only for
many to many relationships!
(But can also be used in one to one and one to many relationships.)

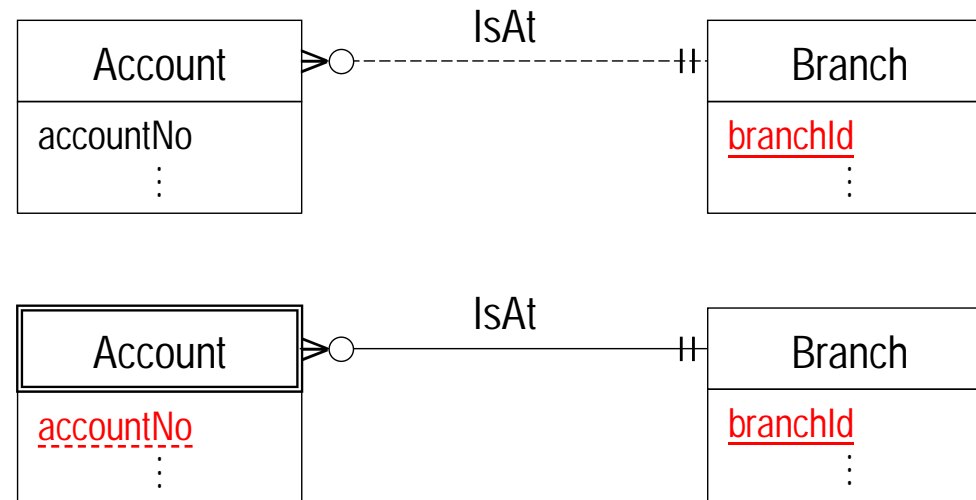
DESIGN CHOICE: STRONG VERSUS WEAK ENTITY

strong entity: When the concept can be **uniquely identified in the application domain** (i.e., it has a **key**).

weak entity: When the concept has **no unique identifier**.

Suppose an account must be associated with exactly one branch and two different branches can have accounts with the same number.

Should Account be a strong or weak entity?



DESIGN CHOICE: ENTITY VERSUS RELATIONSHIP


entity: When the concept **represents something distinct** in the application domain **with several properties**.

relationship: When the concept is **not a distinct application domain concept** and/or has **no property of interest**.

Account as entity



Account as relationship





**What if you want to have several accounts
for a customer **at the same branch**?**

DESIGN CHOICE: ENTITY VERSUS RELATIONSHIP

(CONT'D)

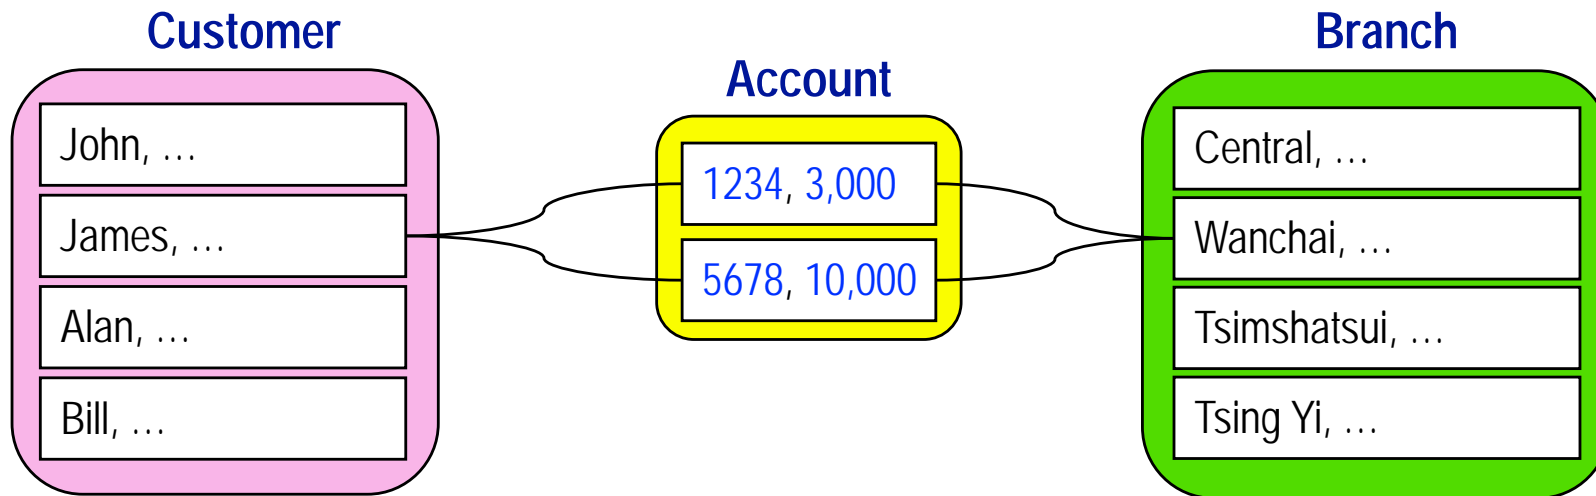

We want to represent the fact that James has two accounts at the same branch.



DESIGN CHOICE: ENTITY VERSUS RELATIONSHIP

(cont'd)

We need to use an entity for Account!

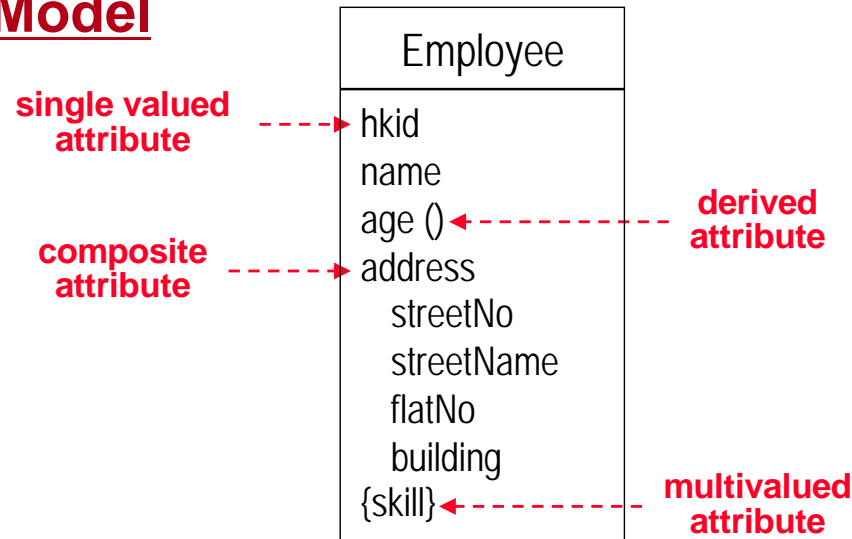


There can be only one relationship instance of a given relationship type between the same two entity instances.

E-R MODEL & DB DESIGN: SUMMARY

Entity-Relationship Model

entity and
attribute




relationship

○ - partial (min-card = 0) | - total (min-card = 1)

(a) one-to-one
(1:1)




(b) one-to-many
(1:N)




E-R MODEL & DB DESIGN: SUMMARY (CONT'D)


(c) many-to-many
(N:M)




(d) explicit
cardinality



(e) relationship
attribute




(f) recursive
(with role names)



E-R MODEL & DB DESIGN: SUMMARY (CONT'D)

(h) generalization/specialization



(i) exclusion constraint

