

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

LECTURE 17

QUERY OPTIMIZATION

QUERY OPTIMIZATION: OUTLINE

Overview

Transformation of Relational Expressions


Estimating Statistics of Expression Results

OVERVIEW

Query optimization is the process of selecting the most efficient query-evaluation plan from among many strategies.

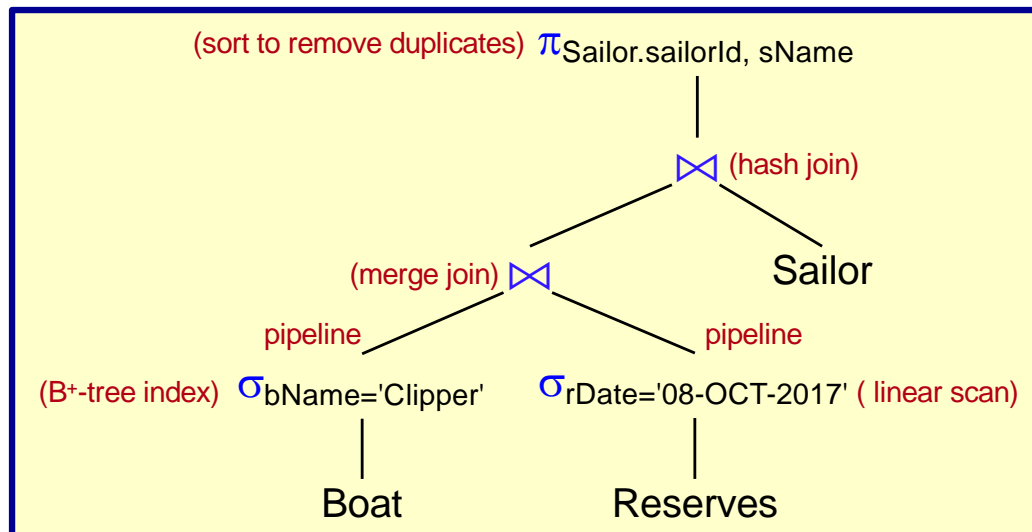
There may be several alternative ways of evaluating a given query.

- Equivalent expressions: Find the ids and names of sailors who have reserved the boat named “Clipper”.
- Different algorithms could be used for each operation.



OVERVIEW (CONT'D)

- An **evaluation plan** defines exactly what algorithm is used for each operation, and how the **execution** of the operations is coordinated.
- 👉 The cost difference between evaluation plans can be enormous (e.g., seconds versus days in some cases).



Most DBMSs allow you to view the evaluation plans for a query.

OVERVIEW (CONT'D)

Cost Based Optimization

1. Generate **logically equivalent evaluation plans** using equivalence rules for relational algebra expressions.
 2. Estimate the **cost of each plan**.
 3. Execute the plan with the **minimum expected cost**.
- Estimation of the plan cost is based on:
 - **Statistical information** about relations.
 - The number of tuples, number of distinct values for an attribute, etc.
 - **Statistics estimation** for intermediate results.
 - To compute the cost of complex expressions.
 - **Cost formulae for algorithms**, computed using statistics.

OVERVIEW (CONT'D)

Heuristic Optimization

1. Perform the **cheap operations first**.
 - Perform selections early (i.e., push them down in the relational algebra tree)
 - Reduces the number of tuples.
 2. **Remove unneeded attributes early**.
 - Perform projections early
 - Reduces the number of attributes.
 3. Try to **utilize existing indexes**.
 - Perform most restrictive selection and join operations before other similar operations.
- Transform the relational algebra tree by **using a set of rules** that typically (but not in all cases) improves execution performance.
 - ☞ **Some systems use only heuristics, others combine heuristics with partial cost-based optimization.**

TRANSFORMATION OF RELATIONAL EXPRESSIONS

- Two relational algebra expressions are **equivalent** if they *generate the same set of tuples* on *every legal* database instance.
 - The order of the tuples generated is irrelevant.
 - We do not care if they generate different results on databases that violate integrity constraints.
- SQL inputs and outputs are **multisets of tuples** (i.e., contain duplicates).
 - Two expressions in the multiset version of the relational algebra are equivalent if they generate the same multiset of tuples on every legal database instance.
- An **equivalence rule** states that two forms of an expression are equivalent.

 **Can replace an expression of the first form by one of the second, or vice versa.**

EQUIVALENCE RULES: SINGLE OPERATIONS

Cascading of projections: $\pi_{a_1}(R) \equiv \pi_{a_1}(\pi_{a_2} \dots (\pi_{a_n}(R)) \dots)$ where $a_1 \subseteq a_2 \dots \subseteq a_n$

Example: $\pi_{\text{sName}}(\text{Sailor}) \equiv \pi_{\text{sName}}(\pi_{\text{sName}, \text{sailorId}}(\text{Sailor})) \Rightarrow$ **only the final projection matters**

➤ Allows unnecessary projections to be ignored.

Cascading of selections: $\sigma_{c_1 \wedge c_2 \dots \wedge c_n}(R) \equiv \sigma_{c_1}(\sigma_{c_2} \dots (\sigma_{c_n}(R)) \dots)$

Example: $\sigma_{\text{age}=20 \wedge \text{rating}>7}(\text{Sailor}) \equiv \sigma_{\text{rating}>7}(\sigma_{\text{age}=20}(\text{Sailor}))$

Commutativity of selections: $\sigma_{c_1}(\sigma_{c_2}(R)) \equiv \sigma_{c_2}(\sigma_{c_1}(R))$

Example: $\sigma_{\text{rating}>7}(\sigma_{\text{age}=20}(\text{Sailor})) \equiv \sigma_{\text{age}=20}(\sigma_{\text{rating}>7}(\text{Sailor}))$

➤ Allows unwanted tuples to be removed early.

Commutativity of joins / Cartesian products: $R \bowtie S \equiv S \bowtie R$

Associativity of joins / Cartesian products: $(R \bowtie S) \bowtie T = R \bowtie (S \bowtie T)$

➤ Allows joins to be reordered: $R \bowtie (S \bowtie T) \equiv (T \bowtie R) \bowtie S$

EQUIVALENCE RULES: MULTIPLE OPERATIONS

Commutativity of selections with projections: $\pi_a(\sigma_c(R)) \equiv \sigma_c(\pi_a(R))$

 **Applicable only if the selection condition includes only attributes of the projection list.**

Legal Example: $\pi_{sName, age}(\sigma_{age=20}(\text{Sailor})) \equiv \sigma_{age=20}(\pi_{sName, age}(\text{Sailor}))$ ($age \in sName, age$)

Illegal Example: $\pi_{sName, rating}(\sigma_{age=20}(\text{Sailor})) \equiv \sigma_{age=20}(\pi_{sName, rating}(\text{Sailor}))$ ($age \notin sName, rating$)

➤ Remove unwanted attributes first or unwanted tuples first?

Commutativity of selections with joins / Cartesian products:

$$\sigma_c(R \text{ Join } S) \equiv (\sigma_c R) \text{ Join } S$$

 **Applicable if c involves only attributes of R .**

Legal Examples:

$$\sigma_{sName='Joe' \wedge rating=7}(\text{Sailor Join Reserves}) \equiv (\sigma_{sName='Joe' \wedge rating=7} \text{Sailor}) \text{ Join Reserves}$$

$$\sigma_{sName='Joe' \wedge boatId=100}(\text{Sailor Join Reserves}) \equiv (\sigma_{sName='Joe'} \text{Sailor}) \text{ Join } (\sigma_{boatId=100} \text{Reserves})$$

➤ Do selections first or joins first?

EQUIVALENCE RULES: MULTIPLE OPERATIONS

(cont'd)

Projection Distributes Over Join: $\pi_a(R \text{ Join } S) \equiv (\pi_a(\pi_{a_1}R) \text{ Join } (\pi_{a_2}S))$

– where

- i. a_1 is the subset of a that belongs to R *plus* the join attribute,
- ii. a_2 is the subset of a that belongs to S *plus* the join attribute.

Example: $\pi_{sName, rDate}(\text{Sailor Join Reserves}) \equiv$
 $\pi_{sName, rDate}((\pi_{sName, sailorId} \text{Sailor}) \text{ Join } (\pi_{rDate, sailorId} \text{Reserves}))$

➤ Allows unwanted attributes to be removed before doing the join.

- If the join attribute is included in the list of projection attributes a , then: $\pi_a(R \text{ Join } S) \equiv (\pi_{a_1}R) \text{ Join } (\pi_{a_2}S)$

Example: $\pi_{sailorId, rDate}(\text{Sailor Join Reserves}) \equiv (\pi_{sailorId} \text{Sailor}) \text{ Join } (\pi_{sailorId, rDate} \text{Reserves})$

JOIN ORDERING

- A good ordering of join operations is important for **reducing the size of intermediate results**.
- Consider finding the best join-order for $r_1 \text{ Join } r_2 \text{ Join } \dots \text{ Join } r_n$.
 - For this expression there are $(2(n-1))!/(n-1)!$ different join orders!
 - For $n = 3$, the number is 12.
 - For $n = 7$, the number is 665,280.
 - For $n = 10$, the number is greater than 176 billion!
 - We also need to consider other operations (e.g., selections) that may exist in the query.
- There is no need to generate all the join orders.
 - Using **dynamic programming**, the **least-cost join order** for any subset of $\{r_1, r_2, \dots, r_n\}$ can be **computed only once** and **stored for future use**.

JOIN ORDERING: DYNAMIC PROGRAMMING

- To find the best plan for a set S of n relations, consider all possible plans of the form:
 - $S_1 \text{ JOIN } (S - S_1)$ where S_1 is any non-empty subset of S .
- Recursively compute the costs for joining subsets of S to find the cost of each plan.
- Choose the cheapest of the $2^n - 1$ alternatives.
- When the plan for any subset is computed, **store it** and **reuse it** when it is required again, instead of recomputing it.

JOIN ORDERING: OPTIMIZATION ALGORITHM

procedure FindBestPlan(S)

if ($bestplan[S].cost \neq \infty$) /* $bestplan[S]$ already computed */

return $bestplan[S]$

if (S contains only 1 relation)

 set $bestplan[S].plan$ and $bestplan[S].cost$ based on best way of
 accessing S

else for each non-empty subset S_1 of S such that $S_1 \neq S$

$P_1 = \text{FindBestPlan}(S_1)$

$P_2 = \text{FindBestPlan}(S - S_1)$

A = best algorithm for joining results of P_1 and P_2

$cost = P_1.cost + P_2.cost + \text{cost of } A$

if $cost < bestplan[S].cost$



$bestplan[S].cost = cost$

$bestplan[S].plan = \text{"execute } P_1.plan; \text{ execute } P_2.plan; \text{ join results}$
 of P_1 and P_2 using A "

return $bestplan[S]$

LEFT DEEP JOIN TREES

- In **left-deep join trees**, the **right-hand-side** input for each join is a **relation**, not the result of an intermediate join.
- ☞ Some optimizers only consider **left-deep plans** since they **facilitate pipelining**. **Why?**



INTERESTING SORT ORDERS

- Consider the expression $(r_1 \text{ Join } r_2 \text{ Join } r_3) \text{ Join } r_4 \text{ Join } r_5$
- An **interesting sort order** is a particular sort order of tuples that could be **useful for a later operation**.
 - Generating the result of $(r_1 \text{ Join } r_2 \text{ Join } r_3)$ sorted on the attributes common with r_4 or r_5 may be useful, but generating it sorted on the attributes common only to r_1 and r_2 is not useful.
 - Using merge-join to compute $(r_1 \text{ Join } r_2 \text{ Join } r_3)$ may be costlier than hash join but may provide an output sorted in an interesting order.
- Need to find the best join order for each subset of the set of n relations, **for each interesting sort order**.
 - Simple extension of earlier dynamic programming algorithms.
 - Usually, the number of interesting orders is quite small and doesn't affect time/space complexity significantly.