

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

TUTORIAL 5

STORAGE AND FILE STRUCTURE

DATA FILE ORGANIZATION

Database \Leftrightarrow a collection of files


File \Leftrightarrow a sequence of records

Record \Leftrightarrow a sequence of fields

Record Organization

- Fixed-Length Records
 - Relative location of records
 - Free lists for deleted records
- Variable-Length Records
 - Byte-string representation
 - Embedded identification
 - Reserved space
 - Pointer method
 - Slotted-page Structure

FIXED-LENGTH RECORDS



- Record access is simple: $n \cdot (i-1)$
 - where n is the (fixed) record length in bytes
 - i is the record number
 - addressing starts at byte 0

Problem: what if there is a deletion?

- use a **free list**

account#, branchName, balance

header			
record 1	A-102	Perryridge	400
record 2			
record 3	A-215	Mianus	700
record 4	A-101	Downtown	500
record 5			
record 6	A-201	Perryridge	900
record 7			
record 8	A-110	Downtown	600
record 9	A-218	Perryridge	700

VARIABLE-LENGTH RECORDS

Byte-string Representation

- Attach an *end-of-record* (\perp) control character to the end of each record.

Problem: deletion and record growth.

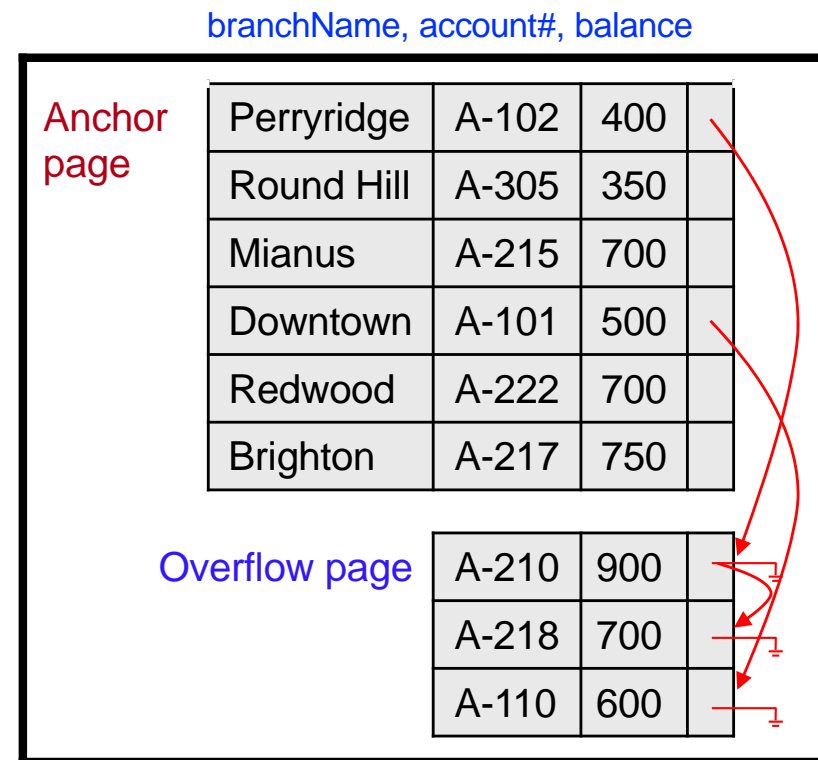
Reserved space

- Use fixed-length records of a known maximum length.

Problem: wastes storage space.

Pointer Method

- Anchor page + Overflow page
 - For Perryridge *insert*: A-210, 900
 - For Perryridge *insert*: A218, 700
 - For Downtown *insert*: A110, 600



VARIABLE-LENGTH RECORDS (CONTD)



Slotted Page Structure

- The **page header** contains:
 - the number of record entries.
 - the location of the end of the free space in the page.
 - the location and size of each record.
- The cost for moving data within a page is low.

FILE ORGANIZATION

Database \Leftrightarrow a collection of files

File \Leftrightarrow a sequence of records

Record \Leftrightarrow a sequence of fields

Heap A record can be placed anywhere in the file where there is space. There is no relationship between the search key value and a record's location.

Sequential Records are stored in sequential order, based on a search key usually the primary key).

Hash A hash function is applied to the search key value of a record; the result specifies in which bucket (page) of the file the record should be placed.

PAGE I/O COST OF OPERATIONS

Operation	Heap File	Sequential File	Hash File
Scan all records	B	B	$1.25^1 B$
Equality search ²	0.5 B	$\log_2 B^3$	1
Range search	B	$\log_2 B +$ # of pages with matches	$1.25^1 B$
Insert	2^4	Equality search + B^5	2
Delete	Equality search + 1	Equality search + B^5	2

B is the number of pages in a file.

- 1 Assumes 80% occupancy of pages to allow for future additions. Thus $1.25B$ pages are needed to store all records.
- 2 Assumes the search is on the key value.
- 3 Assumes binary search is used.
- 4 Assumes the record is inserted at the end of the file – read last page and write it back.
- 5 Assumes insert/delete is in the middle of the file and need to read and write all pages in second half of the file.

EXERCISE 1

Which file organization, heap, sequential or hash, would you choose for a file where the most frequent operations are:

a) search for records based on a range of field values?

sequential file

b) perform inserts and scans, where the order of records does not matter?

heap file

c) search for a record based on a particular field value?

hash file

EXERCISE 2

A file has 10,000 **student** records of fixed-length. Each record has the following fields: **studentId** (9 bytes), **name** (30 bytes), **address** (40 bytes), **phone** (8 bytes), **birthdate** (8 bytes), **gender** (1 byte) and **degreeProgram** (3 bytes). An additional byte is used as a deletion marker.

a) What is the size of a record in bytes?

Record size: $9+30+40+8+8+1+3+1 = 100$ bytes

b) What is the blocking factor bf_r if the page size is 4096 bytes?

bf_r : $\lfloor \# \text{ bytes per page} / \# \text{ bytes per record} \rfloor$
 $= \lfloor 4096 / 100 \rfloor = 40$ records/page

EXERCISE 2 (CONT'D)

c) How many pages are required to store the file:

i. if a sequential file organization is used?

pages: $\lceil \# \text{ records} / bf_r \rceil = \lceil 10000 / 40 \rceil = \underline{250} \text{ pages}$

ii. if a heap file organization is used?

pages: $\lceil \# \text{ records} / bf_r \rceil = \lceil 10000 / 40 \rceil = \underline{250} \text{ pages}$

iii. if a hash file organization is used (assuming 100% page occupancy)?

pages: $\lceil \# \text{ records} / bf_r \rceil = \lceil 10000 / 40 \rceil = \underline{250} \text{ pages}$

iv. if a hash file organization is used (assuming 80% page occupancy)?

pages: $\lceil \# \text{ records} / bf_r * 1.25 \rceil = \lceil 10000 / 40 * 1.25 \rceil = \underline{313} \text{ pages}$

EXERCISE 2 (CONT'D)

d) Consider the query “**Find a student record given a particular student id**”. Assuming that a record with the student id exists in the file, what is the cost, in page I/Os, to answer this query:

i. if a sequential file organization is used?

A binary search can be used.

Search cost: $\lceil \log_2 (\# \text{ of pages}) \rceil = \lceil \log_2 250 \rceil = \underline{8}$ page I/Os

ii. if a heap file organization is used?

A scan of the pages is required, but *on average only half the pages need to be scanned*.

Search cost: $\lceil (\# \text{ of pages}) / 2 \rceil = \lceil 250 / 2 \rceil = \underline{125}$ page I/Os

iii. if a hash file organization is used?

Search cost: $\underline{1}$ page I/O

EXERCISE 3

A school keeps the following file with the records of its students :

`Student(studentId: 4 bytes, name: 10 bytes, deptId: 4 bytes)`

where deptId is the department id to which a student belongs.

There are 10,000 student records and 50 departments.

A page is 128 bytes.

The data file is sorted sequentially on studentId.

- a) What is the size of a record in bytes?
- b) How many records can fit on each page?
- c) How many pages are needed to store these student records?
- d) Given this data file, what is the cost, in page I/Os, to find a particular student given a studentId?

EXERCISE 3 (CONT'D)

Student(studentId: 4 bytes, name: 10 bytes, deptId: 4 bytes)

There exist 10,000 student records and 50 departments.

A page is 128 bytes.

The data file is sorted sequentially on studentId.

a) What is the size of a record in bytes?

Record size: 4 bytes + 10 bytes + 4 bytes = 18 bytes

b) How many records can fit on each page?

$bf_{\text{Student}} = \lfloor 128 / 18 \rfloor = \underline{7}$ records/page

These records will occupy $18 * 7 = \underline{126}$ bytes on a page.

EXERCISE 3 (CONT'D)

Student(studentId: 4 bytes, name: 10 bytes, deptId: 4 bytes)

There exist 10,000 student records and 50 departments.


A page is 128 bytes.

The data file is sorted sequentially on studentId.

c) How many pages are needed to store these student records?

~~# pages: ($\# \text{records} * \text{size of record} / \text{size of page} = 10,000 * 18 / 128 = 1406.25 = 1407 \text{ pages}$)~~

7 records/page occupying 126 bytes **Do not allow records to cross page boundaries!**



EXERCISE 3 (CONT'D)

Student(studentId: 4 bytes, name: 10 bytes, deptId: 4 bytes)

There exist 10,000 student records and 50 departments.


A page is 128 bytes.

The data file is sorted sequentially on studentId.

c) How many pages are needed to store these student records?

$$\begin{aligned} \# \text{ pages: } & \lceil (\# \text{ of records} / \lfloor \text{size of page} / \text{record size} \rfloor) \rceil = \\ & \lceil (10,000 / \lfloor 128 / 18 \rfloor) \rceil = \underline{1429} \text{ pages} \end{aligned}$$

7 records/page occupying 126 bytes 7 records/page



Page size: 128 bytes

EXERCISE 3 (CONT'D)

Student(studentId: 4 bytes, name: 10 bytes, deptId: 4 bytes)

There exist 10,000 student records and 50 departments.

A page is 128 bytes.

The data file is sorted sequentially on studentId.

d) Given this data file, what is the cost, in page I/Os, to find a particular student given a studentId?

pages: 1429

Binary search on studentId: $\lceil \log_2(\# \text{ of pages}) \rceil$
 $= \lceil \log_2 1429 \rceil = \underline{11} \text{ page I/Os}$