

# COMP 3311

# DATABASE MANAGEMENT

# SYSTEMS

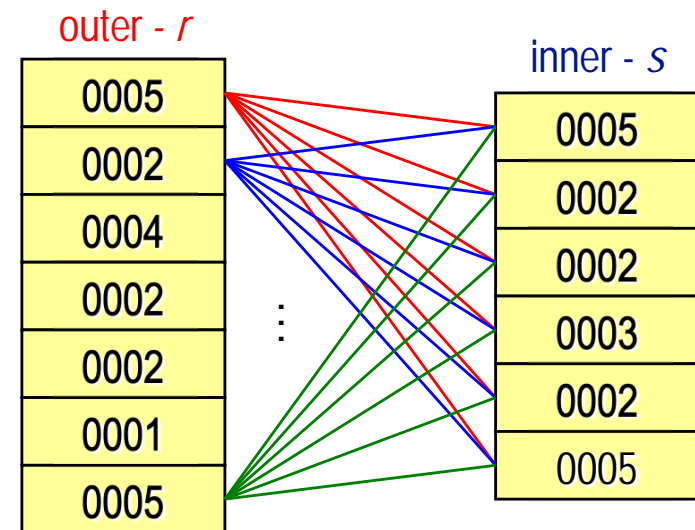
## TUTORIAL 7

## QUERY PROCESSING

# REVIEW: BLOCK NESTED-LOOP JOIN

- Read in the outer relation  $r$  page-by-page.
- For each page of  $r$ , scan the entire inner relation  $s$ .
- Best join page I/O cost:  $b_r + b_s$ 
  - where  $b_r$  and  $b_s$  are the number of pages of  $r$  and  $s$ , respectively, and the inner relation,  $s$ , is small enough to fit in the buffer.
- Buffer pages needed: **at least 3** (1 for  $r$ , 1 for  $s$ , 1 for the output).
  - If there are  $M$  buffer pages, read  $M-2$  pages of  $r$  at a time and use the remaining two buffer pages for  $s$  and the output.

Join page I/O cost:  $\lceil b_r / (M-2) \rceil * b_s + b_r$




## REVIEW: INDEXED NESTED-LOOP JOIN

- An index lookup can replace a file scan if an **index is available** on the **join attribute** of the **inner relation**.
- For each tuple  $t_r$  in the **outer relation**  $r$ , use the index to look up tuples in the **inner relation**  $s$  that satisfy the join condition with tuple  $t_r$ .

Join page I/O cost:  $b_r + n_r * c$

- $n_r$  is the number of tuples in  $r$ .
- $c$  is the cost to traverse the index and fetch all matching  $s$  tuples for one tuple,  $t_r$ , of  $r$ .
- $c$  can be estimated as the cost of a single selection on  $s$  using the join condition.



If indexes are available on the join attribute of both  $r$  and  $s$ , use the relation with **fewer tuples** as the **outer relation** as this will result in fewer index lookups.

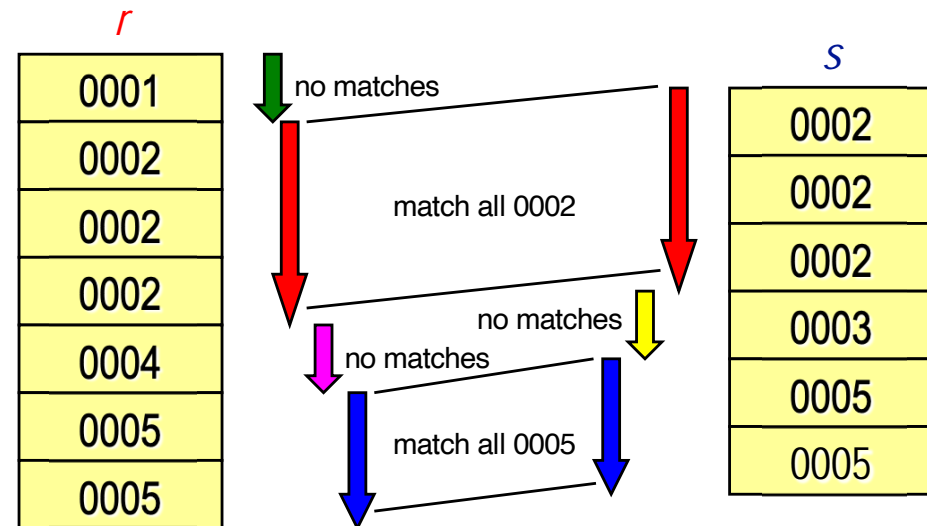
# REVIEW: SORT-MERGE JOIN

👉 **Applicable for equi-joins and natural joins.**

- Sort both relations on the join attribute (if not already sorted).
- Merge the sorted relations to join them.
  - The join step is like the merge phase of the merge-sort algorithm.
  - The main difference is the handling of duplicate values in the join attribute — **every pair with the same value** on the join attribute **must be matched**.
- Each block needs to be read only once.

Join page I/O cost:

$b_r + b_s$  + cost of sorting if the relations are not sorted




# REVIEW: HASH JOIN

👉 Applicable for equi-joins and natural joins.

- A hash function  $h$  is used to place tuples of both relations into  $n$  partitions (buckets) (i.e., a hash file organization).

👉 Partitions the tuples of each of the relations into sets that have the same hash value on the join attributes.

- Only need to compare  $r$  tuples in partition  $r_i$  with  $s$  tuples in partition  $s_i$ .
- Do not** need to compare  $r$  tuples in partition  $r_i$  with  $s$  tuples in any other partition, since:
  - An  $r$  tuple and an  $s$  tuple that satisfy the join condition will have the same join attribute value.
  - Hence, they will hash to the same value  $i \Rightarrow$  the  $r$  tuple has to be in partition  $r_i$  and the  $s$  tuple in partition  $s_i$ !



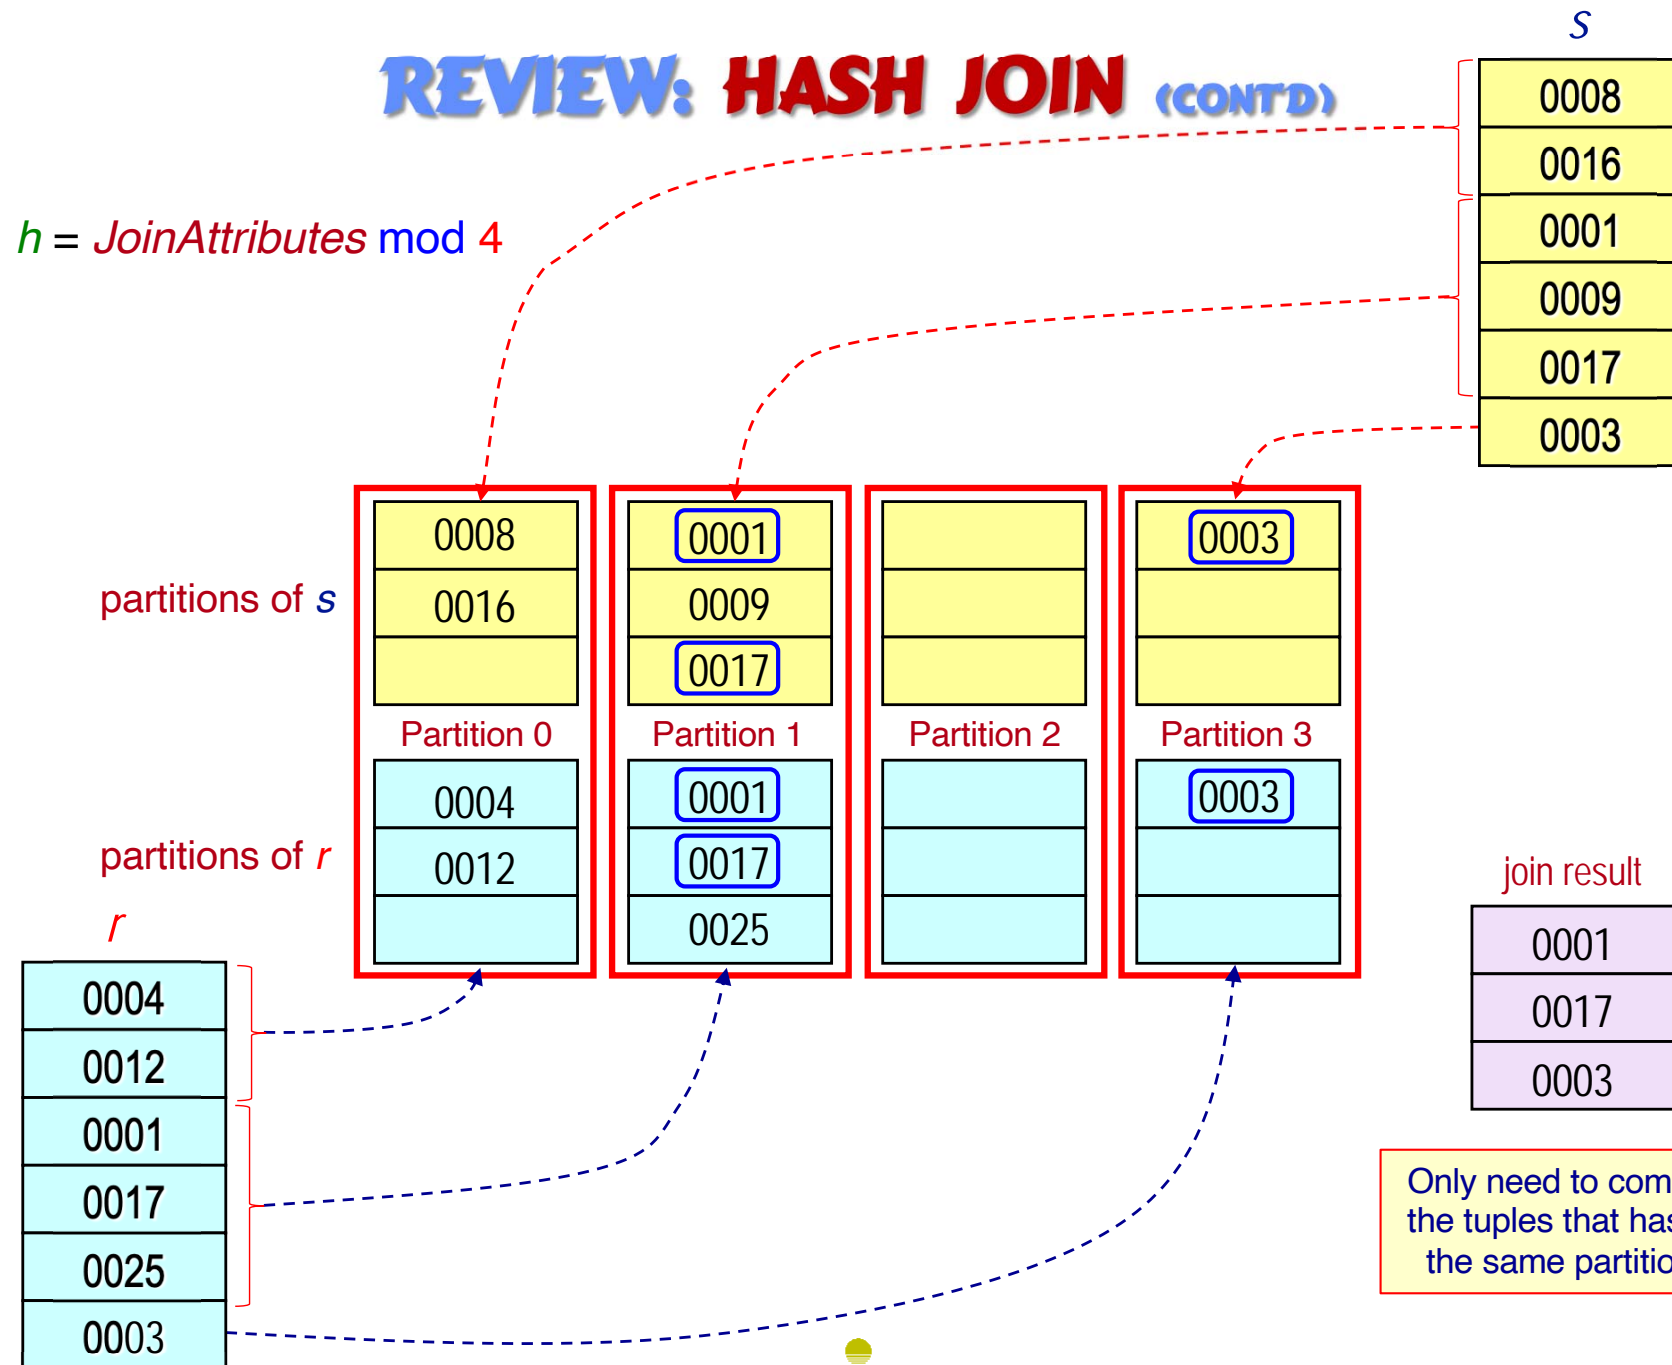
Join page I/O cost:

$$3 * (B_r + B_s)$$

- $\Rightarrow$  1 read and 1 write to create the partitions;  
1 read to compute the join.

# REVIEW: HASH JOIN (CONTD)

$$h = \text{JoinAttributes} \bmod 4$$



## EXERCISE 1

Relations:  $R_1(A, B, C)$  and  $R_2(\underline{C}, D, E)$

$R_1$  20,000 tuples

$bf_{R_1} = 25$  tuples/page

$R_1$  pages: 800

$R_2$  45,000 tuples

$bf_{R_2} = 30$  tuples/page

$R_2$  pages: 1500

Assume:

- 100 buffer pages.
- $R_2$  has a B<sup>+</sup>-tree index with 3 levels on the join attribute  $C$ , the primary key of  $R_2$ .
- $R_1$  and  $R_2$  are not initially sorted on the join attribute.

**Estimate** the number of page I/Os required, *in the worst case*, using each of the following join algorithms for  $R_1 \bowtie R_2$ :

- Optimized block nested-loop join
- Indexed nested-loop join
- Sort-merge join
- Hash join using 10 buckets

## EXERCISE 1 (CONTD)

$R_1$  tuples: 20,000;  $bf_{R_1}$ : 25  
# pages  $R_1$ : 800  
 $R_2$  tuples: 45,000;  $bf_{R_2}$ : 30  
# pages  $R_2$ : 1500  
 $R_2.C$ : 3 level B<sup>+</sup>-tree index  
 $M$ : 100

### a) Optimized block nested-loop join

– The worst-case cost =  $\lceil b_r / (M - 2) \rceil * b_s + b_r$

i. When  $R_1$  is the outer relation

Join page I/O cost:  $\lceil 800 / (100 - 2) \rceil * 1500 + 800 = \underline{14,300}$

ii. When  $R_2$  is the outer relation

Join page I/O cost:  $\lceil 1500 / (100 - 2) \rceil * 800 + 1500 = \underline{14,300}$

### b) Indexed nested-loop join (B<sup>+</sup>-tree index on $R_2.C$ with 3 levels)

– The worst-case cost =  $b_r + n_r * c$

Join page I/O cost:  $800 + (3 + 1) * 20,000 = \underline{80,800}$



## EXERCISE I (CONTD)

$R_1$  tuples: 20,000;  $bf_{R_1}$ : 25  
# pages  $R_1$ : 800  
 $R_2$  tuples: 45,000;  $bf_{R_2}$ : 30  
# pages  $R_2$ : 1500  
 $R_2.C$ : 3 level B<sup>+</sup>-tree index  
 $M$ : 100

### c) Sort-merge join

The worst-case cost = sorting cost +  $b_r$  +  $b_s$

#### i. Sorting cost of $R_1$

Sort page I/O cost:  $800 * 2 * (\lceil \log_{100-1}(800/100) \rceil + 1) = \underline{3,200}$

#### ii. Sorting cost of $R_2$

Sort page I/O cost:  $1500 * 2 * (\lceil \log_{100-1}(1500/100) \rceil + 1) = \underline{6,000}$

#### iii. Merge cost for $R_1$ JOIN $R_2$

Merge page I/O cost:  $1500 + 800 = \underline{2,300}$

Join page I/O cost:  $3200 + 6000 + 2300 = \underline{11,500}$

**Can you improve on this cost?**

## EXERCISE I (CONTD)

$R_1$  tuples: 20,000;  $bf_{R_1}$ : 25  
# pages  $R_1$ : 800  
 $R_2$  tuples: 45,000;  $bf_{R_2}$ : 30  
# pages  $R_2$ : 1500  
 $R_2.C$ : 3 level B<sup>+</sup>-tree index  
 $M$ : 100

### Improved Sort-Merge Join

Previously we first sorted and merged  $R_1$  and  $R_2$  into 1 sorted run each before doing the join. Instead, we first only sort  $R_1$  and  $R_2$  into runs, but do not merge the sorted runs. Then, since we have 100 buffer pages, we can do an  $8 + 15 = 23$ -way merge and join during the merge.

- Create 8 sorted runs of  $R_1$  and write them out.

Sort page I/O cost:  $800 + 800 = \underline{1,600}$

- Create 15 sorted runs of  $R_2$  and write them out

Sort page I/O cost:  $1500 + 1500 = \underline{3,000}$

- Read the 8 sorted runs of  $R_1$  and the 15 sorted runs of  $R_2$  using 1 buffer page *per run* (i.e., a total of 23 buffer pages) and join them during the merge.

Merge page I/O cost:  $800 + 1500 = \underline{2,300}$

Join page I/O cost:  $1,600 + 3,000 + 2,300 = \underline{6,900}$

## EXERCISE I (CONTD)

$R_1$  tuples: 20,000;  $bf_{R_1}$ : 25  
# pages  $R_1$ : 800  
 $R_2$  tuples: 45,000;  $bf_{R_2}$ : 30  
# pages  $R_2$ : 1500  
 $R_2.C$ : 3 level B<sup>+</sup>-tree index  
 $M$ : 100

### d) Hash join (using 10 partitions)

- i. Use  $R_1$  as the **build input** (since it is smaller)

Partition  $R_1$  into 10 partitions, each of size 80 pages. This partitioning can be done in one pass since we use 10 buffer pages for the 10 partitions and 1 buffer page to read  $R_1$  page-by-page. Write a partition page to disk when it becomes full and continue partitioning.

- ii. Use  $R_2$  as the **probe input**

Partition  $R_2$  into 10 partitions, each of size 150 pages. This is also done in one pass like the way we partition  $R_1$ .

- iii. Since we have 100 buffer pages, we read each partition (i.e., 80 pages) of the build input  $R_1$ , in turn, into the buffer. For each  $R_1$  partition, read the corresponding probe partition  $R_2$  into the buffer page-by-page (requires only 1 buffer page) and probe the build partition for matches.

Join page I/O cost:  $3 * (800 + 1500) = 6,900$

## EXERCISE 2

Given relation  $R(\underline{A}, B, C, D, E)$ , organized as an ordered sequential file on search key  $A$ , and the information below, answer the questions.

Tuple size: 200 bytes      Attribute  $A$ : 16 bytes      Page size: 2400 bytes  
Number of tuples: 500,000      Pointer size: 4 bytes

- a) How many pages are required to store  $R$ ?
- b) How many index pages are required if the search key  $A$  is organized using a static, multi-level index (i.e., not a B<sup>+</sup>-tree index)?
- c) For the query `select * from R where A=xxx`, determine the page I/O cost for each of the query evaluation strategies given below.
  - i. linear search      ii. binary search      iii. index search
- d) For the query `select * from R where A>700000`, what is the page I/O cost to answer this query using the index from b) assuming that  $A$  is uniformly distributed on the interval  $[200,000; 800,000]$  and the leaf index pages are chained?

## EXERCISE 2 (CONTD)

tuple size: 200 bytes  
# tuples: 500,000  
attribute A: 16 bytes  
pointer size: 4 bytes  
page size: 2400 bytes

- a) How many pages are required to store R?

$$bf_R = \lfloor \text{page size} / \text{tuple size} \rfloor = \lfloor 2400 / 200 \rfloor = 12$$

$$\text{Number of pages needed: } \lceil \# \text{tuples} / bf_R \rceil = \lceil 500000 / 12 \rceil = 41,667$$

- b) How many index pages are required if the search key A is organized using a static, multi-level index (i.e., not a B<sup>+</sup>-tree index) ?

Since the file is organized as a sequential file on search key A, the tuples are stored in search-key order. Therefore, the index is sparse.

The number of index entries in the leaf nodes is equal to the number of pages of the file (one index entry points to each page of the file).

$$bf_{indexA} = \lfloor \text{page size} / \text{index entry size} \rfloor = \lfloor 2400 / (16+4) \rfloor = 120$$

$$\# \text{pages}_{\text{level1}} = \lceil \# \text{index entries} / bf_{indexA} \rceil = \lceil 41667 / 120 \rceil = 348$$

$$\# \text{pages}_{\text{level2}} = \lceil \# \text{level1 index entries} / bf_{indexA} \rceil = \lceil 348 / 120 \rceil = 3$$

$$\# \text{pages}_{\text{level3}} = \lceil \# \text{level2 index entries} / bf_{indexA} \rceil = \lceil 3 / 120 \rceil = 1$$

$$\text{Total number of index pages: } 348 + 3 + 1 = 352$$

## EXERCISE 2 (CONTD)

tuple size: 200 bytes  
# tuples: 500,000  
attribute A: 16 bytes  
pointer size: 4 bytes  
page size: 2400 bytes

c) For the query `select * from R where A=xxx`, determine the page I/O cost for each of the query evaluation strategies given below.

i. linear search

Query page I/O cost:  $\lceil \#pages / 2 \rceil = \lceil 41667 / 2 \rceil = \underline{20,834}$

ii. binary search

Query page I/O cost:  $\lceil \log_2(\#pages) \rceil = \lceil \log_2(41667) \rceil = \underline{16}$

iii. index search using the index from b)

Query page I/O cost: height of the index + 1 = 3 + 1 = 4

## EXERCISE 2 (CONT'D)

tuple size: 200 bytes  
# tuples: 500,000  
attribute A: 16 bytes  
pointer size: 4 bytes  
page size: 2400 bytes

- d) For the query `select * from R where A > 700000`, what is the page I/O cost to answer this query using the index from b) assuming that A is uniformly distributed on the interval [200,000; 800,000] and the leaf index pages are chained?

Since A is uniformly distributed on the interval [200,000; 800,000], we can **estimate the proportion of the pages** that will be retrieved as

$$(800000 - 700000) / (800000 - 200000) = (100000) / (600000) = 1/6$$

Thus, we expect to retrieve  $1/6^{\text{th}}$  of the relation's pages.

Query page I/O cost:  $\lceil \# \text{index levels} + 1/6 * \# \text{pages} \rceil$

$$\begin{aligned} &= \lceil 3 + 1/6 * 41667 \rceil \\ &= \lceil 3 + 6944.5 \rceil \\ &= \underline{6,948} \end{aligned}$$