

COMP 3311

DATABASE MANAGEMENT

SYSTEMS

TUTORIAL 8

QUERY PROCESSING /

QUERY OPTIMIZATION

REVIEW: MATERIALIZATION AND PIPELINING

Materialization

- Generate results of an expression whose inputs are relations or are already computed and materialize (store) it on disk.
- Repeat until query is processed.

Overall cost: Sum of costs of individual operations + cost of writing intermediate results to and reading intermediate results from disk.

Pipelining

- Pass on result tuples of an operation to parent operations even as the operation is being executed.
- Avoids writing/reading intermediate results to/from disk.

Overall cost: Sum of costs of individual operations.

EXERCISE 1

Consider the relations $R_1(\underline{A}, B, \textit{C})$, $R_2(\underline{C}, D, \textit{E})$ and $R_3(\underline{E}, F)$. Primary keys are underlined, and all foreign keys are in italics and not null.

Assume that:

R_1 has 1,000 tuples in 100 pages.

R_2 has 10,000 tuples in 1,000 pages.

R_3 has 100,000 tuples in 10,000 pages.

a) What is the query result size of $R_1 \bowtie R_2 \bowtie R_3$?

b) Give an efficient **pipelining strategy** to compute $R_1 \bowtie R_2 \bowtie R_3$.

$R_1(\underline{A}, B, C)$	1,000 tuples	100 pages
$R_2(\underline{C}, D, E)$	10,000 tuples	1,000 pages
$R_3(\underline{E}, F)$	100,000 tuples	10,000 pages

EXERCISE 1 (cont'd)

a) What is the query result size of $R_1 \bowtie R_2 \bowtie R_3$?

- Since all joins are on key values, the query result size is equal to the size of the smallest relation R_1 , which is 1,000 tuples.

 $(R_1 \bowtie R_2) \bowtie R_3$ is less costly than $R_1 \bowtie (R_2 \bowtie R_3)$. **Why?**

For $(R_1 \bowtie R_2) \bowtie R_3$

$R_1 \bowtie R_2$ will generate 1,000 tuples; then the join with R_3 will also generate 1,000 tuples.

For $R_1 \bowtie (R_2 \bowtie R_3)$

$R_2 \bowtie R_3$ will generate 10,000 tuples; then the join with R_1 will generate 1,000 tuples.

- Even though query result is the same, the computation for $(R_1 \bowtie R_2) \bowtie R_3$ will be less costly since fewer total tuples are generated.

$R_1(\underline{A}, B, C)$	1,000 tuples	100 pages
$R_2(\underline{C}, D, E)$	10,000 tuples	1,000 pages
$R_3(\underline{E}, F)$	100,000 tuples	10,000 pages

EXERCISE 1 (cont'd)

b) Give an efficient **pipelining strategy** to compute $(R_1 \bowtie R_2) \bowtie R_3$.

- **create an index** on attribute C for relation R_2 .
- **create an index** on attribute E for relation R_3 .

As the condition is equality on a key attribute, **a hash index is best**.

Then, for each of the **1,000** tuples in R_1 , we do the following:

a. To compute $R_1 \bowtie R_2$, search in R_2 to find the match with the C value of R_1 .

Page I/O cost: 2

1 page I/O to find the entry C in the hash index (assuming no overflow buckets).

1 page I/O to retrieve the tuple from the relation R_2 .

b. Then, to compute the join with R_3 for each of the **1,000** result tuples of $(R_1 \bowtie R_2)$, search in R_3 to look up at most **1** tuple which matches the unique value for E in R_2 . Page I/O cost: 2

Again, we need **1** page I/O to access the index and **1** page I/O to retrieve the tuple.

Query processing page I/O cost: $1000 * (2 + 2) + 100 = 4100$



(100 is the cost for reading the **1,000** tuples of R_1)

This plan corresponds to an indexed nested-loop join and ignores the cost of building the indexes.

REVIEW: QUERY OPTIMIZATION

Heuristic Optimization

- Perform selections early.
- Perform projections early.
- Perform most restrictive selections and join operations before other similar operations \Rightarrow create **left-deep join trees**.



EXERCISE 2

Given the following relations and the information about them, process the relational algebra tree for the query below using a **pipelined plan** and answer the following questions.

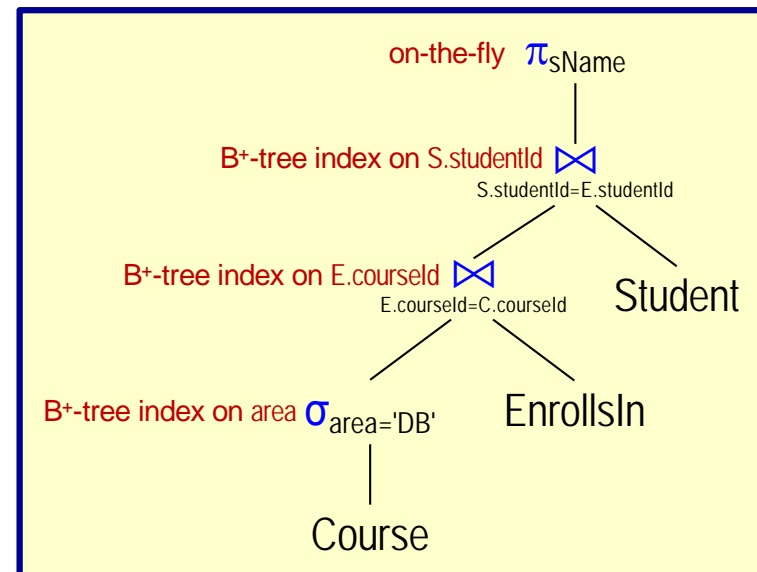
Student(<u>studentId</u> , sName, gender)	1,000 tuples; 100 pages; index on studentId
EnrollsIn(<u>studentId</u> , <u>courseId</u> , year)	6,000 tuples; 600 pages; index on courseId
Course(<u>courseId</u> , cName, area, credit)	200 tuples; 40 pages; index on area 10 different areas, with 20 tuples per area (i.e., 20 different courses per area)

All foreign keys are not null.

All indexes are B⁺-tree clustering indexes with 4 levels.

EnrollsIn tuples are **uniformly distributed**

```
select sName
from Student S, EnrollsIn E, Course C
where S.studentId=E.studentId
and E.courseId=C.courseId
and area= 'DB';
```

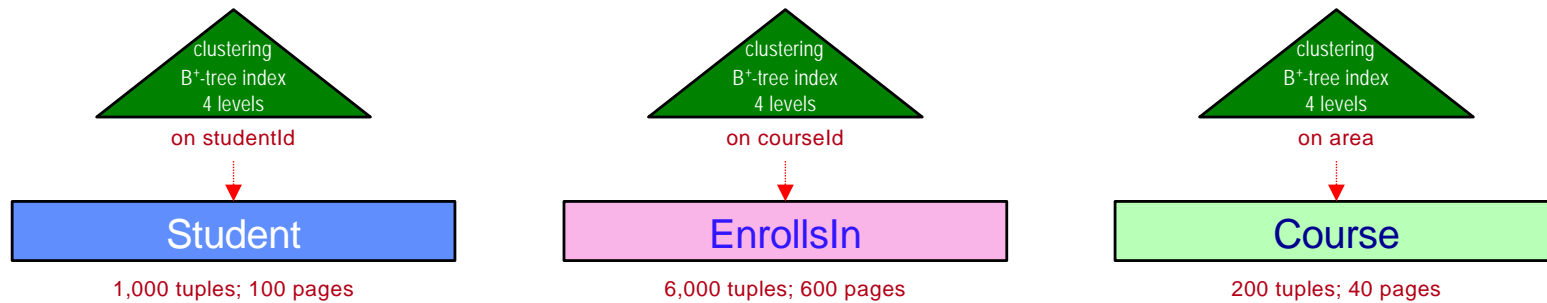


EXERCISE 2 (CONT'D)

Student(studentId, sName, gender)

EnrollsIn(studentId, courseId, year)

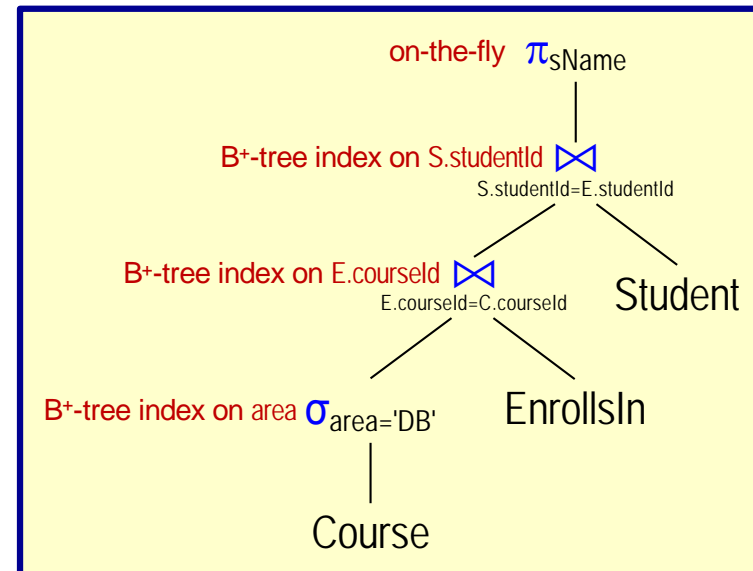
Course(courseId, cName, area, credit)



```
select sName
from Student S, EnrollsIn E, Course C
where S.studentId=E.studentId
and E.courseId=C.courseId
and area= 'DB';
```

Some useful statistics:

- 10 different course areas.
- 20 tuples per course area.



Student(studentId, sName, gender)
 EnrollsIn(studentId, courseld, year)
 Course(courseld, cName, area, credit)

EXERCISE 2 (cont'd)

Student: 1000 tuples; 100 pages; index on studentId
 EnrollsIn: 6000 tuples; 600 pages; index on courseld
 Course: 200 tuples; 40 pages; index on area
 10 different areas 20 tuples per area

a) Estimate the query result size.

Due to the commutativity of selections and joins, we can rewrite the query as $\pi_{sName}(\sigma_{area='DB'}(Student \bowtie_{studentId} EnrollsIn \bowtie_{courseld} Course))$ (i.e., apply the selection after the joins). This will make the estimation clearer.

$EnrollsIn \bowtie_{courseld} Course$

Since courseld is a not null foreign key referencing Course in EnrollsIn, the size of $(EnrollsIn \bowtie_{courseld} Course)$ is the same as the size of EnrollsIn (i.e., 6000 tuples).

$Student \bowtie_{studentId} EnrollsIn \bowtie_{courseld} Course$

Since studentId is a not null foreign key referencing Student in $(EnrollsIn \bowtie_{courseld} Course)$, the size of $(Student \bowtie_{studentId} (EnrollsIn \bowtie_{courseld} Course))$ is the same as the size of $(EnrollsIn \bowtie_{courseld} Course)$, which is the same as the size of EnrollsIn (i.e., 6000 tuples).

Student(studentId, sName, gender)
EnrollsIn(studentId, courseId, year)
Course(courseId, cName, area, credit)

EXERCISE 2 (cont'd)

Student: 1000 tuples; 100 pages; index on studentId
EnrollsIn: 6000 tuples; 600 pages; index on courseId
Course: 200 tuples; 40 pages; index on area
10 different areas 20 tuples per area

$\sigma_{\text{area}='DB'}(\text{Student} \bowtie_{\text{studentId}} \text{EnrollsIn} \bowtie_{\text{courseId}} \text{Course})$

Since the selectivity of the condition $\text{area}='DB'$ is $1/10$ (i.e., 10 different areas), the size of $\sigma_{\text{area}='DB'}(\text{Student} \bowtie_{\text{studentId}} \text{EnrollsIn} \bowtie_{\text{courseId}} \text{Course})$ is $6000/10 = 600$ tuples.

Query result size: 600 tuples

Student(studentId, sName, gender)
EnrollsIn(studentId, courseId, year)
Course(courseId, cName, area, credit)

EXERCISE 2 (cont'd)

Student: 1000 tuples; 100 pages; index on studentId
EnrollsIn: 6000 tuples; 600 pages; index on courseId
Course: 200 tuples; 40 pages; index on area
10 different areas 20 tuples per area

b) Estimate the query page I/O cost using a pipelined plan.

Since we are told to use a pipelined plan and are not told how much buffer memory is available, we can assume that all intermediate results can be kept in memory.

Step 1: $\sigma_{\text{area}='DB'} \text{Course} \Rightarrow \text{result A}$

Strategy: index lookup using B⁺-tree on area

$V(\text{Course}, \text{area}) = 10$ distinct area values.

Thus, there are $200/10 = 20$ Course tuples (or $40/10 = 4$ Course pages) whose area='DB'. Note that Course must be ordered on area due to the clustering index on area.

👉 There are 20 distinct course ids for area='DB'.

To get all these Course tuples requires

4 index page I/Os, since each B⁺-tree index has 4 levels.

4 Course page I/Os, since Course is ordered on area.

Step 1 page I/O cost: 8

Student(studentId, sName, gender)
EnrollsIn(studentId, courseId, year)
Course(courseId, cName, area, credit)

EXERCISE 2 (cont'd)

Student: 1000 tuples; 100 pages; index on studentId
EnrollsIn: 6000 tuples; 600 pages; index on courseId
Course: 200 tuples; 40 pages; index on area
10 different areas 20 tuples per area

Step 2: result A \bowtie EnrollsIn \Rightarrow result B

Strategy: indexed nested-loop join using courseId B⁺-tree index

For each courseId selected from Course, the clustering index on courseId for EnrollsIn is used to retrieve all related EnrollsIn tuples and get the corresponding studentId.

$V(\text{Course, courseId}) = 200$ distinct courseId values.

We are told to assume that courses are uniformly distributed among EnrollsIn tuples. Thus, for each courseId value there are $6000/200 = 30$ tuples (or $600/200 = 3$ pages) in EnrollsIn. Note that EnrollsIn must be ordered on courseId due to the clustering index on courseId.

To get all the EnrollsIn tuples *for each courseId value* requires

4 index page I/Os, since each B⁺-tree index has 4 levels.

3 EnrollsIn page I/Os, since EnrollsIn is ordered on courseId.

Page I/O cost per courseId: 7

There are 20 distinct courseId values for area='DB'.

Step 2 page I/O cost: $20 * 7 = 140$

Student(studentId, sName, gender)
EnrollsIn(studentId, courseId, year)
Course(courseId, cName, area, credit)

EXERCISE 2 (cont'd)

Student: 1000 tuples; 100 pages; index on studentId
EnrollsIn: 6000 tuples; 600 pages; index on courseId
Course: 200 tuples; 40 pages; index on area
10 different areas 20 tuples per area

Step 3: result B ⋈ Student

Strategy: indexed nested-loop join using studentId B⁺-tree index

For each studentId selected from EnrollsIn, the clustering index is used to retrieve the related tuple (in Student) and get the corresponding sName.

To get the Student tuple *for each studentId value* requires

4 index page I/Os, since each B⁺-tree index has 4 levels.

1 Student page I/O, to get the Student tuple.

Page I/O cost per studentId: 5

There are 20 courseId values and 30 EnrollsIn tuples for each courseId value, Thus, there are $20 \times 30 = 600$ studentId values.

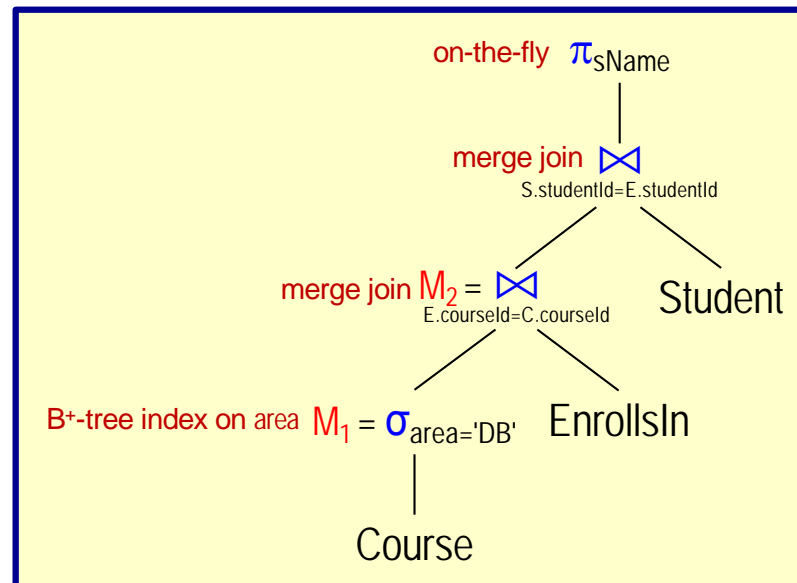
Step 3 page I/O cost: $600 \times 5 = 3,000$

Query processing page I/O cost: $8 + 140 + 3,000 = 3,148$

EXERCISE 3

Suppose the intermediate results $M_1 = \sigma_{\text{area}='DB'}$ and $M_2 = M_1 \bowtie_{\text{courseId}} \text{EnrollsIn}$ are materialized for the relational algebra tree shown below. What is the query processing page I/O cost for this relational algebra tree assuming that merge join is used for all joins?

Assume there are 22 buffer pages and attributes in the same relation all have the same size.



Student(studentId, sName, gender)
EnrollsIn(studentId, courseld, year)
Course(courseld, cName, area, credit)

EXERCISE 3 (cont'd)

Student: 1000 tuples; 100 pages; index on studentId
EnrollsIn: 6000 tuples; 600 pages; index on courseld
Course: 200 tuples; 40 pages; index on area
10 different areas 20 tuples per area

Step 1: Cost to materialize $M_1 = \sigma_{\text{area}='DB'}$

Strategy: index lookup using B⁺-tree on area

From the Exercise 2 analysis we know that to evaluate $\sigma_{\text{area}='DB'}$ requires 8 page I/Os and results in 4 pages of Course tuples.

Since we have 22 buffer pages, we can sort these 4 pages of Course tuples on courseld in memory and materialize the result (i.e., write it to disk) as M_1 .

We need to keep only courseld in the result since only the courseld is needed to join with EnrollsIn.

Since we assume that all attributes in the same relation have the same size and there are 4 attributes in Course, the size of the output M_1 is 4 pages/4 = 1 page.

Step 1 page I/O cost to materialize M_1 : $8_r + 1_w = 9$

Student(studentId, sName, gender)
 EnrollsIn(studentId, courseId, year)
 Course(courseId, cName, area, credit)

EXERCISE 3 (cont'd)

Student: 1000 tuples; 100 pages; index on studentId
 EnrollsIn: 6000 tuples; 600 pages; index on courseId
 Course: 200 tuples; 40 pages; index on area
 10 different areas 20 tuples per area

Step 2: Cost to materialize $M_2 = M_1 \bowtie_{\text{courseId}} \text{EnrollsIn}$

Strategy: merge join

The merge join cost is 1 page I/O to read M_1 and 600 page I/Os to read EnrollsIn. (EnrollsIn is sorted on courseId due to B⁺-tree clustering index.)

We keep only studentId since only studentId is needed to join with Student.

From the Exercise 2 analysis we know that for each courseId value we access 3 pages. Since we only keep the studentId in the result and EnrollsIn has 3 attributes, the size of the result is 3 pages/3 = 1 page for each courseId value. Since there are 20 distinct courseId values for the area 'DB', the join result size M is $1 * 20 = 20$ pages.

With 22 buffer pages, 2 pages can be used for reading M_1 and EnrollsIn and the remaining 20 pages can hold the join result. Since EnrollsIn has a clustering B⁺-tree index on courseId, there is no need to sort it.

We next do an in-memory sort of the result (on studentId) and materialize it as M_2 requiring 20 page I/Os to write M_2 to disk.

Step 2 page I/O cost to materialize M_2 : $1_r + 600_r + 20_w = 621$

Student(studentId, sName, gender)
EnrollsIn(studentId, courseld, year)
Course(courseld, cName, area, credit)

EXERCISE 3 (cont'd)

Student: 1000 tuples; 100 pages; index on studentId
EnrollsIn: 6000 tuples; 600 pages; index on courseld
Course: 200 tuples; 40 pages; index on area
10 different areas 20 tuples per area

Step 3: Cost to compute $M_2 \bowtie_{\text{studentId}}$ Student

Strategy: merge join

20 page I/Os are required to read M_2 and 100 page I/Os to read Student.

Since Student has a clustering B⁺-tree index on studentId, it is already sorted on studentId.

Step 3 page I/O cost: $20_r + 100_r = \underline{120}$

Query processing page I/O cost: $9 + 621 + 120 = \underline{750}$