

# COMP 3311: Database Management Systems

## Tutorial 9 Query Optimization

**Exercise 1:** Sailor(sailorId, sName, rating, age)      6,000 tuples      10 tuples/page  
Reserves(sailorId, boatId, rDate)      1,500 tuples      15 tuples/page  
Boat(boatId, bName, color)      3,000 tuples      20 tuples/page

There is no index on any relation. 15% of boats are red.

- a) Estimate the size in tuples of  $\pi_{\text{boatId}}(\sigma_{\text{color}='red'}\text{Boat})$ .
- b) Estimate the size in tuples of (Sailor Join Reserves Join Boat). Briefly explain the reason(s) for your answer.
- c) i. Calculate the result size in tuples of the join order (Sailor Join Boat) Join Reserves
- ii. Calculate the result size in tuples of the join order (Sailor Join Reserves) Join Boat
- iii. Which join order is better and why?

# COMP 3311: Database Management Systems

## Tutorial 9 Query Optimization

**Exercise 2:** For the query  $\pi_{A,B,C,D}(R \bowtie_{A=C} S)$  assume that the projection routine uses external sorting, eliminates all unwanted attributes during the initial sort pass and removes duplicate tuples on-the-fly during the merge passes. Furthermore, assume the following:

- R is 10 pages; each R tuple is 300 bytes.
- S is 100 pages; each S tuple is 500 bytes.
- The combined size of attributes A, B, C and D is 450 bytes.
- A and B are in R and have combined size 200 bytes; C and D are in S.
- A is a key for R.
- Each S tuple joins with exactly one R tuple.
- The page size is 1024 bytes.
- The buffer size  $M$  is 3 pages.
- Only the (optimized) block nested-loop join method is implemented.

What is the query processing page I/O cost?

- a) Join Cost (using (optimized) block nested-loop and eliminating unwanted attributes during the join)

Join page I/O cost:

Write join result page I/O cost:

- b) Projection Cost (using external sorting)

Sort page I/O cost pass 0:

Merge page I/O cost:

Query processing page I/O cost:

Name: (1) \_\_\_\_\_ / \_\_\_\_\_ Student#: (1) \_\_\_\_\_ Date: \_\_\_\_\_  
Family/Given (PRINT) Given/First (PRINT)

Name: (2) \_\_\_\_\_ / \_\_\_\_\_ Student#: (2) \_\_\_\_\_  
Family/Given (PRINT) Given/First (PRINT)

**NOTE: You are highly encouraged to do this exercise with a partner.**

**Exercise 3:** What is the query processing page I/O cost if merge join is used for Exercise 2 instead of block-nested loop using a buffer of 3 pages?

a) Cost to sort R

Sort pass page I/O cost:

Merge pass page I/O cost:

Page I/O cost to sort R:

b) Cost to sort S

Sort pass page I/O cost:

Merge pass page I/O cost:

Page I/O cost to sort s:

c) Cost to join

Join page I/O cost:

Query processing page I/O cost:

You must upload this completed exercise sheet to Canvas by 11 p.m. today.