

COMP 4021
Internet Computing

Cookies and HTTP

Stateless Protocols

- ❑ HTTP is a stateless **protocol**, and so are all internet protocols (TCP, IP, etc.)
- ❑ The **protocol** does not remember the “state/history” of a conversation
 - An action is fully specified by the command and parameters
 - The action does not depend on result of previous interactions
 - **Stateless**: Send the packet with **id=...** to destination with **ip address=...**
 - ❑ For HTTP: GET <http://www.cse.ust.hk/faculty/dlee/index.html>
 - ❑ The request will return the same page whether you issue it today or tomorrow or from any browser
 - **Stateful**: Must know previous actions to complete current command:
 - ❑ Send **next** packet; retry **previous** packet; sent to **next** address

Stateful Protocols

- ❑ Stateful protocol remembers the state of the interaction
 - Typically requires opening a channel between the communicating nodes and maintaining state variables for each channel
- ❑ In the previous examples, what states are remember?
 - Send **next** packet; retry **previous** packet; sent to **next** address


Stateless vs Stateful Protocols

Get a page given exact URL	Stateless	
Get the current page	Stateful	"current" depends on where you are
Get the "next" page	Stateful	"next" depends on where you are

- ❑ Stateful protocols are more efficient from the user or functional point of view, but stateless protocols are more **scalable** and **recoverable** (i.e., suitable for internet) **WHY?**

HTTP is Stateless

- ❑ **GET** `www.ust.hk/faculty/dlee/index.html`
- ❑ The command and URL are fully specified



- ❑ There is **no** “get next page” in HTTP
 - The “next” page (D) is hardcoded in the previous page (C)
 - `GET .../.../c.html` \Rightarrow Returns `c.html` with a link to `.../.../d.html`
 - Clicking the link `.../.../d.html` will issue `GET .../.../d.html`
- ❑ Question: does this interaction requires memory of any “state”?
- ❑ Question: What if the server or network is down?

Example: Search Engine Result Page

□ Bing's navigation bar for the result pages

BEng in Computer Engineering

ugadmin.ust.hk/curriculum_hb/1415hb/4y/pdf/13-14cpeg2.pdf · PDF file

BEng in Computer Engineering Credit(s) ... COMP 4021 Internet Computing 3 COMP


4111 Software Engineering Laboratory 3 COMP 4311 Principles of Database Design 3

Prev 1 2 3 4 5 6 7 8 Next

- Current page showing results from 31 to 40
- Prev URL:
`http://www.bing.com/search?q=4021&go=&qsn&pq=comp+4021
&sc=0-0&sp=-1&sk=&first=21&FORM=PQRE`
- Next URL:
`http://www.bing.com/search?q=4021&go=&qsn&pq=comp+4021
&sc=0-0&sp=-1&sk=&first=41&FORM=PORE`
- Question: What are the current and next state(s) of this interaction?

Keeping States in Server Application

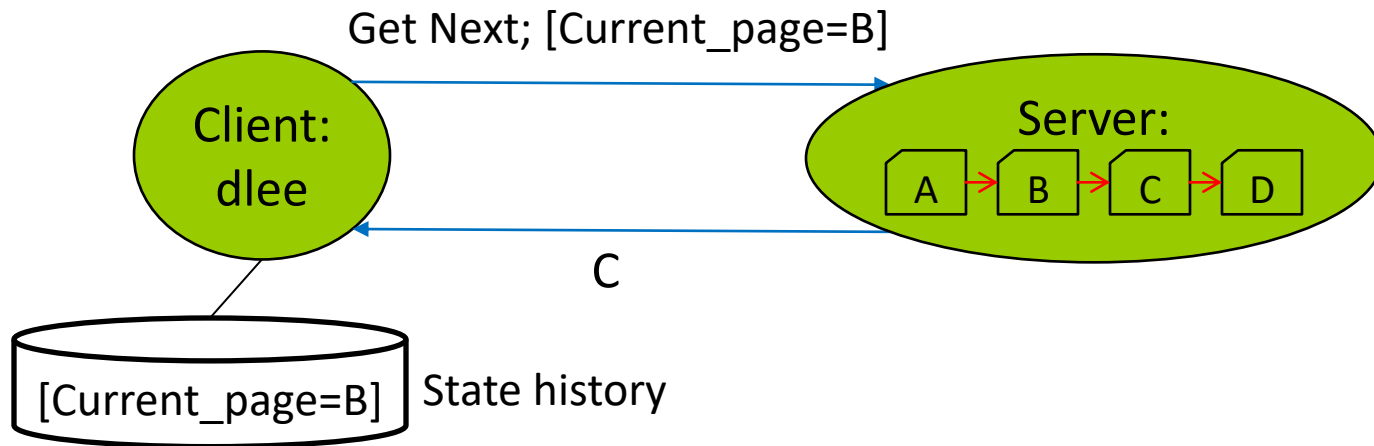
- ❑ The result of a **stateful** request depends on the “state” of the interaction (i.e., the current page)
 - The meaning of “next page” depends on what the “current page” is



- ❑ Server remembers the current state, which consists of two values [user ID, current_page], **for each conversation**
 - State: [user=dlee, current_page=B]
 - Get Next received from user=dlee => returns C

Keeping State in Client

- ❑ Instead of the server keeping the states for all interactions, the client keeps its own state; server does not remember state



- ❑ Client sends the current state to server with each request
- ❑ When the page ("C" in this example) is received, update state [Current_page=C]

Three ways to remember States

- ❑ Hardcoded in webpage (as in the Bing example)
- ❑ Stored in browser (as cookies)
 - Server application embeds JavaScript functions or issues HTTP requests to create/update/delete cookies in browser [to be covered next]
- ❑ Stored in server (e.g., in a file or backend database)
- ❑ A combination of the above

Pros and Cons of keeping states in browser or server side?

What Are Cookies?

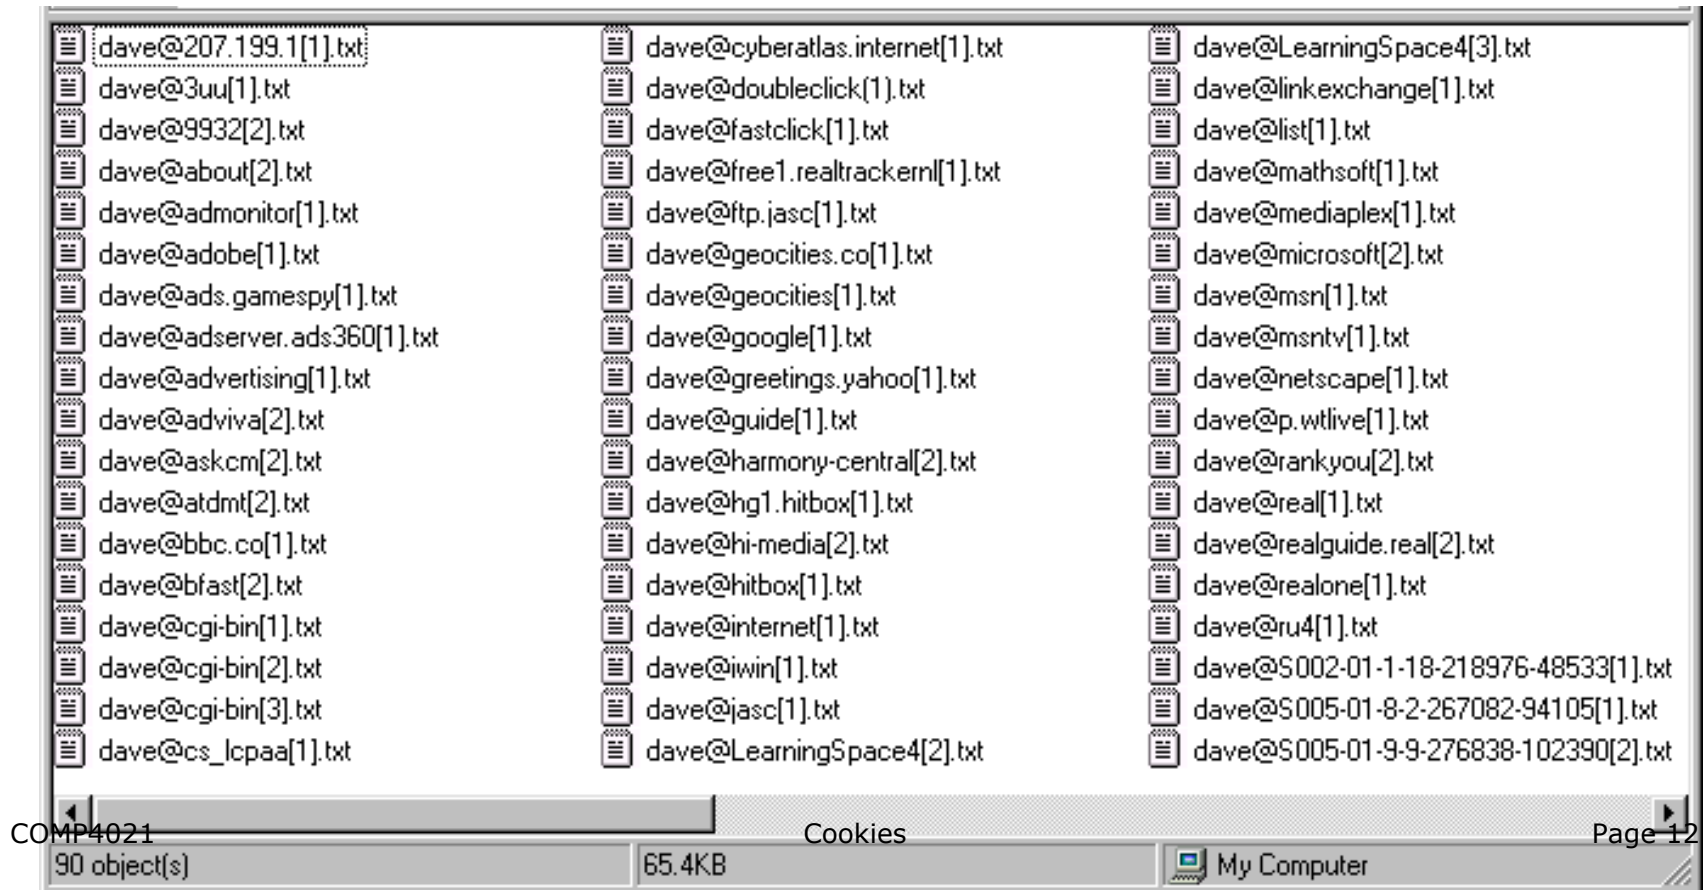
- ❑ Basically, cookies are just text strings that allow the browser/server to remember the “states” of the interaction
 - E.g., Whom am I talking to (user name), visit history, etc.
- ❑ A **website** can store cookies in a browser and later read and modify them
- ❑ If you visit the web site again, then JavaScript code in the web page can read the cookies that were stored earlier
- ❑ Not just websites – files loaded from local disk can also use cookies (If cookies are handled by client-side JavaScript)

Things Stored in Cookies

- ❑ You can store anything you want in a cookie
- ❑ Examples of things commonly stored:
 - The date/ time you last visited the web site
 - How many times you have been to the web site
 - What things you clicked on (e.g. books, etc)
 - The highest score so far – for a game

Example Cookie Data - IE

- For IE, each web site is given its own cookies file
- All cookie data for that web site is stored in that file



Example Cookie Data - IE

The screenshot shows the Internet Explorer 'Cookies' window. A list of cookies is displayed on the left, including 'dave@real[1].txt'. A 'Quick View' window is open for this specific cookie, showing its details. A red bracket highlights the 'Value' field of the 'dave@real[1].txt' cookie.

Cookie Details (dave@real[1].txt):

- Name: RNPProduct
- Value: redir_rp=PN=RealPlayer|PV=6.0.9.374|PT=FREE|DC=NS80|LI=1033|LP=en,*|CO=9|RGN=&US=84U29
- Domain: real.com/
- Path: /
- Expires: 0
- Size: 1411063808
- Attributes: 29550852, 2832372608, 29477427

The status bar at the bottom indicates '1 object(s) selected' and '152 bytes'.

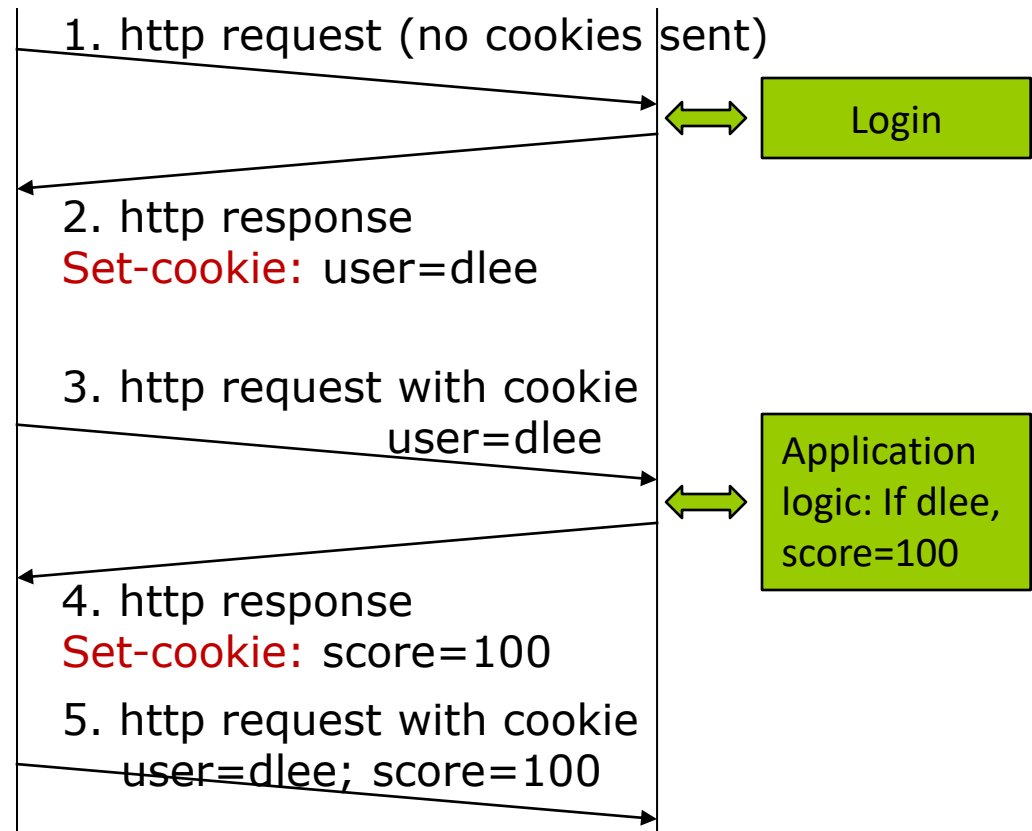
Cookie Limitations

- ❑ A browser can hold up to 300 cookies (totally 4K bytes)
 - If you have one cookie, the cookie can be 4K bytes in size
 - If you have two cookies, the average size is 2K bytes, etc.
 - If you already have a cookie, and your new cookie would make the total size exceeding 4K, the old cookie will be deleted
- ❑ There is a maximum of 20 cookies per web site
 - When you add the 21st cookie, the oldest one would be deleted
- ❑ These limitations are browser dependent, so use a smaller limit if you want to be compatible with more browsers
- ❑ 4K is not really 4K, leave room for overhead, delimiters, etc
- ❑ A cookie is not just name+value, leave room for expire, path, etc.

Browser-Server Interaction Example

- All cookies created by a website are sent to the server **with each request**, whether the server needs/wants them or not !

- Initially, there is no cookie set in browser
- Application sets a cookie (via http response header or JavaScript function)
- Browser sends request and all cookies
- Application sets new cookie
- Browser sends request and all cookies



JavaScript on Cookies

- ❑ JavaScript controls cookies through the `document.cookie` property
- ❑ You can read cookies by reading it
- ❑ You can make a cookie by changing it

Reading the Cookies

- ❑ `document.cookie` is a **string** containing all of the cookies associated with the current web site
- ❑ To see all the cookies do: `alert(document.cookie)`
- ❑ If `document.cookie` has three cookies it will have this structure:
`cookieName1=cookieValue1;cookieName2=cookieValue2;cookieName3=cookieValue3`
 - For example: `name=dave;score=8900;total_time=46`
- ❑ So you need JavaScript string handling to extract the individual data out of the string

Saving/Updating a Cookie

- ❑ Because `document.cookie` is just a string, changing it is straightforward
- ❑ For example, the following JavaScript statement sets two cookies:
`document.cookie = "name=peter;score=260";`

Example JavaScript

- Put the date/time in a cookie, and shows it in the web page

```
<head> <script>  
  var today = new Date();  
  now = today.getHours() + ":" + today.getMinutes() + ":" +  
        today.getSeconds();  
  
  document.cookie = "last_visit=" + now;  
  document.writeIn(document.cookie );  
</script> </head>
```

Same as “write” but with a newline

Cookie Expiry Time

- ❑ It is a good idea to also set a time for the cookie to 'die' (expire)
- ❑ Here a cookie is made with a specific time to expire:
document.cookie = "name=peter;
expires=Tue, 17-Mar-08 00:00:01 GMT";

Deleting a Cookie

- ❑ The way to delete a cookie is to set the date/time of the cookie to a date/time that has already finished (i.e. 1970, or one second ago, or one hour ago)
- ❑ The browser will then automatically remove the cookie
- ❑ For example:
`document.cookie = "name=peter;
expires=Thu, 01-Jan-70 00:00:01 GMT";`

Altering a Cookie

- ❑ What if you have already made a cookie, but now you want to change it?
- ❑ Cookies can be altered by simply
 1. reading the document.cookie string
 2. changing the string as appropriate
 3. copying the string back to document.cookie

Functions for Handling Cookies

- ❑ You can define functions to help with handling cookies
- ❑ The source code for a set of routines for handling cookies is shown on the next few slides
 - `setCookie()`
 - `getCookie()`
 - `deleteCookie()`

Cookie Handling - setCookie

```
function setCookie(name, value, expires, path, domain, secure) {  
    var curCookie = name + "=" + escape(value) +  
        ((expires) ? "; expires=" + expires.toGMTString() : "") +  
        ((path) ? "; path=" + path : "") +  
        ((domain) ? "; domain=" + domain : "") +  
        ((secure) ? "; secure" : "");  
    document.cookie = curCookie; }  
}
```

- This code creates a cookie (using parameters passed to it)
- The expiry time needs to be given to it in milliseconds

Advanced Cookie Properties

- path: URL path must exist in the requested resource before sending the Cookie header
- domain: Host to which the cookie will be sent. Defaults is the host of the current document location
 - A server cannot set a cookie and request it to be sent to a server with a different domain
- secure: A flag (no value) indicating the cookie will only be sent to the server when HTTPS is used

`name=dlee;path=/faculty/dlee;domain=www.cse.ust.hk;secure`

Cookie Handling - getCookie

```
function getCookie(name) {  
    var dc = document.cookie;  
    var prefix = name + "=";  
    var begin = dc.indexOf("; " + prefix); // search ";score"  
    if (begin == -1) {  
        begin = dc.indexOf(prefix);  
        if (begin != 0) return null;  
    } else begin += 2;  
    var end = dc.indexOf(";", begin);  
    if (end == -1) end = dc.length;  
    return unescape(dc.substring(begin + prefix.length, end));  
}
```

cookie not found

Use string functions to extract data from the cookie string, e.g., let's get "score":
`name=dave;score=8900;total_time=46`

Position (or Index) where ";prefix" appears in cookie string

Convert %20 to " ", etc.

Start and end positions of the name cookie in cookie string

Cookie Handling - deleteCookie

- **path** - path of the cookie (must be same as path used to create cookie)
- **domain** - domain of the cookie (must be same as domain used to create cookie)

```
function deleteCookie(name, path, domain) {  
    if (getCookie(name)) {  
        document.cookie = name + "=" +  
            ((path) ? "; path=" + path : "") +  
            ((domain) ? "; domain=" + domain : "") +  
            "; expires=Thu, 01-Jan-70 00:00:01 GMT";    }    }
```

Using the Functions

- ❑ Two examples follow
- ❑ Example 1 – A web counter
 - Each time you visit the page, it adds one to a counter stored in a cookie
- ❑ Example 2 – A name tracker
 - The name of the user is stored in a cookie and is shown every time the page is visited

Example 1 - Web Counter

- ❑ Use a cookie to count how many times someone has visited a particular web page
- ❑ The following script displays the number of times the user has visited, assuming just one person uses the browser)
- ❑ Reload the page to see the counter increment

By the way, you have been here 4 times.

Example 1 - Web Counter

```
var now = new Date();    // create an instance of the Date object
now.setTime(now.getTime() + 365*24*60*60*1000); // expires in 365 days
                                     // getTime() and setTime() work in msec
var visits = getCookie("counter");
if (!visits) {
    visits = 1;    // if the cookie wasn't found, this is the first visit
    document.write("By the way, this is your first time here.");
} else {
    visits = parseInt(visits) + 1; // increment the counter
    document.write("By the way, you have been here " + visits + " times.");
}
setCookie("counter", visits, now); // set the new cookie
```

name, value, expire

Example 2 - Name Tracker

- ❑ The following script asks the user for his/ her name, and "remembers" the input
- ❑ It then welcomes the user each time he/ she accesses the page, without asking again for the name

Welcome to this page, Dave.

Example 2 - Name Tracker


```
var now = new Date();    // create an instance of the Date object

now.setTime(now.getTime() + 365 * 24 * 60 * 60 * 1000);
var username = getCookie("username");
// if the cookie wasn't found, ask for the name
if (!username) username = prompt("Please enter your name:", "");

setCookie("username", username, now);    // set the new cookie
if (username) {
    document.write("Welcome to this page, " + username + ".");
    setCookie("username", username, now);
} else document.write("You refused to enter your name.");
```


Example 2 - Name Tracker

- Altering a cookie, e.g., allow user to change to a new name



```
function changeName() {  
    var now = new Date();  
    // cookie will expire one year later  
    now.setTime(now.getTime() + 365 * 24 * 60 * 60 * 1000);  
    username = prompt("Please enter your name:", "");  
    setCookie("username", username, now); }  

```

Set-Cookie: HTTP Response Header

- ❑ When the server sends a response to the browser, it will send a bunch of HTTP response headers in front of the content

HTTP/1.1 200 OK

Date: Mon, 27 Jul 2009 12:28:53 GMT

Server: Apache/2.2.14 (Win32)

Last-Modified: Wed, 22 Jul 2009 19:15:56 GMT

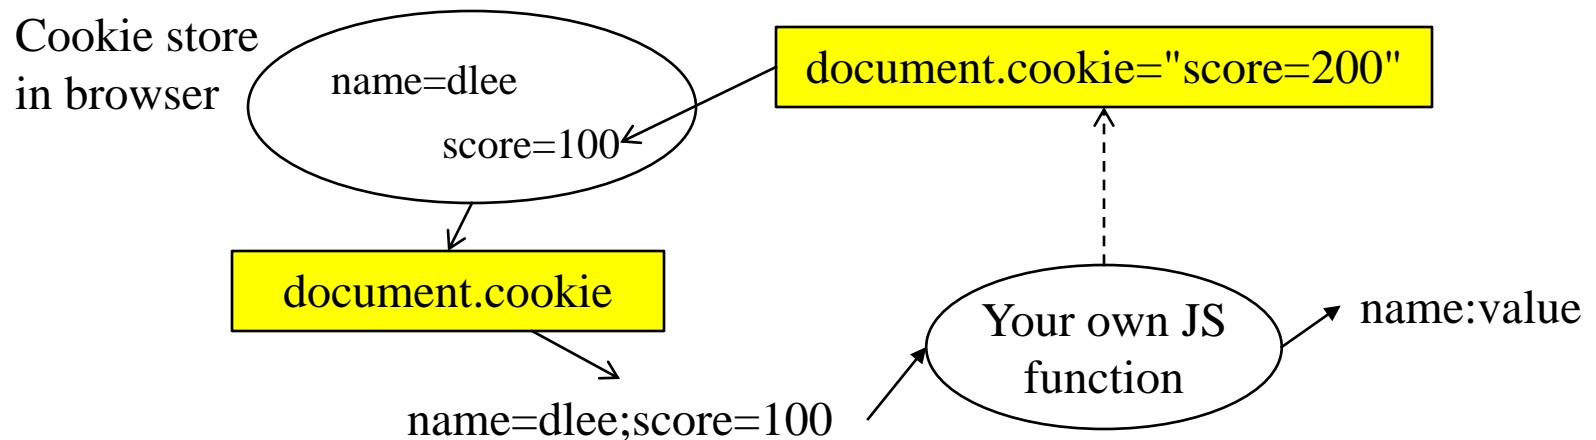
... ..

Set-Cookie: <cookie-name>=<cookie-value>;Expires=<date>

- ❑ Set-Cookie sets cookies in the browser
- ❑ Set-Cookie is a response header, so it must be in the header section of the response (i.e., before any content)

Important Clarification

- Cookies are stored in a cookie store in the browser
- How the cookies are stored in the browser is browser dependent, but they are stored in disk files for persistence
- In fact, the cookie store implementation is irrelevant to most end users, since DOM provides a uniform interface for creating, updating, deleting and reading cookies:



The Cookie Interface

- **document.cookie** returns a string containing the name/value pairs of all cookies in the cookie store, e.g.,
`name=dlee;score=1300;last_visit=22:55:18;`
- **document.cookie** is the only interface for getting cookie values
 - You need to writing your own JS code to get cookie values and implement all error checking functions
 - E.g., to get the value of the “score” cookie, check if "score" exists in `document.cookie`; if yet, get and return the value; if not, return NULL
 - See `getCookie()` function
- **document.cookie** is not a regular string -> See next slide

Creating and Updating Cookie

- To create a new cookie :
 - `document.cookie = "grade=A"`
 - Since grade is not an existing cookie, it is created; document.cookie is:
`name=dlee;score=1300;last_visit=22:55:18;grade=A`
- document.cookie is **not a regular string**, if it is, "grade=A" should replace the old value and document.cookie becomes "grade=A"; now "grade=A" is appended into document.cookie
- Update an existing cookie:
 - `document.cookie = "score=1000"`
 - Since score exists, its value is updated; document.cookie becomes:
`name=dlee;last_visit=22:55:18;grade=A;score=1000`
 - Order of the list is typically based on last-modified-time; latest is last

Wrong Cookie Update

- Although document.cookie returns:
 - `name=dlee;score=1300;last_visit=22:55:18;grade=A`
- If you want to update the values of all four cookies, you cannot do:
 - `document.cooke="name=john;score=1200;visit=00:00;grade=A+"`
- Only the first name=value is recognized, but the rest is ignored
- If original cookie is:
`name=dlee;score=1300;last_visit=22:55:18;grade=A`
- document.cookie becomes:
`name=john;score=1300;last_visit=22:55:18;grade=A`

Correct Way to Update Cookies

- If you want to change all four cookies, do it one by one (you can also add expires or path attributes):
document.cooke="name=john";
document.cookie="score=1200;expires=Thu, 01-Jan-70 00:00:01 GMT";
document.cookie="visit=00:00";
document.cookie="grade=A+"

Take Home Message

- Cookies is a quick way of storing information (or states) about an interaction
 - Number and size of cookies are limited, but enough for most applications
 - Storing cookies as strings make cookies hard to access and maintained but **very easy to send**
- Cookies are not just for displaying friendly messages to the users (e.g., greetings), they are required in many web applications, e.g. gmail requires cookies to work



Cookies are disabled

Your browser has cookies disabled. Make sure your cookies are enabled and try again. [Learn more](#)

Next

Take Home Message (Cont.)

- Cookies are stored in browsers by a web application
 - Switching browsers will lose the cookies
 - Users cannot change the cookies (or understand their meaning), but can delete or refuse to accept them
 - When a user accesses a website, all of the website's cookies (if any) will be sent to the website
- Application can set cookies by (i) triggering a JavaScript function or (2) using Set-Cookie HTTP Response Header